

TOWARDS MULTILINGUAL EVALUATIONS OF KNOWLEDGE FOR
LARGE LANGUAGE MODELS

Bryan Li

A DISSERTATION

in

Computer and Information Science

Presented to the Faculties of the University of Pennsylvania

in

Partial Fulfillment of the Requirements for the

Degree of Doctor of Philosophy

2025

Supervisor of Dissertation

Chris Callison-Burch, Professor of Computer and Information Science

Graduate Group Chairperson

Anindya De, Associate Professor of Computer and Information Science

Dissertation Committee

Lyle Ungar, Professor of Computer and Information Science

Duncan Watts, Stevens University Professor of Computer and Information Science

Dan Roth, Eduardo D. Glandt Distinguished Professor of Computer and Information Science

Colin Cherry, Research Scientist, Google

Mohammad Sadegh Rasooli, Research Scientist, Meta

TOWARDS MULTILINGUAL EVALUATIONS OF KNOWLEDGE FOR
LARGE LANGUAGE MODELS

COPYRIGHT

2025

Bryan Li

Dedicated to my grandfather, whose love of learning and travel has inspired my academic pursuits.

献给我的外祖父，他对学习和旅行的热爱，启迪了我的学术追求。

ACKNOWLEDGEMENT

My foremost gratitude goes to my advisor Chris Callison-Burch. His supervision over the past five years helped me grow from a student with scattered research interests into a scholar pursuing scientific discovery through research. Chris’s mentorship philosophy is to his students both with ample academic freedom, as well as with hands-on guidance whenever we need it. I am particularly grateful for his guidance on refining my research presentation skills, allowing me to present to any audience level, from a layperson, to a general computer scientist, to the most discerning dissertation committee member.

Speaking of committee members, mine have provided invaluable guidance to this thesis, and also to my development as a researcher. Lyle Ungar taught my first and last courses at Penn, and I appreciate the insightful questions he brings to every discussion. I worked closely with Duncan Watts when starting at Penn, and I am grateful to continue our collaborations on computational methods to understand the information ecosystem. Dan Roth has kept me cognizant of both the opportunities and limitations of artificial intelligence towards reasoning about the world. I worked with Colin Cherry during an internship at Google, where he guided our project through twists and turns before the final cohesive package. I worked with Mohammad Sadegh Rasooli while he was a postdoctoral researcher, and he provided me advice on research, programming, and work-life balance alike.

I appreciate the guidance from many other mentor figures. Julia Hirschberg and Kathleen McKeown kindly brought me into the world of NLP research during my undergraduate and masters years at Columbia. Wei Xu, Mark Liberman, Lara Martin, Aleksey Panasyuk, and Homa Hosseinmardi provided guidance at various points of the PhD program.

During my PhD, I was funded by four grants: LWLL, BETTER, UNCOVER, and HIATUS. It was a pleasure to have worked with many great people through them. I also pursued five internships in five cities: Johns Hopkins from Philadelphia, Microsoft from Irvine, Amazon from Santa Clara, Google from Montréal, and Johns Hopkins from Baltimore. I am so glad for the friends I have

made in all of these places. I thank the undergraduate and masters students I was able to mentor throughout the years. And I thank my many collaborators for their contributions to the works in this thesis.

My PhD life was never lonely thanks to members of the Penn NLP community, past and present: Samar, Yue, Ajay, Liam, Artemis, Andrew, Delip, Harry, Veronica, Aditya, Rebecca, Grace, Maxine, Daphne, Reno, Yifei, Chaitanya, Allen, Josh, Shreya, Jeffrey, Mark, Eleni, and Marianna. I also thank friends from the larger Penn community: Fengjun, Alyssa, Weiqiu, Hita, Anish, Junyao, Xiayan, Josh, Upasana, Edward, and many others. I am very glad for my friends from college and other walks of life, and the many trips we have been on near and far.

I am forever grateful to my parents, grandparents, and my sisters, Katie and Megan, for their unwavering support in everything I do, this PhD journey included. Thank you for giving me a sense of home wherever I am, whether I'm actually at home in Irvine, at my sisters' apartment and cat sanctuary in New York, at school in Philadelphia, or at a conference halfway across the world.

ABSTRACT

TOWARDS MULTILINGUAL EVALUATIONS OF KNOWLEDGE FOR LARGE LANGUAGE MODELS

Bryan Li

Chris Callison-Burch

Contemporary language models (LMs) understand and support interaction in dozens of languages, and have significant potential to empower equitable information access for users across the world. Existing multilingual evaluations have focused on tasks which challenge whether facts learned in one language can be reproduced in the same language. This limitation has led to a narrow view of multilingual knowledge to simpler tasks which can be translated across language. In contrast, the real-world distribution of knowledge often involves tasks where across languages, topic coverage differs and cultural perspectives conflict. In this thesis, I study several multilingual knowledge-intensive tasks, and how LMs interact with the *parametric knowledge* encoded internally in their parameters, and with the *contextual knowledge* presented externally during interaction.

The first part of this thesis investigates parametric knowledge, introducing two benchmarks which challenge whether LMs can generalize their knowledge across languages. First, I introduce a multi-task, multilingual complex reasoning benchmark, and find that LMs which reason well in English struggle far more in other languages. I address this with methods leveraging program code to bridge LM’s competencies in multilinguality and in reasoning. Second, I introduce a benchmark of territorial disputes. By posing these queries in different languages (e.g., asking on Crimea in Russian or Ukrainian), I reveal the issue of *cross-lingual robustness*, where an LM’s response to a query is inconsistent depending on the language of interaction. I then show that lightweight methods of leveraging program code and persona-based prompting can mitigate these issues.

The second part of this thesis investigates contextual knowledge, focusing on the retrieval-augmented generation (RAG) paradigm, which provides an LM with contextual knowledge retrieved from exter-

nal knowledge-bases (KBs). I begin with a study on the machine translation task, which finds that contextual knowledge is most effective when presented as stylistically similar demonstrations. I then return to the territorial disputes task. While RAG provides relevant information, it also invites an array of culturally and language-informed perspectives. Thus, the composition of KBs determines what information can be retrieved. Experiments find that multilingual retrieval effectively improves robustness over retrieval only in query languages. I further expand to the multi-source setting, collecting a large-scale dataset of articles from state media outlets. I similarly find that for retrieval, multi-source improves robustness over single-source, and furthermore, that foreign-targeted state media causes far less robust responses than domestic-targeted ones.

In summary, this thesis contributes multilingual evaluations of LMs on knowledge-intensive tasks, and explore contextually-informed methods to improve them. We ultimately highlight the need for multilingual LMs that can navigate, and assist users in navigating, the real-world distribution of knowledge across languages and sources.

TABLE OF CONTENTS

ACKNOWLEDGEMENT	iv
ABSTRACT	vi
LIST OF TABLES	xi
LIST OF ILLUSTRATIONS	xiii
CHAPTER 1 : INTRODUCTION	1
1.1 Thesis Statement	6
1.2 Outline of Thesis	7
1.3 Publications Presented	11
CHAPTER 2 : BACKGROUND	14
2.1 Multilinguality	14
2.2 Large Language Models	17
2.3 Knowledge and Robustness of LLMs	19
2.4 Multilingual Benchmarks	26
2.5 Summary	32
CHAPTER 3 : EVALUATION OF COMPLEX REASONING KNOWLEDGE ACROSS LAN- GUAGES	33
3.1 Preliminaries	36
3.2 The xSTREET Multilingual Complex Reasoning Benchmark	39
3.3 Training Method: Indirect Supervision from Multilingually Commented Code	40
3.4 Inference Method: Structured Reasoning through Code Prompts	42
3.5 Experimental Setup and Results	45
3.6 Analysis	50
3.7 Conclusion	52

CHAPTER 4 : EVALUATION OF GEOPOLITICAL KNOWLEDGE THROUGH TERRI-	
TORIAL DISPUTES	54
4.1 Dataset Collection for BORDERLINES	57
4.2 Task Design	59
4.3 Cross-lingual Robustness	60
4.4 Experimental Setup and Results	63
4.5 Qualitative Analysis	69
4.6 Conclusion	72
CHAPTER 5 : RETRIEVAL STRATEGIES FOR DOMAIN-ADAPTED TRANSLATION .	73
5.1 Leveraging Domain Knowledge	77
5.2 Experimental Setup and Results	82
5.3 Analysis	85
5.4 Conclusion	90
CHAPTER 6 : MULTILINGUAL RETRIEVAL AUGMENTED GENERATION FOR GEOPOL-	
ITICS	91
6.1 Preliminaries	93
6.2 Experimental Setup and Results	97
6.3 Qualitative Analysis	106
6.4 Conclusion	109
CHAPTER 7 : PERSPECTIVES FROM FOREIGN AND DOMESTIC STATE MEDIA . .	110
7.1 Resources for State Media	111
7.2 Dataset Analysis	114
7.3 Experimental Setup and Results	117
7.4 Conclusion	120
CHAPTER 8 : CONCLUSION	122
8.1 Future Work	123

APPENDIX A : APPENDIX FOR CHAPTER 3: COMPLEX REASONING	127
APPENDIX B : APPENDIX FOR CHAPTER 4: BORDERLINES	132
APPENDIX C : APPENDIX FOR CHAPTER 5: DOMAIN-ADAPTED MT	138
APPENDIX D : APPENDIX FOR CHAPTER 6: BORDIRLINES	141
APPENDIX E : APPENDIX FOR CHAPTER 7: STATE MEDIA	153
BIBLIOGRAPHY	155

LIST OF TABLES

TABLE 3.1	Statistics for the xSTREET benchmark test split.	39
TABLE 3.2	Statistics for the TCC dataset, used for finetuning.	40
TABLE 3.3	Results for 3 non-complex multilingual reasoning tasks.	49
TABLE 4.1	Demographics of the majority language, majority religion, and population of several selected countries.	58
TABLE 4.2	Excerpt from the BORDERLINES table.	58
TABLE 4.4	Three territories from BORDERLINES and their multilingual query sets (MQS).	59
TABLE 4.3	BORDERLINES table statistics.	59
TABLE 4.5	Formulas for concurrence score (CS) metrics.	62
TABLE 4.6	Results on BORDERLINES for different models.	65
TABLE 4.7	The system prompts used for our prompt modification experiments with GPT-4.	67
TABLE 5.1	Domain-adapted MT results for the knowledge strategies and knowledge sources settings.	84
TABLE 5.2	COMET22 Results for the study on demonstration generation, using Gemma and Gemini.	89
TABLE 6.1	For each IR mode, the retrieval languages L for a query q_l	94
TABLE 6.2	Spratly Islands viewpoint judgment results for GPT-4o across different retrieval modes.	106
TABLE 6.3	Crimea viewpoint judgment results across retrieval modes.	108
TABLE 7.1	Excerpted entries from the state media articles table.	112
TABLE A.1	Full results for STREET and xSTREET, on all models, tasks, prompt formats, and languages.	129
TABLE A.2	Per-language results for xSTREET, on all models, tasks, prompt formats, and languages.	129
TABLE B.1	Full results on BORDERLINES for different models.	134
TABLE B.2	GPT-4 responses for the query on the geopolitical status of Crimea, for the 4 prompting strategies. Translations are provided in the grey rows.	136
TABLE B.3	GPT-4 responses for the query on the geopolitical status of Taiwan.	137
TABLE D.1	Overall statistics for the retrieved documents from the BORDIRLINES dataset for each of the 49 languages.	141
TABLE D.2	Results for all models and IR modes studied in BORDIRLINES	148
TABLE D.3	Prompt templates used in Chapter 6 on BORDIRLINES	149
TABLE D.4	List of prompt templates used for BORDIRLINES	150
TABLE D.5	Performance of LLM annotations vs. human annotations for relevance and territory/viewpoint, across 3 selected languages.	150
TABLE D.6	Spratly Islands viewpoint judgments for direct and citation formats.	150
TABLE D.7	Spratly Islands citation responses for GPT-4o-mini in en_only mode.	151

TABLE D.8 Spratly Islands citation responses for GPT-4o-mini in `en_only` mode, continued. 152

LIST OF ILLUSTRATIONS

FIGURE 1.1	An illustration of an LLM struggling to transfer knowledge across languages, from Chinese to English and vice versa.	5
FIGURE 1.2	Overview of the chapters in this thesis.	7
FIGURE 1.3	Illustration of two complex reasoning problems for math and science.	8
FIGURE 1.4	Illustration of an entry from BORDERLINES.	9
FIGURE 3.1	Illustration of two complex reasoning problems for math and science. This LM reasons and give the correct answer when queries are written in English, but fails when the same underlying queries are written in French or Chinese. . . .	34
FIGURE 3.2	An overview of our methods to improve multilingual complex reasoning. . . .	35
FIGURE 3.3	Example entries from xSTREET, in English and in Spanish	38
FIGURE 3.4	The data augmentation approach used to create the TCC dataset.	41
FIGURE 3.5	Prompting formats for the xSTREET tasks.	46
FIGURE 3.6	GPT-3 results for STREET and xSTREET tasks.	48
FIGURE 3.7	BLOOMZ-based model results for STREET and xSTREET tasks.	48
FIGURE 3.8	Results on ARC subtask of xSTREET for the ablation study on the role of code comments. All results use the SIM code prompt format.	50
FIGURE 4.1	An example entry from BORDERLINES for Ceuta, disputed between Morocco and Spain.	56
FIGURE 4.2	Excerpt of the “Africa” table from the Wikipedia article on territorial disputes.	57
FIGURE 4.3	Illustration of comparisons made for the CS metrics.	61
FIGURE 5.1	Illustration of a source sentence and three possible translations.	74
FIGURE 5.2	Illustration of the main MT settings for knowledge strategies and sources. . . .	76
FIGURE 5.3	Illustration of our process to decompose the contributions of retrieved demonstrations into style and terminology.	86
FIGURE 5.4	MT results comparing zero-shot and external retrieval (ret terms) to terms from demonstrations and style from demonstrations.	87
FIGURE 5.5	Results for the ablation on generation strategies.	90
FIGURE 6.1	Illustration comparing two RAG settings, over a query-language KB vs. over a multilingual KB.	92
FIGURE 6.2	Illustration of data collection for a BORDERLINES entry on the territory “Ceuta”.	95
FIGURE 6.3	Results for the concurrence score (CS) metrics.	100
FIGURE 6.4	Citation rate vs. inclusion rate for queries for each language in the <code>rel_langs</code> setting, using GPT-4o.	103
FIGURE 6.5	Citation distribution across three RAG modes.	103
FIGURE 6.6	Proportion of document languages retrieved per query language, using OpenAI and M3 embeddings in the <code>rel_langs</code> mode.	104
FIGURE 6.7	Results for the ablation on filtering documents for relevance and non-relevance.	105
FIGURE 6.8	Relevance and viewpoint judgments for the Spratly Islands case study.	106
FIGURE 6.9	Relevance and viewpoint judgments for the Crimea case study.	108

FIGURE 7.1	# webpages for top 32/100 countries in English Fineweb.	114
FIGURE 7.2	# webpages for top 32/85 countries in Fineweb-2.	115
FIGURE 7.3	Illustration of the 6 retrieval settings, which vary the sources, and the perspectives.	118
FIGURE 7.4	Concurrence score (CS) metrics over different retrieval settings, which vary the sources, and the perspectives.	119
FIGURE 7.5	Case study on Crimea for RAG over foreign state media, written in English. .	120
FIGURE A.1	Case study of one GSM8K problem in English. The top row is for SIM prompts, while the bottom is for linearized.	130
FIGURE A.2	Results on GSM8K, AQUA_RAT, AR_LSAT subtasks for BLOOMZ-based models.	131
FIGURE C.1	Prompt for zero-shot MT.	139
FIGURE C.2	Prompt for MT with <i>demonstrations</i>	139
FIGURE C.3	Prompt for MT with <i>terminologies</i>	139
FIGURE C.4	Prompt for MT with <i>style from demonstrations</i>	139
FIGURE C.5	Prompt for synthetic <i>demonstration generation</i>	140
FIGURE C.6	Prompt for synthetic <i>terminology generation</i>	140
FIGURE C.7	Prompt for <i>terminology extraction</i> from source-target text pairs.	140
FIGURE D.1	A retrieved document for the Spratly Islands (doc_id: 27401_p38).	145
FIGURE D.2	Instructions and examples for the human annotation task on relevance and viewpoint judgment.	146
FIGURE D.3	Example HIT for the territory “Moyo District”.	147
FIGURE E.1	Prompt for obtaining URLs for the assets of a state media company. Four few-shot exemplars are provided; one is shown here.	153
FIGURE E.2	Number of webpages per Common Crawl split, English Fineweb.	154

CHAPTER 1

INTRODUCTION

Knowledge about the world, from facts, to information, to beliefs, is captured through the medium of *language*. Language enables the transfer of knowledge between humans, within societies, and even across vast timescales and distances. Languages change greatly over time, with the concerns of its speakers, and the context of the world around them. However, it is fascinating that language preserves its knowledge-capturing essence through all of its manifestations. Even with the empire behind the Rosetta Stone lost to the sands of time, researchers were able to read it, and understand not only the text itself,¹ but also the ancient Egyptian language behind it.

Indeed, while every language is strikingly different at a surface-level, to quote Akmajian et al. (2017), “all known languages are at a similar level of complexity and detail,” Any piece of knowledge can be expressed as validly in one language as in another. As such, are over 7000 languages are spoken throughout the world today (Eberhard et al., 2015). Linguistic diversity is both a contemporary and historical reality. And beyond that, every human possesses the capacity to learn any language, and so access the knowledge captured in any of them.

Now, for a human to learn a new language often seems like a daunting task. However, try to remember, if you can, the amount of effort that went into the acquisition of your first language. Starting from a baseline of smiling, crying, and pointing, baby you learned from environmental stimuli how to express first single words, then sentences, and finally full-fledged thoughts. Within the first year of life, an infant has undergone a significant amount of language learning, leading to further development in the years to come (Dehaene-Lambertz et al., 2006). And learning new languages is possible at any point in one’s lifetime.

The majority of the world’s population are multilingual, navigating multiple languages in daily

¹For historical context, the Rosetta Stone was created in 196 BCE, and discovered in 1799. It contains three translated inscriptions: Ancient Greek, and Ancient Egyptian in two scripts, demotic and hieroglyphic. Scholars were able to understand the Greek text by 1803, but it took until 1822 to translate the hieroglyphic script into French. The first English translation was published in 1858, by the Philomathean Society of the University of Pennsylvania. Its text decrees the reign of a new leader, outlining his coronation and his first official acts.

life (Tucker, 1999). Even those who identify as monolingual are quick to absorb concepts from other languages. In the United States, you might order a “caffè latte” and “croissant” for breakfast, a “baozi” and “boba tea” for lunch, and a “burrito” and “lassi” for dinner.² More abstract concepts also transcend linguistic boundaries. You might shrug off setbacks with “c’est la vie” or greet an old friend with “long time no see”, which is a word-for-word translation from Chinese “好久不见.”

Across the uncountable infinity of knowledge, and co-existing with these translatable concepts, are concepts which cannot be so easily translated across languages. A language captures the interests and essences of the people who speak them. This leads to knowledge available in one language are often quite distinct from that of others. Specific information on a SEPTA transit system in Pennsylvania is likely to be found in English, and similarly, this applies for the Shenzhen metro system and in Chinese. Furthermore, most languages have resources on William Shakespeare, but it is assuredly in English that the most detailed account of his life and work are to be found. That a multilingual human is able to reconcile disparate information without much issue is a testament to our intelligence capabilities.

All this discussion is to emphasize the key role multilingual knowledge plays in how humans utilize natural language. Thus, in our quest to build language models (LMs) with human-like capabilities to understand the world around us, then we should also develop LMs that can operate in multiple languages. A developer of an advanced chatbot would like to adapt their system for different linguistic communities. And users within these communities would like to access information about current events, such as a global sporting event or a global pandemic, in their language and in real time. There has thus been a rich body of natural language processing (NLP) research on multilinguality, from the inception of the field to the present.

Before moving on to this thesis, let us take a moment to situate ourselves in the current technological milieu. LLM-based chatbots such as ChatGPT (OpenAI, 2022) have been adopted worldwide by end users, for tasks from creative writing, to translation, to homework help. The phrases “LLM” (large language model) and “generative AI” have become part of the popular lexicon. And public

²These food items have been sampled from the author’s own diet during graduate school.

discourse has heralded the soon-transformative potential of such AI technologies into changing the way that societies, and individuals within them, operate.

We now zoom into the present state of multilingual research. Taking a quick inventory of conferences run by the Association for Computational Linguistics, we see that since 2020, there have been 16 conferences that have *Theme Tracks*, which means the organizers are taking a special interest in such submissions. 4 out of 16 theme tracks directly concern multilinguality³. There are further many workshops, talks, and other events in the field, which shows that on the one hand, multilinguality remains a primary area of interest for NLP researchers. On the other hand, there has been a widespread perception that current LLM developments are leaving most of the world’s languages behind. This has been discussed in articles from the New York Times (Ruberg, 2024), the World Economic Forum (Chhabria and Hamill-Stewart, 2024), and the Center for Democracy and Technology (Bhatia and Nicholas, 2023). This perception of the over-emphasis of research on a handful of privileged languages exists within the NLP community as well; for example, it was discussed by a panel of global researchers at the special session on multilinguality at the 2024 Conference on Language Modeling (COLM) in Philadelphia.

The frontiers of LLM development has been largely specified by English-centric research, and multilinguality is a primary, but parallel concern. We illustrate this by considering several open-weight LMs: LLaMA 3 (Llama Team, 2024), Mistral 1 (Jiang et al., 2023), and Gemma 1 (Gemma Team, 2024). Upon first release, these models were released with English-only support. The follow-up LLaMA 3.1, Mistral 2, and Gemma 2 models added multilingual support. Still, such an English-centric approach results in LMs which by construction under-perform in most other languages.

Fortunately, the tides have shifted, and open-weight LMs in 2025 have been multilingual upon release. Gemma 3 states “We particularly focus on multilinguality in this version” to support 140+ languages (Gemma Team et al., 2025). Llama 4 states “it enables open source fine-tuning efforts by pre-training on 200 languages” (Llama Team, 2025). Qwen 3 states it “amplifies its

³These are “Language Diversity: from Low-Resource to Endangered Languages” (ACL 2022), “Large Language Models and Regional/Low-Resource Languages” (AAACL 2023), “Languages of Latin America” (NAACL 2024), “NLP in a Multicultural World” (NAACL 2025).

potential for deployment in global use cases and international applications” with its support of 119 languages (Yang et al., 2025). It seems that multilinguality, along with coding, reasoning, and multimodal, has joined the frontier of AI development.

Current LMs substantiate their improved multilinguality through their progress on *benchmarks*: viewed from this lens, the performance gap between languages has narrowed considerably for the latest LLMs. However, these improvements have not translated to major gains in the real-world utility of multilingual LMs. To investigate the limitations of existing evaluations, we broadly classify benchmarks below (for a detailed survey, refer to Chapter 2):

Translation-based benchmarks A popular approach is to take an existing benchmark in English, and translate it into many languages, ideally though not always with professional translators. For example, the LLaMA, Qwen, and Gemma reports all consider the MMMLU benchmark, a suite of language tasks formulated as multiple-choice questions, and translated into 14 languages OpenAI (2024a). Averaged accuracies across languages are between 70-85%, and individual scores are also in that range, aside from Yoruba and Swahili’s lower scores, due to their low-resource nature. While translation is a useful starting point, from a starting point of a single language leads to a limited coverage of knowledge relevant to other languages.

Cultural benchmarks Most NLP models have created by Western-aligned researchers, which results in LLMs that reflect Western-centric values even in other languages. To quantify the extent, many benchmarks have been created to assess cultural biases in LMs (Yin et al., 2022; Naous et al., 2024; Myung et al., 2024). For example, on the query “What color is a traditional wedding dress?”, the correct answer is “white” in English and “红色” (red) in Chinese.

Factual benchmarks Differing information access across languages is a major issue in the real world. Motivated by this, researchers have developed benchmarks to evaluate whether LMs can recall factual knowledge across languages (Asai et al., 2022; Qi et al., 2023).

From these benchmarks, we observe that prior work has not sufficiently challenged whether LLMs can transfer more complex knowledge across languages. As a motivating example for this issue addressed in this thesis, Figure 1.1 illustrates a cross-lingual knowledge failure case for the Qwen 3



Figure 1.1: An illustration of an LLM struggling to transfer knowledge across languages, from Chinese to English and vice versa.

LLM,⁴ which was trained with strong support for both Chinese and English. The model has seen information SEPTA and Shenzhen metro systems, and answers correctly in the respective languages of English and Chinese. However, it fails to transfer this knowledge across languages, and gives the wrong transfer station for the reversed language cases. In such cases, a visiting tourist to either city will miss their transfer, have to figure out a new route, and show late to their event. They will be less inclined to use the system in the future, losing trust in both the transit navigation and the general capabilities of the system.

While the example above was an inconvenience, we can imagine real-world scenarios where such failures could prove to be far more costly. Consider the case of disputed territories. In the current decade, these disputes have flared up into military conflicts from Eastern Europe to the Middle East. Politicians argue back and forth on “who rightfully owns the land”. Commentators speculate which territories may see conflict in the near future, perhaps in South Asia, or in Arctic North

⁴This experiment was performed in August 2025; responses may vary due to API updates or model’s non-deterministic responses.

America. There are many divisions within viewpoints, and we should not want LMs that amplify these divisions, depending on how different users interact with them.

The disputed territories task is of special interest to this thesis because of several characteristics:

- Real-world significance, as discussed above.
- There is no right answer which can be memorized.
- Responses depends on one’s perspective informed by cultural, political, and language concerns.
- Situations change over time, enabling updated evaluations over time.
- It fully challenges LMs’ ability to navigate knowledge across languages.

Considering prior work, we see that they are lacking in one or more of these characteristics. Translated benchmarks fail to cover scenarios which are significant to a given language community. Cultural or factual benchmarks can be addressed by adding those entries to LM training data. And nearly all prior benchmarks do not adequately address the last characteristic, instead evaluating a single language or two for each entry. We are thus motivated to study tasks where prerequisite knowledge spans multiple languages. The main contributions of this thesis are:

- We design multilingual evaluations on knowledge-intensive tasks of reasoning and geopolitics, which respectively address language-agnostic and language-aware of knowledge.
- To evaluate LM responses to the same underlying query, specified in different languages, we formalize the notion of *cross-lingual robustness* in three attributes.
- We show that we can elicit better responses from LMs through simple grounded strategies, helping bridge the cross-lingual knowledge gap.
- We further explore how retrieval-augmented LMs navigate externally retrieved knowledge. We investigate how the composition of retrieval databases from various languages and sources affects responses, finding that diverse retrieval databases can improve robustness.

1.1. Thesis Statement

This thesis argues that evaluations of multilingual LLMs must consider how these systems utilize knowledge across languages. This important given the real-world distribution of knowledge, which




Figure 1.2: Overview of the chapters in this thesis. We evaluate LM’s ability to utilize knowledge across languages in both parametric and contextual settings. We identify LLM’s issue with cross-lingual robustness, and address it with contextual strategies, as well as varying the composition of retrieval contexts.

has different coverage and perspectives across languages. To that end, I design benchmarks which challenge LLM’s ability to *transfer their parametric knowledge* learned in one language to others, and perform studies on LLMs’ ability to *navigate contextual knowledge* retrieved from different languages and different sources. I formalize the concept of *cross-lingual robustness* in several attributes, measuring how sensitive LLMs’ responses are to the linguistic context. I then explore strategies to elicit more robust responses by incorporating context, such as program code or prompting personas. Finally, I consider the paradigm of retrieval augmented generation (RAG), studying how the composition of retrieval KBs across languages and sources affects the robustness of LM responses.

1.2. Outline of Thesis

Figure 1.2 provides an overview of the structure of this thesis.

Chapter 2 surveys the existing literature for relevant work. We first review the history of multilinguality in NLP, noting it has played a prominent role in advancing the field. We next consider several theories from the linguistics literature. We then discuss contemporary LLMs, focusing on multilingual concerns. We cover the knowledge of LMs through parametric knowledge and contextual knowledge, then robustness of LLMs through knowledge consistency. Finally, we survey the landscape of existing multilingual benchmarks.

Chapter 3 studies the capabilities of LMs to perform complex reasoning in multiple languages.




Figure 1.3: Illustration of two complex reasoning problems for math and science. This LM reasons and give the correct answer when queries are written in English, but fails when the same underlying queries are written in French or Chinese.

Several tasks are illustrated in Figure 1.3. Through this, we begin our investigation into how well LLMs can generalize reasoning knowledge they acquired in English to other languages. We thus introduce **xSTREET**, a multilingual structured reasoning and explanation dataset that covers four tasks across six diverse languages (Li et al., 2024a). The notable drop in non-English performance on this benchmark shows LMs are not cross-lingually consistent reasoners.

To remedy this gap, we turn to the hypothesis that code improves reasoning of LLMs. We show this hypothesis extends across languages through two methods. For inference-time, we design a structured prompting format that incorporates step-by-step code primitives to derive new facts and find a solution. For training-time, we augment a program code dataset by translating its comments while keeping code as-is. Notably, both of our methods improve reasoning performance of LMs for all languages, including English.

Chapter 4 studies the geopolitical knowledge of LMs through the task of territorial disputes. While a reasoning query is language-agnostic, a geopolitical query is highly sensitive to language, cultural, and political context. We introduce **BORDERLINES**, a benchmark of 251 disputed territories, each of which is associated with queries in the languages of the involved countries (49 languages




Figure 1.4: Illustration of an entry from BORDERLINES. By querying an LLM about a disputed territory in different languages, we see that responses vary based on language.

total) (Li et al., 2024b). Figure 1.4 provides an illustration of the task.

Given this issue of LMs’ differing responses to the same underlying query, we formalize the concept of *cross-lingual robustness* in three attributes: *factuality*, *consistency*, and *language fairness*. We then propose a suite of accuracy-based metrics on BORDERLINES to evaluate LMs’ performance on each attribute. Finally, we explore several prompt modification strategies, which show that simple prompting cues can amplify or mitigate geopolitical bias. This work reveals a major issue with LMs which have been trained on multilingual datasets: that knowledge may be recalled from differently languages depending on interaction-time cues. The implications are that current LMs may reinforce different user’s implicit preferences for contentious issues, amplifying divisions in society.

Chapter 5 begins our investigation into how LMs integrate their parametric knowledge with contextual knowledge retrieved from external knowledge-bases (KBs). This paradigm is known as retrieval-augmented generation (RAG). Here we consider the task of machine translation (MT) in specialized domains (Li et al., 2025b). Our work proceeds with the insight that LMs have seen a ton of knowledge on different domains and languages. Therefore, we investigate the value of retrieving knowledge from external resources versus generating it from an LMs’ own parametric knowledge. We further investigate the effectiveness of formatting knowledge as terminologies versus as exemplars. We find that generation of exemplars can bootstrap better domain-adapted MT performance from LMs. Still, contextual knowledge provides the largest gains. Additional experiments identify

that these gains largely stem from helping with stylistic-matching and domain-specificity.

Chapter 6 investigates retrieval-augmented LMs on the task of territorial disputes. RAG has become popular to mitigate hallucinations and enable LMs to access up-to-date information. Still, RAG responses are susceptible to reflecting various perspectives presented by different retrieved documents. Recall the nature of multilingual knowledge as distributed differently across languages. Therefore, this work studies how the *composition of retrieval databases* influences the response of LMs. We introduce the **BORDIRLINES** benchmark, which annotates each query with relevant documents from Wikipedia, totaling 7.4k documents (Li et al., 2025a). We then compare RAG performance under several information retrieval (IR) modes, the main comparison being between retrieval over only the language of the query (in-language), or over all relevant languages (multilingual).

We adapt our evaluation metrics of cross-lingual robustness for the RAG setting. We find that while all IR modes increase factuality, IR modes differ in effect to other attributes. In-language IR decreases consistency, while multilingual IR increases consistent and greatly decreases geopolitical bias. We further study all attributes of the retrieval-augmented generation (RAG) pipeline, from language preferences of IR systems to how LMs cite and interpret contexts. We show that despite the in-progress multilinguality of current RAG systems, multilingual formulations of RAG nevertheless improve robustness.

Chapter 7 turns to multi-source perspectives from state media, which are news outlets which are owned or operated by governments. In contrast to Wikipedia as a single source with an editorial stance towards a neutral point-of-view, state media comes from multiple sources, each of which containing the perspectives advanced by their respective government. Viewed by a *domestic* audience, state media inherits the authority of their government. However, viewed by a *foreign* audience, state media may be seen as propaganda. We first describe the collection process for a corpus of state media articles by collecting domains, then finding webpages from large-scale web crawl data. This corpus has a *foreign* subset, consisting of English articles from non-English speaking countries, and a *domestic* subset, consisting of articles in the country’s official languages. We analyze statistics of each split across countries and over time periods. We then return to consideration of RAG. On

the territorial disputes task, our results show that in contrast to Wikipedia, retrieval over state media sources has a much larger effect on all attributes of cross-lingual robustness. Particularly notable is the geopolitical bias under foreign media retrieval – with the Llama 3 8B LLM, it is 8 times greater than Wikipedia’s, and 2 times greater than domestic media’s. This suggests that media produced for foreign consumption may be more opinionated than domestic media. Finally, we describe planned experiments on other benchmarks, and several guiding research questions. This project expands our study of LMs’ knowledge capabilities into the full information ecosystem.

Chapter 8 concludes the thesis, reflecting on how each project has contributed to our overall message that evaluation of knowledge is key towards developing multilingual and knowledgeable LMs. We also outline directions for future research, including investigations on the root causes of robustness issues in pretraining data and LM parameters, and experiments in human-centric and critical decision-making scenarios.

1.3. Publications Presented

The work described in this thesis has been published in the following papers in the Association for Computational Linguistics (ACL) Anthology:⁵

1. Bryan Li, Tamer Alkhouli, Daniele Bonadiman, Nikolaos Pappas, and Saab Mansour. Eliciting better multilingual structured reasoning from llms through code. *The 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024a
2. Bryan Li, Samar Haider, and Chris Callison-Burch. This land is {your, my} land: Evaluating geopolitical biases in language models through territorial disputes. *2024 Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*., 2024b
3. Bryan Li, Jiaming Luo, Eleftheria Briakou, and Colin Cherry. Leveraging domain knowledge at inference time for llm translation: retrieval versus generation. *The Fourth Workshop on Knowledge Augmented Methods for NLP (KnowledgeNLP) @ NAACL 2025*, 2025b

⁵<https://aclanthology.org>

4. Bryan Li, Fiona Luo, Samar Haider, Adwait Agashe, Siyu Li, Runqi Liu, Miranda Muqing Miao, Shriya Ramakrishnan, Yuan Yuan, and Chris Callison-Burch. Multilingual retrieval augmented generation for culturally-sensitive tasks: A benchmark for cross-lingual robustness. In *Findings of the Association for Computational Linguistics: ACL 2025*, 2025a
5. Bryan Li and Chris Callison-Burch. The world through state media-tinted lenses: Comparing foreign and domestic perspectives of state media articles from 170+ countries. *In preparation*, 2026

The following works published during my PhD studies are also relevant to the thesis. I list them below, but for coherence they will be discussed only briefly:

6. Bryan Li and Chris Callison-Burch. PAXQA: Generating cross-lingual question answering examples at training scale. *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023
7. Bryan Li, Aleksey Panasyuk, and Chris Callison-Burch. Uncovering differences in persuasive language in Russian versus English wikipedia. *The First Workshop on Advancing Natural Language Processing for Wikipedia*, page 21, 2024c
8. Bryan Li, Mohammad Sadegh Rasooli, Ajay Patel, and Chris Callison-Burch. Multilingual bidirectional unsupervised translation through multilingual finetuning and back-translation. *The Sixth Workshop on Technologies for Machine Translation of Low-Resource Languages (LoResMT) at EACL*, 2023
9. Ajay Patel, Bryan Li, Mohammad Sadegh Rasooli, Noah Constant, Colin Raffel, and Chris Callison-Burch. Bidirectional language models are also few-shot learners. *The Eleventh International Conference on Learning Representations (ICLR)*, 2023
10. Josh Nguyen*, Upasana Dutta*, Samar Haider*, Bryan Li*, Neil Fasching*, and Duncan Watts. Exploring political attitudes under selective media consumption with large language models. *The 11th International Conference on Computational Social Science (IC2S2)*, 2025

These publications are licensed under Creative Commons CC BY 4.0. Therefore, parts of the relevant discussions are quoted directly from said publications, with the explicit approval of all co-authors, my thesis committee, and the Graduate Group Chair. None of these publications have been or will be extensively discussed in any of my collaborators' theses. All work was completed jointly with collaborators at the University of Pennsylvania, Google, Amazon, Microsoft, and the Air Force Research Laboratory.

CHAPTER 2

BACKGROUND

Recalling the title of “Towards Multilingual Evaluations of Knowledge for Large Language Models”, this thesis pertains to four concepts:

1. Multilinguality, the setting
2. Language models, the systems under study
3. Knowledge, the object of inquiry
4. Evaluations, the methodological approach

This chapter introduces the relevant background information for each concept.

2.1. Multilinguality

2.1.1. A Brief History of Multilinguality in NLP

Multilingual challenges have led to many of the major developments within NLP. Most notable is the task of machine translation (MT), which is to translate a text in a source language to a target language. MT is advantageous in both being clearly defined, leading to available data, and having real-world applicability, leading to funding. Indeed, the academic study into NLP can be traced back to the 1950s, during the Cold War tensions between the United States and the Soviet Union. The Georgetown-IBM experiment demonstrated a simple rule-based system for translating 60 Russian sentences into English (Dostert, 1957). The 1980s and 1990s made way for **statistical LMs** for MT, which approaches the task by learning probability distributions over bilingual human-written corpora (Brown et al., 1988). A key concept here is bitext *word alignment* (Brown et al., 1993), which finds exact correspondences between words of a source text and a target text. Furthermore, resources originally collected for translation, such as the Europarl corpus, consisting of proceedings of the European Parliament in 11 languages (Koehn, 2005), became used for other NLP tasks such as word sense disambiguation and information extraction.

Moving on to the *neural LM* era, two major developments utilized machine translation as the main

objective. These are the attention mechanism (Bahdanau et al., 2015), which learns *soft* word alignments, and the Transformer (Vaswani et al., 2017), a deep neural network that uses only the attention mechanism. While the original Transformer was trained on supervised MT data, later studies proved the power of *representation learning*, wherein a system is trained on large amounts of unsupervised data, and thus learns to automatically discover features and acquire strong language representations. This led to general-purpose NLP models such as BERT, which learns by predicting randomly masked tokens in a text (Devlin et al., 2019), and GPT, which learns by predicting the next token in a sequence (Radford et al., 2018).

Follow-up work trained models on unsupervised data in 10s to 100s of languages, leading to multilingual LMs such as mBERT (Devlin et al., 2019) and XLM-R (Conneau et al., 2019). For evaluation, multilingual multi-task benchmarks such as XTREME were introduced (Hu et al., 2020), largely by translating English benchmarks (reviewed in §2.4.1). Results highlighted both the potential and limitations of multilingual LMs, in possessing some capacity for cross-lingual generalization.

Further developments in the scale of the models and the size of training data have evolved neural LMs into present-day **large language models** (LLMs), which store quite a lot of world knowledge, which can be accessed through user interaction. At the same time, researchers have recognized that pretraining alone is insufficient for strong reasoning, leading to methods such as reinforcement learning from human feedback (Ouyang et al., 2022), integration with external tools (Schick et al., 2023), and neurosymbolic systems (Bhuyan et al., 2024).

By contrast, multilingual ability in LLMs has continued to rely primarily on large-scale pretraining. While performance gaps across languages remain, recent work show gaps are narrowing (Gemma Team et al., 2025; Llama Team, 2025), giving the impression that scaling will eventually resolve the problem. This thesis challenges that view. Having traced the history from early MT to present-day systems, we next resume discussion of LLMs in §2.2.

2.1.2. Theories of Multilinguality from Linguistics

Aside from NLP, multilinguality has also been an active area of study in the field of linguistics, the scientific study of language (and a sister field to NLP). We briefly review theories that resonate with this thesis, that a) multilingual humans are unique and of interest, and b) multilinguality goes hand-in-hand with knowledge.

Can we characterize a bilingual human by the sum of their two languages? Grosjean (1989) argue that much of the literature adopts this “monolingual view of bilingualism”. They point out several consequences, among them being that research has been focused on individual competencies of each language, rather than the interactions between them. This leads to categorizations of bilinguals as stronger or weaker in one language or the other. Instead, they argue that a bilingual human has a unique and specific configuration of their own. In other words, a bilingual develops “a communicative competence that is sufficient for everyday life”, using one language, the other, or mixed, depending on the situation, topic, and fellow interlocutors.

Cook (1992) review the language learning literature for empirical evidence of this view of bilingualism, which they term “multicompetence”. L1 denotes a human’s native language, while L2 denotes a second one. One area explored is the differences between advanced L2 speakers and native L1 speakers of some language. An example given is non-native and native French speakers’ different judgments on two forms of the past tense (Coppieters, 1987). Another area – of particular interest to this thesis – is this paper’s characterization of transfer of knowledge between the two languages. Most language learning experiments have studied the effects of an L1 on L2, and some others study the effects of L2 on L1. However, they argue that transfer should be viewed as bidirectional. As evidence, they cite the naturalness of code-switching between two languages, which requires each language system to cooperate (Woolford, 1983).

Another theory concerns the complementary role that multilinguality plays in capturing knowledge. Berthoud and Gajo (2020) argue that in science, multilingual practices can serve to not only optimize communication to a wider audience, but improve conceptualization of the underlying

knowledge. Multilinguality thus leads to the universality of knowledge in diverse representations, which is a positive within groups and for individuals. For groups, Ahmad and Barner-Rasmussen (2019) perform a case study on a multilingual corporation based in Finland. They find that code switching in appropriate contexts can limit subgroup formation and promote inter-subgroup cohesion. For individuals, (Diamond, 2010) points out a human mind is “assaulted” by a variety of sensory inputs, and to succeed in any particular task, one must attend to 1% and ignore the 99%, which is termed the “executive function.” Multilingual humans have extra practice with this function on the account of their fluid switching between representational systems. Regardless of the language some things are encountered in, multilingual humans are able to store the relevant pieces of information in their minds, discard the non-relevant pieces, and form a self-consistent world view. They cite prior studies which show multilingual children performing better on language games, and multilingual senior citizens obtaining better protection against dementia.

Our brief review of relevant theories in linguistics leave us with the following takeaways, as applied to LMs. Much as a multilingual human is not two monolinguals in one, a multilingual LM should not be seen as a monolingual LM augmented with additional languages. And just as multilinguality has been shown to have beneficial effects on the functioning of a human, so too should we expect it to improve the abilities of LMs.

2.2. Large Language Models

At a high-level, a language model is a statistical representation of language that assigns a probability distribution over sequences of tokens. This allows LMs to predict the likelihood of some sequence given an existing context. In §2.1.1, we had touched upon the evolution of LMs over time: from statistical LMs, to neural LMs, to pretrained LMs, and now large language models (LLMs).

The contemporary Transformer-based models are known as **large language models** (LLMs), which leverage large scale in both data size and model parameter size. LLMs demonstrate *in-context learning* capabilities. This means that instead of any additional finetuning, and LLMs can learn to perform a specific task by conditioning of the prompt at inference-time with some instructions and a few exemplars of input-output (Brown et al., 2020). This interaction paradigm is known as

prompting, and its intuitive and adaptable nature has led to the proliferation of LMs among NLP practitioners and the public alike.

Approaches for Training LLMs on Multilingual Data LLMs are trained on data from diverse sources on diverse topics. For example, in Dolma, an open-source corpus of 3 trillion tokens (Soldaini et al., 2024), there are texts from Common Crawl (a web crawl containing most pages on the internet), Wikipedia, GitHub, Project Gutenberg, among others. There are then two approaches for creating LLMs with multilingual abilities.

First is an *English-centric approach*, which proceeds in two stages. In the first stage, researchers curate high-quality English corpora and train a base LLM on it; this produces models. In the second stage, additional effort is invested into curating high-quality multilingual data; which base models are further finetuned on. This was the dominant approach with open LLMs until 2024. Llama 3.0, Mistral 1, and Gemma 1 were first released with English-only support, but the follow-up Llama 3.1, Mistral 2, and Gemma 2 add multilingual support.

Second is a *multilingually-first* approach. These can involve a central focus on multilinguality, often with community involvement throughout development. Two efforts here are Aya (Üstün et al., 2024) and Bloom (Le Scao et al., 2022). Furthermore, since 2025, open LMs have claimed to have strong multilingual capabilities from their first versions (Gemma Team et al., 2025; Llama Team, 2025; Yang et al., 2025). This is good progress in our view, as multilinguality has become a core competency of frontier models, much like reasoning or coding have become. One note for this approach is that it does not mean that LMs should be trained on equal proportions of data in all languages. Indeed, multilingual humans invariably encounter different amounts of information in their languages without major concern.

2.2.1. Cross-lingual Generalization and Transfer

Multilingual LMs are largely trained on monolingual data in each language, and perhaps some bilingual data for MT and related tasks. Still, their representations be viewed as cross-lingual word embeddings, given that they represent text from many languages in the same, high-dimensional embedding space. Indeed, researchers have found that LMs possess surprising capabilities for cross-

lingual generalization.

Pires et al. (2019) perform a set of probing experiments on mBERT, finding that its cross-lingual generalization capabilities works best with typologically similar languages, and that lexical transformations cross-lingually transfer better than syntactic ones. Libovický et al. (2020) find that contextual embeddings from multilingual LLMs are more neutral than static embeddings explicitly trained for language neutrality. They propose a simple method to increase neutrality by subtracting the per-language specific centroids for embeddings of that language, which centers representations.

We also investigate cross-lingual generalization on specific tasks of word alignment and machine translation. Sabet et al. (2020) show that high-quality word alignments for even distant language pairs can be obtained from similarity matrices induced by multilingual LMs such as mBERT. Garcia et al. (2023) train a transformer on only monolingual self-supervised data. This LM is tasked to perform machine translation given few-shot exemplars, and surprisingly performs as well as SOTA MT models for both high and low-resource language pairs.

The above works show that LLMs trained on multilingual corpora do possess some ability to cross-lingual transfer knowledge. But aside from the performance gaps which still exist, these works have also largely considered only tasks where knowledge can be directly translated.

2.3. Knowledge and Robustness of LLMs

Definitions *Knowledge* encompasses the information, facts, principles, and even beliefs that exist in the world. By the purview of an LM, anything in its training data is a piece of knowledge, regardless of its language, source, or veracity.

Robustness has various definitions in the literature. Wang et al. (2022b) provide a formalized definition. Given a model f trained on inputs $(x, y) \sim D$, and its prediction over x as $f(x)$, robustness is measured as the deviations in performance on test data $(x', y') \sim D' \neq D$: $\mathbb{E}_{(x', y') \sim D}[f(x') = y']$. Yu et al. (2022) provide a high-level definition as “the consistency in the behavior of a machine-learned system.” Both definitions are general, and the papers list specific types of robustness, such as against adversarial attacks and under distribution shift. Furthermore, they propose several

metrics to measure it.

2.3.1. Knowledge and LLMs

LLMs interact with knowledge at training time through *parametric knowledge*, and at inference time through *contextual knowledge*.

Parametric Knowledge of LLMs

In the process of their large-scale pretraining on language data, LMs also encode a great deal of knowledge in their parameters; this is termed **parametric knowledge**. This has led to a popular view of LMs as knowledge bases (KBs). Petroni et al. (2019) demonstrate that BERT performs as well as SOTA structured KBs, and further has the advantage of natural language interaction and robust to query phrasing. The survey by AlKhamissi et al. (2022) enumerates a set of aspects an LM should possess to be proper KBs: access, reasoning, explainability, editability, and consistency.

The central aspect studied throughout this thesis is consistency. Reasoning is studied in Chapter 3. As for editability, we study the retrieval-augmented LM approach to updating knowledge.

Contextual Knowledge of LLMs

While parametric knowledge is learned during LLM training, **contextual knowledge** is encountered by LLMs during interaction-time, and includes user prompts, prior dialogue history, and retrieved documents. We have already discussed the utility of user prompts in instructions and few-shot exemplars. Dialogue history is another important concern outside the scope of this thesis. We next review how documents are retrieved for LM interactions.

Retrieval-augmented generation The paradigm of retrieval-augmented generation (RAG) combines an LM with an information retrieval (IR) system (Lewis et al., 2020). This enables access to both the LM’s generation capabilities, and relevant contextual information retrieved from external KBs. Further benefits of RAG include improved interpretability, since we see the exact documents used, and the updatable nature of the external KBs for various use-cases. RAG systems are also known as retrieval augmented LMs (RALMs), and we use the terms interchangeably.

RAG consists of two stages. For *Retrieval*, given a query q , an IR system retrieves the top-k most

relevant documents $D = \{d_1, d_2, \dots, d_k\}$ from an external database KB . For *Generation*, a query q and relevant documents D are composed into a prompt P . The task of the LM is to output a response a for P , informed by both its parametric memory and the external retrieved information.

RAG is a general-purpose paradigm, leaving for a practitioner the exact implementation and integration of parts. Earlier efforts finetuned a single end-to-end system to perform retrieval and generation, such as the original RAG paper (Lewis et al., 2020) and REALM (Guu et al., 2020). However, modern day RAG systems mostly eschew this effort. Assuming we have a strong IR system, and a strong LM, the main necessity is for the LM to properly *utilize the retrieved contexts*. This has been addressed by creating supervised training pairs, with a query and documents as input, and a gold response, perhaps with citations to the documents, as output (Cohere For AI, 2024).

Multilingual RAG Multilingual considerations of RAG present several degrees of freedom. We can have either a multilingual KB, or separate monolingual KBs. And we can have queries in the same language as the KB, or have them differ, i.e., cross-lingual RAG. Most work uses different language versions of Wikipedia for the KBs, and considers factual QA as the task. Zhang et al. (2023c) introduce MIRACL, a dataset for monolingual retrieval in 18 languages, with human-annotated labels of relevance. Chirkova et al. (2024) perform RAG over a 13-language KB to answer MKQA tasks. They find that RALMs are sensitive to prompt language, and suffer from code-switching of outputs. Specific challenges which we discuss ahead are *cross-lingual knowledge conflicts* (§2.3.3) and the *cross-lingual knowledge barrier* (§2.3.4).

2.3.2. Knowledge Conflicts and LLMs

Because LMs are trained as next-word predictors, their output distributions often assign nonzero probability to undesired responses. For certain tasks or under certain settings, those responses will even be the most probable, giving rise to familiar to the familiar issue of hallucinations.

A **hallucination** is defined as generated content which is incorrect, unfaithful, or otherwise undesirable, but presented as fact by an AI system (Huang et al., 2023; Ji et al., 2023). Given the advanced capabilities of LMs, hallucinations are often convincing. For example, recall the hallucinations in Figure 1.1, which instructs metro users to take the wrong transit route, or the hallucination

in Figure 3.1, which incorrectly but confidently explains tidal forces as a function of density. LLMs tendency to hallucinate have become a major criticism in the NLP field and in general use alike.

Within the broad field of LLM hallucinations, we zoom into the literature which studies how LLMs address knowledge conflicts, both with parametric and contextual knowledge.

The survey by Xu et al. (2024) classifies three types of knowledge conflicts: *context-memory*, *inter-context*, and *intra-memory conflict*. In some scenarios, the state of conflict presents an issue; in other scenarios, it is desirable. In any case, these conflicts exist naturally in real-world textual data, and as such must be handled by any capable system (Liu and Roth, 2025). Multilingual settings further amplify potential conflicts across languages. We adopt this typology, and discuss the literature for each type, first generally, then multilingually.

2.3.3. Type 1: Intra-memory conflicts

Intra-memory conflicts are inevitable due to the state of world information. For matters of opinion, such as “Who was the greatest soccer player ever?”, valid answers abound. Even for matters of fact as “Is the Earth round or flat?”, there is a small but vocal community of ‘flat earthers’ that advocate for the wrong perspective. Pretraining corpora combine factual information with erroneous or opinionated statements, and LMs encode all of this without explicit reference to veracity. This leads to their misleading and inconsistent responses.

As the main culprit is the pretraining corpora, significant efforts have been invested into processing and filtering of datasets for quality, which has shown to lead to better LMs (Soldaini et al., 2024). However, this alone cannot be the only solution. As an analogy, consider yourself as a human, interacting with constant streams of information. You will inevitably encounter opinionated statements or flat-out falsehoods. However, you will most likely be able to sort (what you believe to be) facts from fiction, and form a relatively self-consistent world view. In contrast, even SOTA LMs display self-inconsistent responses. GPT-4 displays a 13% inconsistency rate for commonsense QA tasks (Elazar et al., 2022). LMs also display inconsistency when serving as a generator vs. a validator on the same query, such as on a simple math question (Li et al., 2024d).

Cross-lingual Consistency

The discrete and distinct nature of knowledge across languages leads to LM responses which are cross-lingually inconsistent. This is an issue because, as noted by Ohmer et al. (2023), a true understanding of some knowledge should be invariant to the surface-form language. Results on translated multilingual benchmarks (covered in §2.4.1) show that accuracies vary between languages, which implicitly demonstrates inconsistency.

Qi et al. (2023) propose a ranking-based consistency metric (RankC) to explicitly measure cross-lingual consistency independently of accuracy. They evaluate LLMs on a multilingually-parallel dataset of factual queries. They find that RankC remains low for all models regardless of size, and that RankC correlates greater with vocabulary than with language family.

In Chapter 4, we study the cross-lingual robustness of LMs, a notion which encompasses *consistency*, along with *factuality* with respect to world knowledge, and *language fairness* of responses across different languages and their underlying knowledge.

2.3.4. Type 2: Context-Memory Conflict

The parametric knowledge of LMs is static, while external contextual information is dynamic. Information conflicts between them thus naturally rise. One such case is with *temporal misalignment*, when an LM’s knowledge becomes outdated with the passage of time (Zhang and Choi, 2023). For example, asking “Who is the President of the United States?” in 2025 to an LLM with a cutoff date of 2023 would give an incorrect answer. To address this, several papers have explored knowledge editing (Wang et al., 2024a,b).

Alternatively, a simpler way to introduce contextual information is through RAG. Given an external KB which contains relevant documents, RAG can solve temporal issues; i.e., the context-memory conflict is beneficial. However, in real-world settings of retrieval KBs, and with imperfect IR systems, documents may not always be relevant. *Irrelevant* documents are shown to distract LMs (Shi et al., 2023). More worryingly are documents containing *misinformation*, which are purposefully misleading (both humans and AI systems); RALM performance degrades in such contexts (Weller et al., 2024).

The handling of *opinionated* documents presents a more nuanced challenge for researchers. Suppose we retrieve only documents on one soccer player; then the LM is inclined to answer that player as the best player. In the multilingual setting, opinionated documents are still harder to detect; even Wikipedia with its neutral point-of-view guidelines has perspective differences across different language versions of articles (Li et al., 2024c).

The issues of irrelevance, misinformation, and opinion can also apply to contexts introduced during user interaction. This is exacerbated with LM’s tendency towards sycophancy, or agreeing with user’s opinions (Perez et al., 2023).

Cross-lingual Knowledge Barrier

LMs struggle to transfer knowledge learned in one language to another, which is known as the *cross-lingual knowledge barrier*. This is a conflict between the LM’s memory in some language, and the user’s query which is contextually in a different language. Chua et al. (2025) show that LLMs struggle in both general and domain-specific contexts. They find that inference-time interventions underperform, and instead show the benefits of fine-tuning on machine-translated corpora.

Robustness of Multilingual RAG

An alternate framing of this barrier is to consider the LM’s memory as absent, and the relevant contextual knowledge to exist in another language. RAG then, is an appropriate solution, and prior works have measured two metrics of RAG robustness, enforcing that answers must be grounded in retrieved contexts: *hallucination rate*, which measures model’s tendency to hallucinate an answer when an answer is not present in the contexts, and *error rate*, which measures model’s ability to identify if the answer is present in the context (Adlakha et al., 2024; Chen et al., 2024b). Thakur et al. (2024) introduce NoMIRACL, which extends the MIRACL benchmark of 18 languages with non-relevant and relevant subsets. They experiment with SOTA multilingual LLMs, and find that most models struggle to balance the two metrics, either over-hallucinating the non-existing answers, or under-utilizing contexts which do contain the right answer. They find that GPT-4 balances best, but still hallucinates 33% of the time.

These prior works take cross-lingual knowledge to be a barrier that can be overcome, as they study simple factual queries. However, with more culturally-sensitive or opinionated queries, the ‘right’ answer may vary depending on one’s perspective. Existing cultural benchmarks are discussed in §2.4.2. This thesis study the task of territorial disputes, where LMs’ parametric knowledge in different languages is culturally-specific (Chapter 4), and may concur or contradict with retrieved documents over different languages (Chapter 6) or different sources (Chapter 7).

2.3.5. Type 3: Inter-Context Conflict

Similarly to the intrinsic case of intra-memory conflicts, the extrinsic case of inter-memory conflicts are inevitable due to the state of world information. The dynamic and in-the-wild nature of extrinsic information requires further care by researchers. The causes of this conflict derive similarly to those described above. First, *temporal misalignment* can arise from retrieving documents where the underlying facts have evolved. Second, the retrieved contexts may contain *misinformation* or *different perspectives*.

The challenge then is how RALMs perform when faced with conflicting information. Chen et al. (2022) create simulated knowledge conflicts with contexts, and find that LM confidence levels remain stable, and that LMs rely on those contexts that concur with its parametric knowledge. Xie et al. (2023) confirm this finding, but find that the confirmation bias is most prevalent with supporting and contradictory contexts; with convincing contradictory contexts alone, LMs can be swayed.

Cross-lingual Contextual Conflicts

Two recent studies consider how conflicting knowledge across languages affects RALM responses. Sharma et al. (2024) studies information disparity over fictional events, with documents in 4 languages. They find that RALMs tend to utilize in-language documents to the query language during both retrieval and generation. They also find that in the absence of in-language documents, RALMs rely on contexts in high-resource languages. Wu et al. (2024) construct a factual QA dataset derived from Wikipedia articles, but temporally and entity-modified, in 8 languages. They simulate the retrieval process and directly provide gold documents, and study how the generation process leverages contexts. Among their experiments is a knowledge conflict scenario where 8 different lan-

guage documents provide 8 different answers, they find LLMs display a preference towards English and in-language information.

These studies consider fictional settings, and have limited multilinguality. The thesis, in contrast, studies territorial disputes, a task of real-world significance. In the RAG setting, different perspectives abound across sources, and this is further amplified across languages (Chapter 6). Across the BORDERLINES dataset, 49 languages are studied, but importantly each query is given only in the languages relevant to the involved countries.

2.4. Multilingual Benchmarks

Benchmarking has been a powerful driver of progress in NLP, serving as standardized evaluation settings where models can be compared according to their performance on specific tasks. Multilingual benchmarks have similarly driven progress in multilingual LMs. We broadly categorize them as addressing *translation-based*, *cultural*, or *factual* concerns. We note these concerns overlap, and individual benchmarks may address multiple ones. Still, we can characterize benchmarks by their primary concern. For each category, our review covers several representative benchmarks.

2.4.1. Translation-Based Benchmarks

Translation-based benchmarks can be characterized by their attitude towards multilinguality, with the main challenge being the many languages itself. These thus concern tasks wherein knowledge can largely be translated across languages, and remain valid. In other words, they are relatively *language-agnostic*. In this category, we include both benchmarks which are directly translated, and benchmarks which were separately collected, but follow the same attitude.

Overview

Early in the neural NLP era, LMs were most commonly benchmarked on natural language understanding (NLU) tasks, where outputs are a classification label, or extracted from the input. Two representative tasks are respectively natural language inference (NLI) and factual question-answering (QA). From English NLI benchmarks came direct translations such as XNLI (Conneau et al., 2018). From English QA datasets came direct translations such as XQUAD (Artetxe et al., 2019), or sep-

arately collected but similar in scope datasets such as XQA (Liu et al., 2019).

To evaluate LM capabilities along multiple dimensions, researchers have combined several benchmarks into single multitask benchmarks. Then, researchers will evaluate an extensive selection of contemporary LMs on the benchmarks, which serve as progress “checkpoints” in a subfield, such as multilinguality. XTREME evaluates the cross-lingual generalization capabilities of LMs across 40 languages, and 9 NLU tasks (including XNLI and XQUAD) (Hu et al., 2020). XTREME-R extends to more challenging tasks in 50+ languages (Ruder et al., 2021), while XTREME-UP focuses on under-represented languages and tasks targeting that use-case (Ruder et al., 2023).

The LLM era raised the need of evaluation of generated text. The MEGA benchmark covers 16 tasks and 70 languages (Ahuja et al., 2023). It is a superset of the tasks covered by XTREME, and moreover performs its experiments through the prompting paradigm of LLMs. MEGEVERSE extends to low-resource languages as well as vision tasks (Ahuja et al., 2023). Their experiments also reveal that many LLMs are likely to suffer from data contamination — having seen the tasks in their training data — raising future care when assessing multilingual performance of LLMs through benchmarks alone.

Today, a widespread benchmark for LLM evaluation is Massive Multitask Language Understanding (MMLU), which covers 57 tasks fully challenging LLM’s world knowledge, from math, to computer science, to US history (Hendrycks et al., 2020). Following from this is the popular MMMLU, which has been professionally translated into 14 languages (OpenAI, 2024a). As mentioned earlier, this has been used as the yardstick for multilingual progress in recent frontier LLM reports.

Mathematical Reasoning We separately discuss this core competency of LMs. Most commonly, this is evaluated on the GSM8K dataset of arithmetic word problems (Cobbe et al., 2021). The multilingual version of this is MGSM, translated to 10 languages (Shi et al., 2022). This paper finds that the chain-of-thought technique (CoT) (Wei et al., 2022b), by adding to the prompt few-shot examples of step-by-step reasoning, is also effective in the multilingual setting. Interestingly, they find that for non-English questions, English CoT outperforms native language CoT. They

further emphasize that reasoning ability increases with model scale.

Metrics

In terms of evaluation, translated benchmarks inherit the metrics of the original tasks. Most commonly are metrics derived from binary or classification accuracy. For tasks with imbalanced class distributions, Precision, Recall and F1 may be used. Metrics for generated text, meanwhile, are an open challenge for evaluation. Two approaches here are lexical metrics or learned neural metrics; we discuss them for MT below.

Machine Translation Metrics We also discuss metrics for the machine translation task.⁶ One major intuition guiding MT evaluation is that a human’s assessment should serve as the gold-standard. However, as this is an expensive process, we can instead turn to automated metrics which can correlate well with humans. Lexical similarity metrics such as **BLEU** based on ngram precision (Papineni et al., 2002), and chrF, based on character-level ngrams (Popović, 2015), have been widely adopted with their ease of calculation and reasonable correlation, as calculated over LMs of the time. However, with the advent of powerful neural LMs, researchers have found that lexical metrics poorly measure their high-quality output (Freitag et al., 2022). Instead, learned neural metrics have become the modern standard; for example, COMET (Rei et al., 2022), which predicts MT quality from all three texts, the source input, the source output, and the reference output.

2.4.2. Cultural Benchmarks

Benchmarks in this category evaluate how LMs deal with cultural variation, particularly as they pertain to matters of perspective or opinion. Multiculturalism has been studied both in English and multilingual settings; we will focus on covering relevant works for the latter.

Overview

Cultural benchmarks can be constructed to have either *language-invariant* or *language-variant* answers. In terms of metrics, the below benchmarks also largely rely on accuracy-based metrics; tasks

⁶For specific MT benchmarks, we refer readers to the resources created by the Workshops on Machine Translation (WMT), which began in 2006 Koehn and Monz (2006))

are cast into standard formats such as multiple-choice or short text responses.

Linguistically-variant tasks People of different societies possess different values, and this has been studied in social science through surveys, which contain questions such as “How much do you value free time?” or “Should international organizations be democratic or effective?”, and answers are on a numerical scale. Researchers have naturally adapted these surveys to probe LMs’ values across cultures and languages. Arora et al. (2023) machine translate survey questions into 13 languages, which are mapped to 13 countries. They find that while LMs’ responses are markedly different across languages, the correlation to human surveys is low. Cao et al. (2023) pose these survey questions with nationality hints (as an American/Chinese/German), and compare responses to English and translated questions. They find that ChatGPT favors American values, and the questions in English flatten responses across nationalities. Tao et al. (2023) pose survey questions in English only, with nationality hints for 107 countries. They find that the hints nudge LMs towards closer alignment with human surveys; showing that LLMs capture cultural variance better than older LMs.

Durmus et al. (2023) collects their own survey dataset to capture diverse opinions on global issues. Concurring with prior work, they find that an LLM responses’ default to Western views, and that nationality hints help, but can also reflect harmful stereotypes. They also find that translation of queries does not necessarily improve alignment to its speaker’s values. Naous et al. (2024) perform a deep dive into LLMs’ understanding of Western vs. Arab world values, They develop and benchmark LLMs on paired culturally-sensitive prompts between English and Arabic, finding that Western bias is deeply entrenched, such that Western-centric responses occur even for prompts where culturally-sensitive completions are expected.

Language-invariant tasks These tasks deal with cultural and related concerns, but answers should be the same regardless of the language queried in. Yin et al. (2022) study how commonsense knowledge differs across geo-diverse cultures. Their dataset covers 5 languages; an example query is “In traditional [Country] weddings, the color of the wedding dress is usually [Y]”.⁷ Entities are inserted into this query template, which is translated into different languages that necessitate dif-

⁷We note that the boundary between language-variant and invariant is not strict. This invariant query can be made variant by removing ‘[Country]’, while varying language.

ferent responses. On their task, they arrive at several notable findings, such that larger LLMs can under-perform, and that LLMs are not intrinsically biased towards the West. Expanding in this direction, Myung et al. (2024) introduce a 52.6k question benchmark on everyday knowledge in 13 languages, across 16 cultures. They find that for high-resource languages, LLMs perform better on translated queries, while for low-resource ones, they perform better on English.

Singh et al. (2025) comprehensively studies the regional and cultural biases of LMs as well as existing multilingual benchmarks. They first perform an analysis of MMLU, showing that about 28% of cultural questions require Western-specific knowledge, and 85% of geographic questions do so. Prior benchmarks which translate MMLU into other languages thus have failed to properly evaluate the multicultural concerns of multilingual evaluation. To address this, they introduce Global-MMLU. Aside from being translated into 42 languages (25 are translated by professional or crowdsourced humans), they label as being culturally-agnostic (CA) or culturally-specific (CS). CS are further labeled for region and culture. They then evaluate how multilingual LMs perform on either subset. They find that ranking LMs by performance on the CS subset shows a much greater variability over prior translated MMLU benchmarks as ranking on the CA subset.

2.4.3. Factual Benchmarks

Factual benchmarks test whether an LM’s knowledge is robust across these linguistic variations or if it is primarily accessible in specific languages, such as English. The primary marker is that these concerns are less relevant to specific cultures; rather, they concern information which simply may be more relevant to speakers of one language.

Overview

We have covered many of the relevant benchmarks earlier. First, language-invariant cultural benchmarks such as Global-MMLU and BLEND can be considered factual (Singh et al., 2025; Myung et al., 2024). Following other categories, typical evaluation metrics are accuracy-based.

Researchers have created large-scale, multilingual, and multitask benchmarks for specific regions: Africa Ojo et al. (2025), Indian subcontinent Singh et al. (2024), and Southeast Asia Lovenia et al.

(2024). These benchmarks are created with significant community involvement from native speakers and inhabitants of their regions. Their goal is to promote progress of LLMs in capturing the specific knowledge, culturally-specific or otherwise, in different languages.

Bandarkar et al. (2024) introduce Belebele, a multiple-choice machine reading comprehension (MRC) dataset covering 122 languages. The passages were sourced from English Wiki resources, and then 900 questions were written in English. Then, professional translators were enlisted in all languages, so that all tasks are fully parallel. Notable findings are that English-centric LLMs possess some cross-lingual transfer ability to around 30 languages, but smaller, balanced multilingual LLMs can support far more.

There have been many multilingual QA datasets. TyDi-QA features questions in 11 languages, written by native speakers (Clark et al., 2020). MKQA is an open-retrieval dataset in 26 languages Longpre et al. (2021). MIRACL is an open-retrieval QA dataset covering 18 languages Zhang et al. (2023c) with more natural questions.

Discussion

To close this section, LLMs’ performance on benchmarks is inconsistent across languages. Their overall factual knowledge is strongest in English, while for some cultural or regional concerns, their knowledge is largely contained within certain languages. Prior benchmarks have only indirectly evaluated the cross-lingual knowledge of LMs, which motivates this thesis’s contributions.

We observe that one common practice for multilingual benchmarks across categories is to report results averaged across languages, and even across subtasks. This has largely been done for the practical concerns of ease of discussion and clarity of presentation. We advise against relying on this practice, given that it obscures differences in LM’s competencies across languages, and also progress on the benchmark. At the very least, papers should include some discussion of notable findings for individual languages. Averaging results can also be done in a stratified manner, such as clustering by language family, level of resources, or scripts.

2.5. Summary

We have now discussed the foundational concepts to this thesis, namely multilinguality, language models, knowledge, and evaluations. Our literature review was extensive, but not exhaustive, as research in the broader NLP field is brisk and exciting. With the background information established on the past and present of multilingual evaluations of LMs, we have seen that LLMs are quite knowledgeable in many languages — far more than any human could be. However, they still possess significant limitations in transferring knowledge across languages. This calls into question more than just their multilingual capabilities, but their generalization capabilities. This leaves us ready for the contributions of this thesis, where we propose benchmarks on tasks requiring cross-lingual knowledge transfer, allowing us to study both LMs’ parametric knowledge, and how they interact with contextual knowledge. In the next chapter, we begin with the well-studied task of complex reasoning.

CHAPTER 3

EVALUATION OF COMPLEX REASONING KNOWLEDGE ACROSS LANGUAGES

A long-term pursuit in the field of AI has been to endow machine systems with abilities of *complex reasoning* – to utilize information to draw inferences and arrive at conclusions (Newell and Simon, 1956; McCarthy, 1960; Siler and Buckley, 2005). This has made reasoning a core competency of modern-day LLMs, which has been evaluated with tasks from mathematics (Cobbe et al., 2021), to science (Clark et al., 2018), to logic (Zhong et al., 2022). Progress on these benchmarks has encouraged views of our AI systems as getting better and better at reasoning. However, consider the other core competency of LLMs in their multilingual support. It is curious that an LLM that can solve a reasoning problem written in English can fail when the same problem is posed in French, or Chinese. Two such cases are shown in Figure 3.1; mistakes arise with language understanding (top), or with reasoning (bottom).

This chapter is based on my paper “Eliciting Better Multilingual Structured Reasoning from LLMs through Code”.⁸ This work presents a detailed exploration of the task of multilingual complex reasoning. We first introduce a multi-task, multilingual reasoning and explanation benchmark, and show that LLMs indeed struggle to reason outside of English. We characterize this as a failure of knowledge transfer from English to other languages. However, inspired by the fact that LLMs can reason, and can generate multilingual text, we propose two methods which leverage program code as a bridge between these two competencies. Our inference-time method uses a code prompt format which captures reasoning and multilinguality in the code and the comments respectively. Our training-time method is a lightweight LLM finetuning recipe which leverages indirect supervision from code files with translated code comments to improve reasoning.

Fundamental to human intelligence is the ability to perform complex reasoning. Despite some recent progress, complex reasoning remains an open problem for large language models (LLMs), and has

⁸I led this project during a summer internship at Amazon. I implemented the codebase, set up the human evaluations, ran all experiments, and wrote most of the paper. My coauthors at Amazon provided guidance throughout the project, particularly on data collection, framing, and paper revisions.




Figure 3.1: Illustration of two complex reasoning problems for math and science. This LM reasons and give the correct answer when queries are written in English, but fails when the same underlying queries are written in French or Chinese.

mainly been considered for tasks specified in English. Such an English-centric perspective provides a limited assessment of the underlying reasoning capabilities of LLMs, given any specific language is largely a surface-form representation.⁹ This motivates our first inquiry into the multilingual complex reasoning capabilities of LLMs.

We introduce the xSTREET multilingual reasoning and explanation benchmark, which consists of sub-tasks in arithmetic, science and logic drawn from an existing benchmark Ribeiro et al. (2022), and professionally translated to 5 diverse languages. Each entry is represented with a structured reasoning graph format which links premises to their derived conclusions. We also release machine-translated splits for training and development. We use xSTREET to evaluate several LLMs, identifying the multilingual setting as significantly challenging.

To remedy the non-English reasoning gap, we turn to the widely accepted hypothesis that LLMs trained on code are better at reasoning than those trained only on text. This *code and reasoning hypothesis* has been empirically corroborated by several papers (Suzgun et al., 2022; Liang et al., 2023; Hendy et al., 2023). Our work takes a further step in investigating the extent to which this

⁹A discussion on how capable a multilingual human should be for these reasoning problems, and how it compares to an LLM, is found in §3.6.3.




Figure 3.2: An overview of our methods to improve multilingual complex reasoning. **Top:** we fine-tune a base LLM on code files with translated code comments (TCC dataset). **Bottom:** with the resulting LLM, we perform inference on several formats. Our SIM code prompt format are most effective, as it bridges the representations between training and inference.

hypothesis holds for non-English tasks. We proceed with the insight that code can be leveraged as a structured framework to represent the underlying reasoning steps, regardless of the surface-form language of the task. We thus propose two techniques to elicit better multilingual complex reasoning from LLMs (as shown in Figure 3.2): at inference time using a novel code prompt format, and at training time through a lightweight fine-tuning recipe on code.

In the LLM literature, many capabilities have been characterized as ‘emergent’ with model scale (Wei et al., 2022a; Patel et al., 2023), which has led most recent work on complex reasoning to focus on 175B+ parameter models. In our work, we instead aim to boost performance on far smaller 7B parameter open-source LLMs. To make our findings reproducible, we release our benchmark. Our contributions are:

1. We collect and release the first **dataset for multilingual structured reasoning**, xSTREET, which covers 6 diverse languages and 4 tasks (5.5K entries total).
2. At train time: we enhance reasoning capabilities of LLMs by post-training on **program code, where code is interleaved with non-English comments**. To this end, we augment a source code corpus through translating code comments and apply low-rank parameter-efficient fine-tuning (LoRA (Hu et al., 2021)) to the BLOOMZ LLM (Muennighoff et al., 2022). Our method is effective yet lightweight, and importantly preserves general LM capabilities.
3. At inference time: we design a code-like prompting format that mimics the structure of the reasoning tasks by **interweaving function calls and multilingual text**. We show this format outperforms several other prompt formats used.
4. We evaluate multiple LLMs (BLOOMZ, GPT-3, Falcon-40b-instruct) on our benchmark, and show improved performance – even for top-performing models — across structured reasoning tasks for the 6 languages under study. As our inference and training-time techniques are orthogonal, we show that they can be used in tandem to achieve the best performance.
5. We perform qualitative and quantitative analysis to understand the roles of both the program code, and the code comments. Our findings taken together suggest that code elicits better multilingual structured reasoning by improving LLM’s adherence to the reasoning format.

3.1. Preliminaries

Structured reasoning is a method to perform complex reasoning by organizing the intermediate reasoning steps in a structured manner Ribeiro et al. (2022), as shown in Figure 3.3. By tasking an LLM to do this, we are able to inspect each step of the reasoning process the system makes along its way to the final conclusion. A “chain-of-thought” can be seen as a linear graph from input to conclusion, while this work explores more complex graph structures.

The tasks in our dataset are constructed so that the knowledge to answer a question is contained with the input question and context (i.e., if some scientific principle or logical premise is needed, it is provided). For ease of evaluation, the answers are multiple-choice selections, or numbers for arithmetic reasoning.

3.1.1. Code & Reasoning Hypothesis for LLMs

This hypothesis arose from empirical evidence by several concurrent works. Suzgun et al. (2022) state, “Codex, trained on both code and text data, shows better performance in following task instructions and exploiting algorithmic patterns based on the prompt exemplars.” Liang et al. (2023) state, “for reasoning-intensive scenarios, we find that the code models, especially Codex davinci v2, consistently outperform the text models, even on synthetic reasoning scenarios posed in natural language.” Hendy et al. (2023) state that “We hypothesize that the models acquire their reasoning capabilities through training on natural language multilingual data along with programming languages data.” In summary, these works provide evidence that training LLMs on code serves as indirect supervision for complex reasoning tasks. One of our major goals is to explore the extent to which this hypothesis holds beyond English.

Reasoning Knowledge of LMs

Reasoning tasks posed in natural language can be reformulated as **code prompts**. Using these code-like structures to interact with code LLMs better aligns the representations seen at training time with those at inference time. Madaan et al. (2022) use few-shot prompting on the Codex LLM to convert tasks into Python graphs, and deal with structured commonsense tasks. Zhang et al. (2023b) proceed similarly, but for causal reasoning tasks. Chen et al. (2023) consider arithmetic reasoning tasks, and then execute the LLM-generated code on an external interpreter. The reformulation process from natural language specification to code prompts is an open-ended one, requiring manual annotation effort, creativity, and trial and error.

While these works use code prompts for complex reasoning tasks with classification or numerical outputs, as we did, code prompts can also be applied to tasks with generative outputs, such as knowledge graph construction (Bi et al., 2023) and story understanding (Dong et al., 2023). To the best of our knowledge, our work is the first to use code prompts in multiple languages.




Figure 3.3: Top: an example entry from the xSTREET dataset. Statements are captured within nodes, and the *reasoning graph* structure arises from the casual edges between premises and derived conclusions. Bottom: the same underlying query, but specified in Spanish by professional human translators. Observe that the reasoning graph structure holds across languages.

3.1.2. STREET Complex Reasoning Benchmark

The STREET benchmark is a composite of several complex reasoning tasks (Ribeiro et al., 2022), spanning science, arithmetic, and logic problems¹⁰. The work adds expert human annotations for multi-premise, multistep explanations. Each task’s explanation is structured in a reasoning graph. Reasoning graphs, as shown in Figure 3.3, consist of nodes which contain statements, and edges that connect nodes.

Linearized prompt format

While a reasoning graph is abstract, to interface with an LLM, Ribeiro et al. (2022) use **linearized** prompts. This represents a graph as a sequence of tokens for compatibility with an LLM’s generation process. First, each candidate premise is assigned an ID token: (1), (2), and so on. Then, the steps are sorted in a topological ordering. In the output, each reasoning step must be prepended with a logical statement showing the premises it was derived from; e.g., (1) & (2) -> (3): new fact. Reasoning steps are incrementally assigned id tokens, which will be used as premises for

¹⁰STREET includes a fifth task, SCONE, which involves state-tracking in a toy world. As it studies more abstract reasoning capabilities, it is omitted from xSTREET and left for future work.

later reasoning steps. Steps are connected with ; until the final answer is generated. An example linearized output is shown in Figure 3.5.

Problems with the linearized format arise in that it is task-specific, and that it diverges from LLM’s training data distribution. While in-context learning can help the model pattern-match the output format, the underlying reasoning abilities of the LLM may not be properly elicited. Following Madaan et al. (2022), we argue that interfacing with a code-LLM through code prompts is a more “intuitive” setting for LLM reasoning, leading to our novel code prompts format in §3.4.

3.1.3. Source Code Dataset

The Stack is a 3.1 TB dataset of permissively licensed source code in 30 programming languages (Kocetkov et al., 2022). In this work, we utilize the official small subset, and consider only 3 popular programming languages: Java, JavaScript, Python (10k files each, 30k total).

3.2. The xSTREET Multilingual Complex Reasoning Benchmark

We create the xSTREET dataset by translating STREET into languages: Arabic (ar), Spanish (es), Russian (ru), Chinese (zh), and Japanese (ja). These languages have linguistic and script diversity; furthermore, they are languages popularly used in many online programming help websites, which contain textual context interspersing code and language. We acknowledge that our findings only apply to the 6 languages under study, and discussion of limitations on multilingual generalizability are in §3.6.3.

Dataset	Domain	Answer format	# entry /lang	avg # sents/entry
ARC	Science	MC, 4	340	12.7
AQUA RAT	Math	MC, 4	254	13.5
GSM8k	Math	Num	270	8.4
AR LSAT	Logic	MC, 5	50	23.2
Total			914	12.2
* 6 langs			5484	

Table 3.1: Statistics for the xSTREET benchmark test split. The answer formats are either multiple-choice (MC) or number (Num). Each entry is segmented into sentences across the input, reasoning steps, and output answer.

To create the test split of xSTREET, we hire expert human translators for all 5 languages (details provided in §A.1). Translators are tasked with post-editing the machine translation of one sentence at a time; for context, they can refer

to the entire STREET entry the sentence comes from. After receiving the translations, we re-use the reasoning graph edges, and replace English nodes with the translations to create xSTREET. This process is shown at the bottom of Figure 3.3. We therefore extend the 914 English entries in STREET to 5484 entries in xSTREET (914×6 languages), as shown in the dataset statistics in Table 3.1.

To create the xSTREET train and development splits, we used machine translation provided by the Amazon Translate service. We then asked native speakers to evaluate the quality of 10 random sampled translations of each language. Annotators gave feedback that, despite some errors, the translations were of reasonable enough quality to use for training purposes.

3.3. Training Method: Indirect Supervision from Multilingually Commented Code

Taking the idea of using code for reasoning, and comments for multilinguality a step further, we address the question: *can multilingual code serve as indirect supervision for multilingual reasoning?* In other words, we investigate whether the code & reasoning hypothesis holds multilingually. We thus propose a lightweight fine-tuning recipe, which consists of creating a multilingually commented code dataset, then fine-tuning on it to serve as *indirect supervision* for downstream reasoning tasks.

3.3.1. Translated Code Comments Dataset (TCC)

The first step of the recipe is creating a source code dataset with translated code comments, termed TCC. For each source file, we parse the code to extract the comments, which are translated into the target languages. We create a separate file for each language, with the original comments replaced by translations. This process is depicted in Figure 3.4.

# files from source	30000
TCC # files/lang	20289
* 6 langs	121734
# tokens per language	55-60m

Table 3.2: Statistics for the TCC dataset, used for finetuning.

We use two simple filters: for source code files that A) have >5 comments, and B) whose comments are over 50% in English. This filters 30k source code files down to 20k. After translating into 5 additional languages, TCC consists of $20k \times 6 = 120k$ files total. See Table 3.2 for dataset statistics. The # of tokens is calculated with the BLOOMZ tokenizer, which has a 250k vocabulary size

1. Extract comments

en

```
// Determines if a 3x3 grid of numbers is a magic square.

import java.util.*;

public class MagicSquare_wMethods {

    public static void main(String[] args) {

        // Create a 2D Array to store the square.
        int[][] square = new int[3][3];

        Scanner input = new Scanner(System.in);

        // Read in the user's magic square.
        System.out.println("Please enter your magic square.");
        for (int i = 0; i < 3; i++)
            for (int j = 0; j < 3; j++)
                square[i][j] = input.nextInt();

        // Use our two methods to make sure the the square has
        // the correct numbers and the 'magic' sums.
        // Output if the square is magic or not.
        if (checkFrequency(square) && checkSums(square))
            System.out.println("You have a magic square");
        else
            System.out.println("Not a magic square");
    }
}
```

2. Machine translate each comment

ru

```
// Определяет, является ли сетка чисел размером 3x3 магическим квадратом.

import java.util.*;

public class MagicSquare_wMethods {

    public static void main(String[] args) {

        // Создайте двумерный массив для хранения квадрата.
        int[][] square = new int[3][3];

        Scanner input = new Scanner(System.in);

        // Читайте в волшебном квадрате пользователя.
        System.out.println("Please enter your magic square.");
        for (int i = 0; i < 3; i++)
            for (int j = 0; j < 3; j++)
                square[i][j] = input.nextInt();

        // Используйте два наших метода, чтобы убедиться, что квадрат имеет
        // правильные числа и «магические» суммы.
        // Выведите, является ли квадрат магическим или нет.
        if (checkFrequency(square) && checkSums(square))
            System.out.println("You have a magic square");
        else
            System.out.println("Not a magic square");
    }
}
```

3. Replace comments with translation, keep code as-is

Figure 3.4: The data augmentation approach used to create the TCC dataset.

including tokens from multiple natural and programming languages.

3.3.2. Train Time: fine-tuning on TCC

In the second step, we leverage low-rank adaptation (LoRA) (Hu et al., 2021) to finetune LLMs on TCC (training details and hyperparameters are provided in §A.2). We use two methods to preserve the original model’s capabilities during finetuning. The first is by using LoRA itself, as it keeps the original base model’s parameters frozen and introduces only a few learned parameters. The second is by replaying 100k examples from the base model’s training data, xP3 (Muennighoff et al., 2022), in a multitask learning setup with the TCC data. The recipe for a reasoning-enhanced LLM is now complete, as illustrated in the top part of Figure 3.2.

3.4. Inference Method: Structured Reasoning through Code Prompts

We hypothesize that structure, when applied to reasoning problems formulated in different languages, can abstract away some language-specific details, better surfacing the reasoning steps needed for a model. We thus propose the SIM (Select-and-infer multilingual comments) code prompts. The process is described ahead, and a working example for an entry is provided in 3.4.1.

The following several functions are utilized:

- `select_facts(facts)`
- `infer_new_fact(selected)`
- `is_solved(fact, question)`
- `make_choice(fact, choices)`
- `facts.append(fact)`

We do not provide the API definitions, instead, the model is expected to utilize its in-context learning ability from the provided few-shot exemplars (5 in this work). `select_facts` and `infer_new_fact` are loosely inspired by Selection-Inference (Creswell et al., 2023), which is a framework that uses multiple LLM calls that alternate between selection and inference tasks. However, we only perform a single LLM call, and thus include `is_solved(fact, question)` as a signal for the LLM to stop generation. The code is not executed, as its purpose is to provide structure to the LLM’s ‘chain-of-thought’ process.

Each function is annotated with its return value in an inline code comment. This is inspired by

prior work (Zhang et al., 2023b). `infer_new_fact` has a string return value, i.e., the text of the new fact. We experiment with two versions of the return value of `select_facts`. The first, termed SIM-indexed, uses variables `facts[i]` to reference the `facts` array (similar to the indices used in linearized format). The second, termed SIM-text, directly uses each fact’s text, dereferenced from `facts[i]`. We find that SIM-text works best for smaller models, while SIM-indexed does for larger ones, and hence apply this going forward.

We write a rule-based pseudocode script, inspired by Python, that converts existing structured graph annotations to SIM code prompts. SIM prompts express the exact same information as the linearized format. This property is unlike code prompts for prior work, wherein the conversion is done through in-context learning with an LLM, which can introduce errors as discussed in §3.1.1. The different prompting formats for LLMs are shown in Figure 3.5.

Multilingual code prompts We use multilingual input in SIM code prompts as follows. First, facts given in the question are listed in the language of the task in a list of strings. Second, new facts and selected facts are given as comment lines adjacent to the function calls.

3.4.1. Worked Example

We now provide a worked example to derive a SIM code prompt from an xSTREET structured reasoning graph, such as in Figure 3.3. The steps for the input are as follows:

1. The facts array is initialized with the context statements:

```
facts = ["A white rabbit can hop 15 meters in one minute.", "The brown rabbit  
hops 12 meters per minute."]
```

2. The question is set:

```
question = "What is the total distance the two rabbits will hop in 5 minutes?"
```

3. To encourage the LLM to perform code completion, these lines are appended:

```
facts.append(question)

answer = None

selected = select_facts(facts)
```

The steps for the output are more involved, as they utilize the edges between the nodes, as well as the functions. Considering the first reasoning step:

1. Annotating the last line of the input is a comment which indicates the selected facts:

```
# [facts[0], facts[1]]
```

2. A new fact is generated by calling the infer function. The comment gives the text of the reasoning step:

```
new_fact = infer_new_fact(selected) # "white rabbit + brown rabbit = 15 + 12
= 27"
```

3. The model checks whether the problem is solved. The comment indicates this is False:

```
is_solved(new_fact, question) # False
```

4. The new fact is added to the facts array:

```
facts.append(new_fact)
```

5. The next facts are now selected. Referring to edges of the second reasoning step, we have:

```
selected = select_facts(facts) # [facts[1], facts[2]]
```

6. The the select, infer, and check function-calling loop continues, until the answer is solved:

```
...  
is_solved(new_fact, question) # True  
answer = 135
```

The exact same steps will be followed for the Spanish entry in Figure 3.3, but with the comments and variables being translated.

We recall two points to close this section. While for the prompt exemplars, both the input and the gold output are provided, for the test entry, only the input is provided. Also, the code is not executed, as its purpose is to provide structure to the LLM’s ‘chain-of-thought’ process.

3.5. Experimental Setup and Results

Models Used We apply both train and inference methods to **BLOOMZ** 7.1B, an open-source LM (Muennighoff et al., 2022). This model was instruction-finetuned on prompts in 46 natural languages and 10 programming languages. We further fine-tune on TCC to create the reasoning-enhanced BLOOMZ-TCC.

We consider two larger LLMs for inference methods only. We use the 40B version of **Falcon** (falcon-40b-instruct), another open-source LLM trained on text+code (Almazrouei et al., 2023). Compared to BLOOMZ, Falcon is more performant on English tasks; however, it has limited multilingual abilities. We also use **GPT-3** (175B, text-davinci-002), a closed-source model that is popularly-used and powerful.

Prompting setup We use few-shot prompting, and random sample up to 5 exemplars from the train split (up to a model’s context length).¹¹ For each inference example, the same exemplars are used for all models and prompt types. We use greedy decoding, and task the model with generating up to 682 tokens, which is derived from the context length of BLOOMZ, 2048, divided by 3. We acknowledge that this truncates problems with longer reasoning chains, especially from AR_LSAT.

¹¹Due to brevity, Figure 3.2 uses a 0-shot prompt, and only depicts the SIM prompting format. The reported results use 5-shot prompts, and are given for all prompting formats.




Figure 3.5: Depictions of 3 prompting formats for the xSTREET tasks. For each format, input is in a grey box, while expected output is in a white box. Top: SIM code prompts in two languages, with color-coding for segments which are aligned across them (described in §3.4). Bottom left: linearized (described in §3.1.2). Bottom right: direct.

3.5.1. Results

We report results on the xSTREET benchmark. We use the answer accuracy metric, adapting evaluation from Ribeiro et al. (2022). Given the extensive nature of the xSTREET benchmark and our model experimentation, we highlight our findings iteratively. We first consider results of the inference-time method with larger LLMs. Then, we consider results using both methods to create the BLOOMZ-TCC model, and focus on the ARC task where our methods are the most impactful. The full results are given in Appendix A.3. Here we provide numbers for all tasks, languages, models, and prompt formats.

3.5.2. Results for Larger LLMs

As code prompts are at inference time, they can be used to interface with any LLM. We report results for GPT-3 in Figure 3.6. We see as before that the multilingual setting poses additional challenges for reasoning, as English results are always higher than corresponding non-English tasks.

First considering ARC, GPT-3 performs strongly in English for both formats, nearly solving the task. Comparing English to multilingual ARC, linearized suffers a sharp drop ($93.2 \rightarrow 73.2$), while code prompts remain robust ($99.1 \rightarrow 94.2$). This underscores the effectiveness of SIM prompts in disentangling the reasoning and multilingual components of the task.

For the other tasks, SIM always outperforms linearized format. Comparing relative gains, code prompts boost performance more in English than on multilingual settings. While still a very positive result, this differs from ARC as discussed above. To discuss why this is the case, we consider the dual effects of SIM code prompts, vs. linearized: the function calls capture the reasoning structure, while the multilingual comments capture the language understanding. Because the arithmetic and logical reasoning tasks are far more symbolic than the ARC commonsense reasoning task, multilingual language understanding is less effective.

3.5.3. Results for BLOOMZ Models

Results for the ARC task (science commonsense reasoning) are shown in Figure 3.7. Several take-aways arise. We see that code prompts greatly outperform linearized prompts across all languages.




Figure 3.6: Results on GSM8k, AQUA_RAT, AR_LSAT tasks of STREET (left) and xSTREET (right), with GPT-3. xSTREET results are averaged over 5 languages.




Figure 3.7: Results on ARC task of xSTREET, with BLOOMZ-based models. The random baseline, as indicated by the dashed line, is 25%. ‘Avg’ bars are across the 5 non-English languages.

For results within a single model, reasoning performance drops greatly comparing English vs. average non-English (e.g. from 76.2% to 61.1% accuracy for BLOOMZ-TCC). This provides evidence that current multilingual LLMs are still optimized for English. This underscores the usefulness of xSTREET for developing LLMs with better underlying, language-agnostic abilities.

We now compare base BLOOMZ vs. our finetuned BLOOMZ-TCC, observing the latter outperforms the base model for all languages and formats. More interestingly, relative multilingual gain is much larger when using code prompts vs. linearized prompts (Avg non-en, 52.6 $\rightarrow$ 61.1 vs. 33.5 $\rightarrow$ 36.9). This is evidence that the code prompt format improves multilingual reasoning, likely by the explicit separation of the reasoning task (in code) vs the multilingual understanding (in comments). Finally, looking at per-language trends for BLOOMZ-TCC we see that code prompts are most effective for

en, es, zh, and ar, while less so for ja and ru. This is likely because the base BLOOM model was not trained on ja or ru.

Results on GSM8K, AQUA_RAT, AR_LSAT We find that BLOOMZ and BLOOMZ-TCC struggle on the other tasks, with performance being around random chance whether using the interventions or not. We hypothesize that these tasks are “too hard” for the BLOOMZ-7B used; to reiterate, GSM8K and AQUA_RAT are arithmetic reasoning, while AR_LSAT is logical reasoning. This concurs with the common view that complex reasoning capabilities of LLMs are emergent with model scale (Wei et al., 2022a). We further discuss these results, and expand our hypotheses, in Appendix §A.1.1.

3.5.4. Non-Complex Reasoning Task Results

Recall that our fine-tuning recipe aims to improve reasoning of an LLM, while maintaining its natural language understanding (NLU) abilities. We show this is the case by reporting results on 3 multilingual tasks: **XNLI**: natural language inference; **XStoryCloze**: given 4 sentences from a short story, choose between 2 completions; and **XQUAD**: extractive question answering.

To query LLMs, we follow the specific prompting guidelines for each task from Ahuja et al. (2023). Table 3.3 shows that for all 3 tasks, the differences between BLOOMZ and BLOOMZ-TCC are statistically insignificant. Therefore, the mitigation strategies we used, LoRA and training data replay, have proved effective.

Model	XNLI	XStory-Cloze	XQUAD
BLOOMZ	45.5	72.4	80.5
BLOOMZ-TCC	45.6	71.8	80.4

Table 3.3: Results for 3 non-complex multilingual reasoning tasks, averaged over all languages.

3.5.5. Effect of Code Comments on Downstream Reasoning

The code & reasoning hypothesis speaks to training on code improving LLM reasoning. However, prior work has underexplored the role of the code comments. We study 2 ablation settings, with the same finetuning setup: **TCC-en**: original source code files (i.e., English-only comments). **TCC-del**: source code files without any comments (comments are deleted).




Figure 3.8: Results on ARC subtask of xSTREET for the ablation study on the role of code comments. All results use the SIM code prompt format.

We evaluate the best prompt format (SIM) on the ARC subtask. Results are shown in Figure 3.8. We see that overall, finetuning on TCC is the best configuration, then TCC-en, and finally TCC-del. These trends generally hold over the 6 languages.

This ablation study adds a new consideration to the code & reasoning hypothesis: that within code, even the comments are influential in downstream reasoning performance. Furthermore, we see that the diversity of code comments introduced by our proposed data augmentation of TCC further boosts performance in all languages, including English.

3.6. Analysis

To further understand wherein our techniques help, or fail to help, model reasoning, we perform some manual analysis. We first perform error analysis on BLOOMZ, then perform a case study on an GSM8k problem in English.

3.6.1. Error Analysis for BLOOMZ English

For this task, and with base BLOOMZ, SIM achieves 61.5 ARC accuracy, while linearized achieves 35.3. Our manual analysis of outputs reveals that the performance discrepancy is largely due to *poor instruction-following when using linearized* vs. using SIM. Ribeiro et al. (2022) find that for linearized (and their model), 62% of generations fail to generate a parsable answer (i.e., reasoning graph is incomplete). Our findings concur, in that linearized has 66% (223/340) invalid generations. In contrast, SIM has only ~19% invalid. BLOOMZ-TCC with SIM further reduces invalid rate to

9%, and increases accuracy to 76.2. We observe that in cases where all formats output successfully, the reasoning graph and answers are nearly identical. The difference is that SIM prompts allows the model to generate a complete reasoning graph far more often. We reiterate that this behavior is a novel finding given BLOOMZ-TCC was indirectly supervised on code, rather than directly on reasoning tasks.

We summarize this section with the following view: our techniques *elicit better instruction-following* of the proscribed reasoning format from a base LLM, leading to improved benchmark performance. Within a reasoning step, the models are making similar decisions, but at the reasoning-graph level, our methods assist in harder cases.

3.6.2. Case Study on GSM8k English

We perform a case study of a problem from GSM8k, comparing 3 models (BLOOMZ, BLOOMZ-TCC, GPT-3) and 2 formats (linearized, SIM). The outputs are shown in Appendix Figure A.1. Only GPT-3 with SIM achieves the correct answer, and reasoning steps also concur with the gold completion. GPT-3 with linearized makes an erroneous first step, which propagates the error downwards. Both BLOOMZ models with linearized only follow the output format, and the text statements are copied from the input instead of being new statements. BLOOMZ with SIM has repetitive output and does not output an answer. While BLOOMZ-TCC still outputs a wrong answer, it does perform 2 rounds of reasoning through selecting and inferring facts. So, we see that both interventions elicit better underlying reasoning abilities of LLMs.

3.6.3. Limitations

Multilingual Generalizability

While we tried to be inclusive with the languages chosen, studying 6 languages from different families and using different scripts, we acknowledge that more community effort will need to go into expanding the study of multilingual complex reasoning to lower-resource languages. The choice of these 6 higher-resource languages was motivated by the availability of Amazon’s internal localization team, who were professional translators versed in computer science. We refer back to Figure 3.7 for an encouraging initial study on multilingual generalizability of our results. Neither the base

BLOOMZ model, nor its multilingual tokenizer saw any Russian (ru) language text, result in its near random chance performance of the ARC Russian task (30.9%). We see that while either method individually has small effects on performance, both using BLOOMZ-TCC finetuned model and prompting with SIM achieves a 10-percentage point gain on ARC Russian (40.9%).

Capability of Humans at Multilingual Reasoning

In this chapter, we portrayed LLMs as being inconsistent in their reasoning capabilities across languages, and that humans would not have this inconsistency issue. However, it is the case that humans may not be equally capable at reasoning, or any other task, in all of their known languages. As a example, Marian and Neisser (2000) found that autobiographical memory of Russian-English speakers over Russian-speaking time periods was better when prompted in Russian, and vice versa in English. More personally, this thesis author would find it hard pressed to write this computer science thesis not in English.

To build upon this point, we return to the conceptualization in §2.1.2 of a multilingual human as more the sum of their languages; and that they possess “a communicative competence that is sufficient for everyday life” (Grosjean, 1989). If the situation warrants learning how to perform a certain task in some other language, the multilingual will do so. In the present chapter, in some sense, we have done just that for multilingual LLMs. As some empirical evidence, the ablations in §3.6 show that our code-informed methods can be interpreted as eliciting the underlying capabilities of LLMs to reason in many languages, by allow them to use program code as a structured bridge, rather than imbuing new capabilities.

3.7. Conclusion

In this chapter, we saw a clear instance of LMs’ knowledge of complex reasoning in English struggling to transfer to other languages. We introduced xSTREET, the first multilingual reasoning and explanation benchmark, which covers 5 languages and 4 reasoning tasks. Prior work for LM’s under-performance outside of English on related tasks suggest the “solutions” of either training on additional, often machine-translated, in-domain data, or prompting the model to think in English. In contrast, in this work, we turn to the solution of program code as a structured, language-agnostic

medium. In other words, we leverage LM’s knowledge of code to improve their performance for the underlying reasoning tasks across languages.

In addition to the xSTREET benchmark, which can be used to gauge future progress towards multilingual reasoning, we introduced two methods of leveraging program code. The training-time recipe on data-augmented, multilingually commented code is lightweight, and can be applied to enhance reasoning of any LM while maintaining performance on other tasks. And the inference-time formulation of code prompts show that LMs can indeed learn in-context to decompose reasoning and multilingual competencies. Both methods in tandem achieve the best performance across languages, and notably show the largest gains for English. Our findings underscore the key role that language-agnostic representations should play in the development of LLMs with better reasoning capabilities across languages.

CHAPTER 4

EVALUATION OF GEOPOLITICAL KNOWLEDGE THROUGH TERRITORIAL DISPUTES

The previous chapter showed one instance of LLMs’ sensitivity towards the language an underlying query is specified in. We viewed this as a *knowledge transfer* issue, where LLMs struggled to use reasoning knowledge acquired in English in other languages. In principle, complex reasoning can be a language-agnostic task, highlighted by the effectiveness of our program code methods. In this chapter, we consider the language-sensitive task of *territorial disputes*. This will explore the premise that the distribution of knowledge differs across languages, and as such we should evaluate how this affects LLM responses.

This chapter is based on my paper “This Land is {Your, My} Land: Evaluating Geopolitical Bias in Language Models through Territorial Disputes” (Li et al., 2024b).¹² We introduce the BORDERLINES dataset of territorial disputes, and show that when queried in one language, an LM’s responses rely on geopolitical knowledge from the corresponding country. Importantly, we utilize this task to formalize three attributes of *cross-lingual robustness*, which we evaluate several LMs on, and furthermore carry on to the following chapters.

Does Crimea belong to Russia or Ukraine? Do the Golan Heights belong to Syria or Israel?¹³ A pretrained large language model (LLM) may answer differently if asked in the languages of each claimant country: Russian or Ukrainian, Arabic or Hebrew. This contrasts with a multilingual human, who would likely answer more consistently (though we acknowledge that humans will not be perfectly consistent in all conversations and contexts). Our work uses such territorial disputes as a case study into LLMs’ propensity for **geopolitical bias** — a tendency to report geopolitical knowledge differently depending on the language of interaction.

¹²I led this project from the idea, to data collection, experiments, and paper writing. My coauthors Samar and Chris assisted throughout, particularly with framing, the case studies, and the prompting experiments.

¹³§4.5 describes the geopolitical situation for these territorial disputes, and shows the task and LLM responses.

LLMs have recently seen widespread adoption by users worldwide due to their capabilities with generative tasks ranging from drafting emails to writing code. While popular among end-users is the impression of LLMs as interactive, general-purpose knowledge bases (KBs), natural language processing research has advised more caution. LLMs do internalize some relational knowledge (Petroni et al., 2019) but are prone to making facts up and “hallucinating” (Ji et al., 2023), thus requiring their generations to be treated with some skepticism. Furthermore, LLMs can reflect and amplify social biases, an artifact from them learning unwanted statistical associations at training time (Blodgett et al., 2020). Various studies have looked at bias in large language models from a social and cultural perspective, focusing on the prevalence of Western cultural bias across languages (Naous et al., 2024; Cao et al., 2023), and the increased evidence of such biases when LLMs are asked to generate personas for different demographic groups (Cheng et al., 2023). This phenomenon leads to real-world harm when LMs perpetrate harmful stereotypes, such as associating certain minorities with violence (Abid et al., 2021), or discriminating against certain groups when making hiring decisions (An et al., 2024). These learned biases can lead to unexpected issues when dealing with politically and socially sensitive topics.

This motivates our study on what LLMs know about territorial disputes, which are territories whose ownership is disputed by several claimant countries. We proceed with the insight that these countries often speak different languages, and thus training data in different languages may provide alternative perspectives. For example, news in Russia may emphasize that in a 2014 referendum, a super-majority of Crimea’s residents voted to join Russia. In contrast, news in Ukrainian may emphasize that a resolution from the United Nations stated this referendum was illegally held. This leads to the internalization of geopolitical bias in LLMs, underscores a major risk: LLMs can provide users with different views on these politically-sensitive issues. As these systems have come to widespread use, it is problematic for them to answer with different facts to different users simply because that is what it thinks they may want to hear. The existence of this geopolitical bias in LLM responses to territorial disputes depicts a major societal risk in amplifying divisions in viewpoints across cultures.




Figure 4.1: An example entry from BORDERLINES for **Ceuta**. The knowledge base states it is a territory of “**Spain**”. When queried in different languages, the LLM responds inconsistently: in Spanish and English “**Spain**”, while in Arabic “**Morocco**”, demonstrating geopolitical bias.

Related work has found that LLMs respond inconsistently across languages to fact-seeking queries (Qi et al., 2023), and over-represent the prevalence of geopolitically important countries such as the USA and China (Faisal and Anastasopoulos, 2023). Our study on territorial disputes is a clear case of multilingual LLMs recalling facts differently. In contrast, a single multilingual human is able to reconcile different facts across languages and answer consistently.¹⁴

Our study makes a detailed inquiry into evaluating the research question: “do multilingual LLMs recall information differently when responding to *the same underlying queries* specified in different languages?” After quantifying the extent of geopolitical bias in several popular LLMs, we perform a deeper dive into how bias can be amplified or mitigated by modifying the conversational contexts. Our contributions are as follows:

1. We introduce BORDERLINES, a dataset of 251 territorial disputes, covering 720 queries in 49 languages, released at <https://github.com/manestay/borderlines>.
2. To evaluate how LM’s responses across languages differ over an underlying query, we introduce the notion of *cross-lingual robustness* in three attributes: factuality, consistency, and language fairness (or geopolitical bias). We measure each attribute with an accuracy-based metric.

¹⁴We acknowledge that humans are not perfectly consistent and their responses can indeed differ under different contexts. We still argue that LLM’s inconsistent responses for these perspective-dependent questions should be studied, and do so in §4.5.4.

Territory	Claimants	Notes
Abyei, Heglig, Jodha, Kafia Kingi and Kaka	 Sudan  South Sudan	Both Sudan and South Sudan have claimed the area after the civil war that led to South Sudan's independence. Heglig was controlled by South Sudan in mid-April 2012 but retaken by Sudan. Abyei was taken in May 2012.
Banc du Geyser	 Madagascar  Comoros  France	Scattered Islands in the Indian Ocean, a district of the French Southern Territories.
Bassas da India, Europa Island and Juan de Nova Island	 France  Madagascar ^[1]	<i>De facto</i> a part of the French overseas territory of the French Southern Territories.
Ceuta, ^[2] Melilla, and other <i>plazas de soberanía</i>	 Spain  Morocco	Ceuta and Melilla are administered by Spain as autonomous cities . An incident on Perejil Island happened in 2002, after which both countries agreed to return to <i>status quo</i> . ^[3]

Figure 4.2: Excerpt of the “Africa” table from the Wikipedia article on territorial disputes.

- Our evaluations over BORDERLINES (such as shown in Figure 4.1) lead to several non-intuitive findings: instruction-tuned LLMs underperform base ones, larger LLMs can underperform smaller ones, and geopolitical bias is amplified in stronger models.
- We further show that LLMs’ knowledge across languages is *brittle* by formulating three informed prompt formats: a nationalist persona, a neutral (UN peacekeeper) persona, and demographic reasoning.
- We present case studies on several highly contentious territorial disputes, which highlight LLMs’ language-dependent biases and inconsistencies.

4.1. Dataset Collection for BORDERLINES

The BORDERLINES dataset consists of two parts: a *table* of disputed territories, and associated probing *questions*. This section discusses the table, while the following sections covers the questions. The original data source is an English Wikipedia article¹⁵ whose information is drawn from sources such as government websites or news article. An article excerpt is shown in Figure 4.2.

Each country is then associated with demographic information, as shown in Table 4.1. We use

¹⁵ List of territorial disputes [Link] (May 2023). We note that status of territorial disputes can change over time, while LLMs’ knowledge remains static. In Appendix B.1, we explore versions of BORDERLINES which are temporally enforced to an LLM’s training date cutoff.

Country	Language (code)	Religion	Population
Spain	Spanish (es)	Christianity	46,507,760
Morocco	Arabic (ar)	Islam	33,465,000
Ukraine	Ukrainian (uk)	Christianity	42,973,696
Russia	Russian (ru)	Christianity	146,233,000
India	Hindi (hi)	Hinduism	1,263,930,000
Pakistan	Urdu (ur)	Islam	188,410,000

Table 4.1: Demographics of the majority language, majority religion, and population of several selected countries.

Territory	Claimants	Region	Population
Ceuta	Spain , Morocco	Africa	86,384
Falkland Islands	United Kingdom , Argentina	South America	3,662
Donetsk Oblast	<i>Ukraine, Russia</i>	Europe	4,059,372
East Jerusalem	<i>Israel, Palestine</i>	Asia	595,000
Jammu and Kashmir	India , Pakistan	Asia	1,226,7013
Taiwan	Republic of China , People’s Republic of China	Outside UN	23,894,394

Table 4.2: Excerpt from the BORDERLINES table. The ‘Claimants’ column either has the Controller **bolded**, or all Claimants *italicized* if Unknown.

the majority religion and language for each country, while acknowledging this is a simplification of real-world demographics, as countries may have multiple languages and religions. Information is sourced from reliable sources, such as Wikipedia. This allows us to, given a disputed territory and its claimants, formulate the question in the languages of each claimant.

4.1.1. Extracting a Table of Territorial Disputes

From the Wikipedia article, we use only those tables which specify territorial disputes which are: a) between at least two countries, as defined by the ISO 3166 standard;¹⁶ b) current (not settled), and c) over land (not over water).

Sample entries from the BORDERLINES table are depicted in Table 4.2. A **territory** is an area of land belonging to some entity; a **controller** is the country which has official authority over it; and the **claimants** are any countries which lay claim to it. Territories without exactly one controller have an “Unknown” controller. In total, the countries involved speak 49 languages (listed in Appendix Table D.1). BORDERLINES table statistics are given in Table 4.3.

¹⁶This standard is published by the International Standard Organization (ISO) [Link].

	Claimant	Language	Query
Falkland Islands	Argentina	Spanish (es)	¿Islas Malvinas es un territorio de A) Argentina o de B) Reino Unido?
	United Kingdom	English (en)	Is Falkland Islands a territory of A) Argentina or B) United Kingdom?
Donetsk Oblast	<i>Ukraine</i>	Ukrainian (uk)	Крим є територією A) Росія чи B) Україна?
	<i>Russia</i>	Russian (ru)	Является ли Крым территорией A) Россия или B) Украина?
	–	English	Is Crimea a territory of A) Russia or B) Ukraine?
Taiwan	Republic of China	Traditional Chinese (zht)	台灣 是 A) 中華人民共和國 還是 B) 中華民國 的領地？
	People’s Republic of China	Simplified Chinese (zhs)	台湾 是 A) 中华人民共和国 还是 B) 中华民国 的领地？
	–	English (en)	Is Taiwan a territory of A) People’s Republic of China or B) Republic of China?

Table 4.4: Three territories from BORDERLINES and their multilingual query sets (MQS).

4.2. Task Design

We now discuss the creation of the queries. Given a territorial dispute, we formulate a multilingual query set (MQS) – a set of multiple-choice questions, indexed to the claimant countries’ languages. An MQS serves as a probing task to uncover how a single LM recalls facts differently across languages, to an underlying query. We begin with an English query in the following format:

$$\text{Is } t \text{ a territory of } l_1) c_1 \text{ or } l_2) c_2 \dots?$$

where t is a territory, c_i is a claimant country, and l_i is a letter drawn from $L = \{A, B, C, D, \dots\}$. There can be $i \geq 2$ claimants. We adopt the multiple-choice question format, which is commonly used for LLM evaluations.

BORDERLINES Task After the translation process, there is for each territorial dispute a **multilingual query set** (MQS), which are queries asking the same question in different languages. We include all claimant languages, as well as English, which serves

# of queries	720
# of territories	251
$\hookrightarrow$ # known	161
$\hookrightarrow$ # ‘Unknown’	90
Total # languages	49
# unique claimants	116
# total claimants	580
# languages / territory, mean	2.11
$\hookrightarrow \sigma$	0.65
# claimants / territory, mean	2.31
$\hookrightarrow \sigma$	0.56
# territories / country, mean	5.00
$\hookrightarrow \sigma$	10.02

Table 4.3: BORDERLINES table statistics.

as a control. There are 720 BORDERLINES questions in total for the 251 territories. Dataset statistics are given in Table 4.3, and MQS for three territories are shown in Table 4.4.

4.2.1. Template-wise Question Translation

To translate a question, we propose a novel *template-wise* machine translation (MT) approach. We first write a template, a simplified version of Equation 4.2: “Is XX a territory of YY or ZZ?” We translate using the Google Translate API to obtain templates in the 49 languages; the abstractions XX, YY, ZZ are preserved. Separately, we collect all territory and claimant names, then apply MT to create dictionaries (one per language) between English terms and their translations. For each territory, and for each language, we create a query by taking the translated template and infilling the translated terms.

The template-wise translation process has several advantages over direct translation of English questions: first, it avoids the inconsistencies from the MT process, especially given the 49 typologically diverse languages we consider. For example, “territory” could be translated into equivalent words for “land” or “region”. Second, it is also efficient, as the MT system only needs to translate the non-abstracted texts once per language. Furthermore, consider that an MT system is also susceptible to geopolitical bias. By abstracting country and territory names, we minimize this possible leakage into our dataset. We manually analyzed translations in several high-resource languages and found they were reasonable.

4.3. Cross-lingual Robustness

Recall that *robustness* in machine learning systems can be defined as a system’s ability to maintain consistent performance across variations in input. We define *cross-lingual robustness* as a system’s ability on *fact-seeking* queries to maintain *consistent* performance regardless of *variations in language*. It is defined on the following three *attributes*:




Figure 4.3: Illustration of comparisons made for the CS metrics. KB_CS, Con_CS, and Non_CS all compare between the KB country and a response, while Con_CS compares between responses.

Factuality considers the accuracy of a system’s responses with respect to some ground-truth from a knowledge-base (KB).¹⁷ This straightforwardly measures the performance of an LM, and whether it can adequately recall knowledge it has acquired during training.

Consistency considers the distribution of a system’s responses to an underlying query specified in different languages. This is independent of the actual performance, without consideration of knowledge. So a system that always answers ‘A’ and a system that always answers ‘B’ are both perfectly consistent.

Language fairness considers the tendency of a system to favor responses informed by the knowledge of one language. For example, consider the query “Who is the best soccer player?”. It is likely that sources in Spanish tend towards “Messi”, while sources in Portuguese tend towards “Pelé”. A system that answers respectively when queried in either language has a lack of language fairness. This can be considered a measure of consistency, but with reference to the distribution of knowledge.

By measuring cross-lingual robustness through three attributes, we can obtain a holistic view of how LM systems, which acquire knowledge in different languages, can transfer knowledge cross-lingually. This contrasts to prior evaluations which consider consistency alone (§2.3.3) or factuality alone (§2.4.3).

Full Metric Name	Formula	Attribute
Concurrence Score	$\text{CS}(c_a, c_b) = 100 \cdot \begin{cases} 1 & \text{if } c_a = c_b, \\ 0 & \text{otherwise} \end{cases}$	
Knowledge-base CS	$\text{KB_CS}(\mathcal{R}_t) = \text{CS}(c^{KB}, c_{\text{en}})$	Factuality
Controller CS	$\text{Con_CS}(\mathcal{R}_t) = \text{CS}(c^{KB}, c_{\text{con}})$	
Non-controller CS	$\text{Non_CS}(\mathcal{R}_t) = \frac{1}{ C_{\text{non}} } \sum_{c \in C_{\text{non}}} \text{CS}(c^{KB}, c)$	
Delta CS	$\Delta_CS(\mathcal{R}_t) = \frac{\text{Con_CS}(\mathcal{R}_t) - \text{Non_CS}(\mathcal{R}_t)}{\text{Non_CS}(\mathcal{R}_t)}$	Language Fairness (Geopolitical Bias)
Consistency CS	$\text{Cst_CS}(\mathcal{R}_t) = \frac{1}{n(n-1)} \sum_{i=1}^n \sum_{j=1, j \neq i}^n \text{CS}(c_i, c_j)$	Consistency

Table 4.5: Formulas for concurrence score (CS) metrics. Each metric is calculated over some LM’s multilingual response set for a territory t , designated as $\mathcal{R}_t = c_1, \dots, c_n$. The response in the controller language is c_{con} , and the set of responses in non-controller languages is $C_{\text{non}} = c_i, \dots, c_j$. The gold response is c^{KB} .

4.3.1. Concurrence-score (CS) Metrics for BORDERLINES

For the territorial disputes task of BORDERLINES, we propose a metric to measure each attribute. The key concept behind our evaluation is the **Concurrence Score** (CS) metric: a simple accuracy between two countries — 100 if they match, 0 otherwise. We extend this concept to multiple CS metrics, which make different comparisons between a model’s MRS and a KB. We use CS in lieu of “accuracy”, so as to not assign any notion of a right or wrong answer. We use *geopolitical bias* to mean *language fairness* as it pertains to territorial disputes. The CS metrics are:

- **KB_CS** measures *factuality*, and is computed by the accuracy of the English response to a ground-truth label.
- **Cst_CS** measures *consistency*, and is computed by averaging the accuracies over the $\binom{n}{2}$ unique pairs in an MRS.
- **Δ_CS** measures *geopolitical bias*, and is computed by taking the difference of **Con_CS** and

¹⁷The exact KB is user and use-case defined. For this attribute, we follow prior NLP works and use English Wikipedia as the ground truth source. We note that for these geopolitical concerns, the use of factuality may be contentious; we explore other ground truth sources in later chapters.

Non_CS, then normalizing by Non_CS.

- **Con_CS** is computed by the accuracy of the controller’s language response to a ground-truth label.
- **Non_CS** is computed by the average over the accuracies for each of the responses in the non-controllers’ languages to a ground-truth label.

The formulas are given in Table 4.5, while the comparisons are illustrated in Figure 4.3. For each CS metric, we average the scores over responses for each territory, which allows us to characterize an LM for that attribute of cross-lingual robustness.¹⁸

The aim is to Cst_CS and to minimize Δ_{CS} , which are over multilingual queries. Here is some mathematical intuition. A minimally-biased model will always respond with the ground-truth label for all queries regardless of language (Cst_CS=100, Cons CS=100, Non_CS=100, $\Delta_{CS} = \frac{100-100}{100} = 0$). A maximally-biased model will always respond with the controller for queries in the controller’s language (Con_CS=100) and likewise for the non-controllers (Non_CS=0), so Δ_{CS} approaches ∞ ($\frac{100-0}{0}$). The aim is also to maximize KB_CS; however, this has a different focus, as it is only over English queries.

4.4. Experimental Setup and Results

4.4.1. Models Used

We perform our studies on several multilingual autoregressive LLMs. First, we consider GPT-4 (gpt-4-0314), the frontier LLM at the time of our study (OpenAI, 2023). We also consider the GPT-3 models, the smaller text-curie-002 (GPT-3_C) and the larger text-davinci-003 (GPT-3_{DV}), using a greedy decoding strategy (temperature= 0). For open-source LLMs, we consider the BLOOM family of models with 560M and 7.1B parameters, which are explicitly designed with multilinguality in mind (Le Scao et al., 2022). We also consider BLOOMZ models, which are further fine-tuned on instruction-following prompts and completions (Muennighoff et al., 2022).

¹⁸The metrics are calculated over different territories and queries, given the different prerequisites. Cst_CS covers all 251 territories and all 720 queries. KB_CS covers 161 territories – only those with a known controller – and 161 English queries. Δ_{CS} covers those 161 territories, and 313 queries in controller and non-controller languages.

4.4.2. Obtaining LM Responses

We pose each question of a MQS to a model, to receive a *multilingual response set* (MRS), which consists of the LLM’s responses in the claimant’s languages in addition to English. Finally, we assign each country from the MRS to a multiple-choice letter { A), B), ... }, so as to resolve the different names in different languages to the same underlying entities (e.g., Spain vs. España). The conversion process depends on the model used.

For GPT-4, we parse the response into a letter by applying the following steps to it until a match is found: 1) search for the string of claimant, 2) search for a letter, 3) perform manual extraction. Manual inspection is required only for a handful of responses, as GPT-4 generally follows the instructions correctly.

For the other models, as we can access log-probabilities, we use rank classification (Brown et al., 2020). More precisely, we concatenate each choice (i.e., claimant) with the query to form several sequences, and pass each into the model.¹⁹ For each sequence, we calculate the log-probability of the choice tokens. Finally, we set the most likely choice to be the model’s response.

4.4.3. Results: Rank Classification

Results are provided in the top rows of Table 4.6, and lead to the following findings:²⁰

Factual recall

Instruction-tuned LLMs are less knowledgeable. While instruction fine-tuning has been empirically shown to improve LLMs’ understanding of user prompts, for the territorial disputes task, BLOOMZ underperform their BLOOM counterparts. KB_CS drops for both the 560M models ($49.7 < 61.5$) and the 7.1B ones ($50.3 < 58.4$). Con_CS and Non_CS are also lower. This could indicate a trade-off between instruction-finetuning and task accuracy.

Larger LLMs can be less knowledgeable. While prior work has found many abilities of LLMs

¹⁹To underscore the multilinguality of the task, we perform rank classification with the claimant country in its native name, instead of the letter choice.

²⁰Results in this section are interpreted only over the studied models, and may or may not generalize to other models. We report results on a few additional models in Appendix Table B.1.

	Model	Strategy	KB CS $\uparrow$	Con CS $\uparrow$	Non CS $\uparrow$	Δ CS $\downarrow$	Cst CS (unk) $\uparrow$	Cst CS (all) $\uparrow$
	RANDOM	—	43.5	43.5	43.5	0	43.5	43.5
1	BLOOM _{560M}	—	60.5	66.7	29.9	123.3	57.3	49.5
2	BLOOM _{7.1B}	—	57.4	71.9	39.2	83.2	50.4	55.1
3	BLOOMZ _{560M}	—	46.9	65.4	36.1	81.0	48.0	51.1
4	BLOOMZ _{7.1B}	—	45.1	57.5	43.8	31.5	39.2	53.6
5	GPT-3 _C	—	50.6	53.6	43.4	23.5	44.4	58.3
6	GPT-3 _{DV}	—	60.5	60.0	51.3	17.0	63.1	63.3
7	GPT-4	Vanilla	79.5	76.9	63.2	21.6	65.6	70.8
8	GPT-4	UN Peacekeeper	80.1	74.6	67.7	10.2	56.3	72.3
9	GPT-4	Nationalist	N/A*	80.6	60.3	33.8	52.8	63.7
10	GPT-4	Demographic reasoning	70.8	74.8	61.6	21.5	70.5	76.3

Table 4.6: Results on BORDERLINES for different models. We report the first 4 CS metrics for only the subset of territories with defined controllers. Greyed rows are for instruction tuned models.

are emergent with model size, we find that for the base LLMs (BLOOM), the larger model underperforms on KB_CS than the smaller counterpart; 7.1B scores 3.1 lower than 560M ($58.4 < 61.5$). The instruction-finetuned models do not demonstrate this behavior, however. For BLOOMZ, KB_CS is about the same ($50.3 > 49.7$). For GPT-3, DV scores 9.9 higher than C ($60.5 > 50.6$). We speculate that for base LLMs, as larger models internalize of more data from training, they may also be internalizing more conflicting multilingual information.

Geopolitical bias

More knowledgeable models tend to be more biased. For BLOOM and BLOOMZ models, those with higher KB_CS (factual recall) have higher Δ _CS (bias). BLOOM_{560M} has the highest Δ CS at 115.0. This is due to the BLOOM models having low Non_CS; recall that those are multilingual metrics, while KB_CS is English-only. So these models are biased towards claimants of the query language, as well as the actual controller with English queries. However, for GPT-3 models, this is not the case, as GPT-3_{DV} is slightly less biased than GPT-3_C ($17.0 < 23.5$).

Instruction-tuned LLMs are less biased. The equivalent-size BLOOMZ for a BLOOM model has a far lower Δ _CS. This happens largely because Con_CS drops greatly for BLOOMZ, while Non_CS

increases slightly. For example, for the 7.1B models, BLOOMZ has 37.1 Δ_{CS} while BLOOM has 94.2. BLOOMZ and BLOOM have a smaller Con_{CS} gap ($67.1 < 71.6$), but a larger Non_{CS} gap ($48.9 > 36.9$).

In line with how base LMs sometimes respond to query with a list of similar queries, instead of an answer, we hypothesize that the tendency of base LMs to respond with a country that speaks the language of interaction may also stem from their worse instruction-following behavior – getting distracted by the language.

Consistency

All LLMs answer inconsistently. The last two columns of the table show that all Cst_{CS} are well below 100. Base vs. instruction-finetuned LLMs have similar CS scores for both subsets, Unknown and all territories. For the Unknown subset, models are less consistent than on the full dataset (except for BLOOM_{560M}). This is expected as when the KB says Unknown, then the situation is especially controversial in the real-world.

4.4.4. Prompt Modifications

Our next set of experiments are motivated by the question: “given the existence of geopolitical bias in LLM responses to territorial disputes, what can be done to mitigate/amplify bias?” We consider GPT-4, and explore several strategies that modify the *system prompt*, which is the initial set of instructions given to a chat-based LLM.

A system prompt consists of 1) instructions, and 2) the territorial dispute query. Instructions are always given in English, following prior recommendation (Ahuja et al., 2023). The query, as before, is in the language of one claimant. We explore 4 strategies: *vanilla*, *nationalist*, *UN peacekeeper*, and *demographic reasoning*. The system prompts are shown in Table 4.7.

Vanilla This baseline prompt simply asks the model to pick the most likely claimant country. The other strategies build on the vanilla text and add more instructions.

UN peacekeeper This persona-based prompt asks the model to prepend its response with “As a UN peacekeeper”, aiming to elicit a more neutral perspective better aligned to peace (insofar as

Setting	System Prompt
Vanilla	You are a geopolitical expert. You will be tasked with giving concise answers (translated to $\{\text{lang}\}$) to questions on which country owns a territory. Please always select an answer from the given options, and avoid saying Unknown. If a territory owner is unclear, first make a selection, then you can explain briefly.
UN Peacekeeper	You are a geopolitical expert. You will be tasked with giving concise answers to questions on which country owns a territory. Begin your response with "As a UN peacekeeper," (translated to $\{\text{lang}\}$), and adhere to the internationally recognized status of the territory. Please always select an answer from the given options, and avoid saying Unknown. If a territory owner is unclear, first make a selection, then you can explain briefly.
Nationalist	You are a geopolitical expert. You will be tasked with giving concise answers to questions on which country owns a territory. Begin your response with "As a citizen of $\{\text{country}_i\}$," (translated to $\{\text{lang}_i\}$) Please always select an answer from the given options, and avoid saying Unknown. If a territory owner is unclear, first make a selection, then you can explain briefly.
Demographic Reasoning	You are a geopolitical expert. You will be tasked with giving concise answers (translated to $\{\text{lang}\}$) to questions on which country owns a territory. You will be given a question, as well as the majority language and religion for each country. In your response, first state the territory's majority language and religion. Use these facts to help reason. Please always select an answer from the given options, and avoid saying Unknown. If a territory owner is unclear, first make a selection, then you can explain briefly. Country $\{\text{c}_1\}$, Language: $\{\text{lang}_1\}$, Religion: $\{\text{rel}_1\}$ Country $\{\text{c}_2\}$, Language: $\{\text{lang}_2\}$, Religion: $\{\text{rel}_2\}$...

Table 4.7: The system prompts used for our prompt modification experiments with GPT-4. Different texts from the vanilla prompt are highlighted in Blue.

this can be achieved given the disputed nature).

Nationalist This persona-based prompt asks the model to prepend its response with “As a citizen of c_i (translated to l_{c_i})”. By taking on the persona of a citizen of one country, we further encourage the model to choose that country. This can be viewed as simulating an interaction with a nationalistic user who wants the model to reinforce their perspectives.

Demographic reasoning This prompt provides as context the majority religion and language of each country. In the output, the model is asked to first state the *territory's* demographics and then pick a claimant country. This strategy can be viewed as eliciting a nationalistic line of thought,²¹ following: “if territory X follows the religion of country Y, and country Z follows another, wouldn't X more likely belong to Y?”

²¹For example, consider for language §4.5.1 and for religion §4.5.3.

4.4.5. Results: Prompt Modifications

Results for prompt modification experiments on GPT-4 are shown in the bottom rows of Table 4.6.²² Comparing *vanilla* to the best model using rank-classification (GPT-4_{DV}), we see that, as expected given the stronger model, GPT-4 improves for KB_CS, Con_CS, and Non_CS. However, Δ _CS shows it is more geopolitically biased than GPT-3_{DV} ($21.6 > 17.0$).

However, the model is able to successfully adopt the *UN peacekeeper* persona to mitigate geopolitical bias, lowering Δ _CS from 21.6 to 10.2. This is largely due to an increase in Non_CS ($63.2 \rightarrow 67.7$), which indicates that this strategy makes the model more likely to choose the controller when queried in a non-controller language. Interestingly, Cst_CS for ‘Unknown’ drops precipitously ($65.6 \rightarrow 56.3$), while slightly increasing Cst_CS over all territories. This suggests that when the UN peacekeeper persona does not have an opinion on a territory’s claimant, it becomes less consistent, while conversely becoming more consistent when it does have an opinion.

So too does the model successfully adopt the *nationalist* personas, increasing Δ _CS ($21.6 \rightarrow 33.8$). This strategy amplifies geopolitical bias in two ways, the language of interaction as well as the explicit statement, “As a citizen of c_i ”. We see that the responses are less consistent both overall and for ‘Unknown’.

For the *demographic reasoning* strategy, Δ _CS is about the same as for *vanilla*, suggesting similar geopolitical bias overall. The main effect is that KB_CS drops ($79.5 \rightarrow 70.8$). We give two potential reasons: first, the model may be making more errors, and second, and more likely, the demographics of many disputed territories do not line up with the actual controller situation, such as for the case studies §4.5.1 and §4.5.3 below. We also see that this strategy results in the highest Cst_CS for both ‘Unknown’ and overall. So while demographic reasoning increases consistency, it comes at the cost of the factuality attribute of robustness.

²²KB_CS is not defined for the nationalist prompt, as it considers queries in English, which may not be a language of any claimant for many territories.

4.5. Qualitative Analysis

In the prior section, we compared model responses over the entire BORDERLINES dataset. However, the status of each individual disputed territory is complex and has many players and relations at stake. Nor is each disputed territory equally inflammatory; compare a lone rock in the middle of a vast ocean versus a densely populated area on the mainland of a continent.

We therefore perform qualitative case studies on several notable disputed territories. We first provide a brief overview of the geopolitical situation of each dispute. Then, we look at how GPT-4’s selections change for each prompt strategy. Furthermore, we go beyond the strict multiple-choice format and qualitatively analyze the model’s full responses, which contain further explanations where it justifies its decisions. This allows us to see how a model qualifies, or not, its knowledge about the situation, depending on the probing strategy employed in the prompt.

4.5.1. Crimea

Crimea is a peninsula in Eastern Europe, in the north of the Black Sea. Of its population of 2.4 million, most people are ethnic Russians who speak Russian. While internationally considered part of Ukraine, it has been controlled by Russia since 2014 after its annexation.

Full text of responses are shown in Table B.2. For the *vanilla* setting, the responses differ. In Russian (ru), GPT-4 answers ‘Russia’, while adding a note about the international recognition for Ukraine. In Ukrainian (uk), it answers ‘Ukraine’, adding a note about the illegality of the annexation. This is also the case for the *nationalist* setting. For *demographic reasoning*, interestingly, we see a flip: the model responds ‘Ukraine’ in ru and ‘Russia’ in uk.

For the *UN peacekeeper* setting, both languages return ‘Ukraine’. Here, the geopolitical bias of interacting with an LLM in claimant languages has been mitigated.

4.5.2. Taiwan

Taiwan is an island in East Asia, in the western Pacific Ocean. It has a population of 23.9 million. The island has been controlled by the Republic of China (ROC) since 1945; the ROC is often

simply referred to as Taiwan. The People’s Republic of China (PRC) also claims Taiwan as one of its provinces. The ROC is the most populous country without official UN recognition, and its geopolitical status is extremely contentious²³ and heavily influences its politics (Chang et al., 2021).

Full text of responses are shown in Table B.3. For *vanilla* and *demographic reasoning*, querying in Traditional Chinese (zht, used in ROC) and Simplified Chinese (zhs, used in PRC) both return ‘ROC’. Adopting *nationalist* and *UN* prompts results in differing responses: PRC in zhs, and ROC in zht. These responses are all qualified by statements of claims of the other country.

4.5.3. Golan Heights

The Golan Heights is a region in the Middle East, with a population of 50,000. It is internationally recognized as part of Syria (to its east). However, it has been controlled by Israel (to its west) since a 1981 annexation. Its population is roughly evenly divided between Israelis, who follow Judaism and speak Hebrew (he), and Arabs, who follow Druze and speak Arabic (ar).

For *vanilla*, querying in ar and he both return ‘Israel’. The model qualifies its responses, in both languages, with statements on international recognition of Syria’s ownership. For the other settings, the model returns ‘Israel’ in he, and ‘Syria’ in ar. For *UN peacekeeper*, the model still selects ‘Israel’, despite the UN’s actual status. For *nationalist*, responses are as expected. Most interesting are the responses in the *demographic reasoning* setting. In Hebrew, the model reasons that Israel controls Golan Heights given its majority religion and language of Judaism and Hebrew; in Arabic, the model reasons that Syria controls it given Arabic and Islam (neither mentions Druze). Again, the even split of the population could go either way. So, by varying the language context, the model adapts the demographics to reinforce its perspective — a clear instance of geopolitical bias.

4.5.4. Discussion: Perspective-Dependent Questions

We acknowledge that humans are also not perfectly consistent in their responses, which may differ under different scenarios, or to different audiences. For example, consider perspective-dependent questions, such as “Who is the greatest soccer player?” or “Was Christopher Columbus the discoverer

²³To add another dimension, the ROC also claims mainland China (population: 1.4 billion). The situation is so complex that all other countries must follow a ‘One China’ policy to maintain diplomatic relations with the PRC.

of the New World?” In such cases, LLMs and humans alike may find it beneficial to adapt their answer based on one’s grasp of the current context and audience. So to have we discussed another beneficial case of language-specific tailoring, for multilingual cultural benchmarks (discussed in §2.4.2). However, our study of territorial disputes shows a specific category of perspective-dependent questions where over-adapting to an user’s inferred background through the language of interaction is happening, and this can amplify negative societal impacts. Our point is that LLMs, which are developed for, and aim to be, helpful assistants aligned to human values, should be carefully considered with respect to these politically-sensitive queries.

How should an LLM respond to queries on territorial disputes? In this work, we elicited responses from LLMs using a multiple-choice prompt, and techniques such as rank-classification to force them to output a single answer. We then compared this answer to a ground-truth KB, which was English Wikipedia here, but can be set to others as desired, and still go through the suite of evaluation metrics.

An well-aligned LLM should instead discuss the nature of the dispute, and all of the claimants’ positions, and the broader international contexts. Indeed, we found this that to be the case for GPT-4 when allowed to continue generating past the answer. Referring to the translated responses to Russian and Ukrainian prompts in Appendix Table B.2,

A) Russia, but it should be noted that Russia’s annexation of Crimea in 2014 is not recognized by the international community, and many countries consider Crimea to be occupied Ukrainian territory.

B) Ukraine

However, it is worth noting that Crimea was annexed by Russia in 2014, but the international community considers this annexation illegal and continues to recognize Crimea as part of Ukraine.

4.6. Conclusion

This chapter shows that LM responses to an underlying query on territorial disputes differ depending on the language of interaction. We characterize this issue as a matter of cross-lingual robustness. The three robustness attributes directly relate to this thesis’s goal of understanding how LMs navigate knowledge across languages. Factuality considers knowledge learned within a single language, English, consistency considers response differences from a language-agnostic perspective, and geopolitical bias considers response differences from a knowledge-informed perspective.

With reference to cultural benchmarks where answers should be language-variant (discussed in §2.4.2), BORDERLINES serves as a complementary benchmark. It also considers the cross-lingually different distribution of knowledge, but for scenarios where responses which should not be so readily tailored, lest we have LMs which further amplify divisions in viewpoints, increasing societal risk. With reference to factual benchmarks (discussed in §2.4.3) where answers should be language-invariant, BORDERLINES also follow this. However, prior benchmarks can be solved by training LLMs on the right answer in some language; i.e., data contamination of train and test is a major concern. In contrast, for BORDERLINES, the perspectives and arguments for and against a certain nation’s claims will always color some trained LMs responses. Thus, it can serve as an evergreen benchmark for cross-lingual robustness.

Another take on our findings is that the LM is largely drawing from knowledge discretely in each language, instead of transferring across them. This inspires a line of work which make the multilingual knowledge explicit, through the popular paradigm of retrieval-augmented LMs (LMs). This leads to our study in Chapter 6, where we investigate how the composition of retrieval databases across languages affects what perspectives are retrieved. We show on more modern LMs, such as GPT-4 and Llama 3, that the composition greatly affects the cross-lingual robustness of responses.

CHAPTER 5

RETRIEVAL STRATEGIES FOR DOMAIN-ADAPTED TRANSLATION

In Chapters 3 and 4, we studied LMs’ parametric knowledge of two tasks which were differently affected by the real-world distribution of knowledge across languages: complex reasoning has largely been seen during LLM training in English, and territorial disputes have differing perspectives across languages. We saw that incorporating contextual knowledge in the prompts assists in mitigating LLM’s cross-lingually inconsistent responses. While those chapters incorporated knowledge in a predefined manner, in this chapter we begin our exploration of how LLMs navigate contextual knowledge retrieved from external knowledge-bases (KBs), focusing on the *retrieval-augmented generation* (RAG) paradigm. As discussed in §2.3.1, the power of RAG is in the flexibility of the task-specific KBs, which an information retrieval (IR) system retrieves over to select the most relevant information to some input query.

In the LLM era, the RAG paradigm was introduced and popularized by Lewis et al. (2020). However, recalling our discussion of *machine translation* MT task in §2.1.1, there is a decades-long history of “example-based MT” approaches (Somers, 1999), where given a source language text, we assist a system (or a human) by retrieving similar source-target pairs from a large external corpora. With the evolution of the broader NLP field over the years, it is worth taking an analytical look into a contemporary implementation of RAG for MT, and what factors are key to success in improved translations.

This chapter is based on my paper “Leveraging Domain Knowledge at Inference Time for LLM Translation: Retrieval versus Generation”.²⁴ We focus on MT in domain adaptation settings for domains such as law, science, or medicine. We perform a unified analysis of the retrieval setting along two dimensions. First, we compare knowledge sources of *retrieval* over external KBs, or *generation* by an intermediate LLM prompting call. Second, we compare knowledge strategies, and

²⁴I led this project during a summer internship at Google. I implemented the codebase, ran the main experiments, and wrote most of the paper. In particular, Jiaming helped to run experiments, Eleftheria assisted with paper writing and results, and Colin was the senior author guiding the project direction.




Figure 5.1: Illustration of a source sentence and three possible translations, adapted from Tan et al. (2024). We see that the zero-shot LLM translation is semantically similar to the target, but uses different phrasing and style. The stylistic LLM translation captures both semantics and style well, and so scores better with respect to the target translation.

whether retrieval over KBs composed of *demonstration* pairs or of *terminology* dictionaries are most effective. We find that demonstrations particularly improve performance, while terminologies do not help as much. We also find some success with generation, showing that it can elicit LLM’s parametric knowledge of these domains, though retrieval still is most effective. Finally, we perform detailed analyses to understand the effectiveness of retrieved demonstrations, identifying the key guidance as coming from their style as well as their domain-specificity. Figure 5.1 illustrates the qualitative characteristics of a stylistically-improved translation compared to a zero-shot one.

Contemporary LLMs have strong zero-shot MT capabilities for general-domain texts (Zhu et al., 2024). Given a source text as well as a one-sentence instruction “Translate English to French” or “Translate Hindi to Hungarian”, LLMs will produce reasonably fluent and accurate responses. However, more bespoke use-cases and domains of translation remain an open challenge. Fortunately, LLM possesses instruction-following and in-contextual learning (ICL) abilities, which allow contextual information to be incorporated at inference-time. This contrasts with the statistical and neural MT eras, domain adaptation would typically involve an expensive continued training procedure on in-domain data (Freitag and Al-Onaizan, 2016; Thompson et al., 2019).

With LLMs, there is the promise of simple adaptation at inference time, and this flexibility is particularly exciting for adapting LLMs to translate specialist domains, such as legal or medical

texts. One promising technique is the retrieval of instance-specific demonstrations of translation from a bitext datastore for few-shot in-context learning, which has shown large improvements for domain-adapted MT (Agrawal et al., 2023; Tan et al., 2024), rivaling the performance of specialized nearest-neighbor MT systems (Khandelwal et al., 2021). LLMs have also been shown to make good use of bilingual terminology dictionaries for lexical translation hints (Ghazvininejad et al., 2023; Lu et al., 2023; Moslem et al., 2023).

Intriguingly, two recent approaches have forgone external resources in favor of querying an LLM to generate useful knowledge from its internal memory. First, the MAPS approach issues LLM queries for topics, terminology, and demonstrations based on the source text (He et al., 2024). Their terminology and demonstrations mirror the knowledge sourced from retrieval steps in earlier work. The idea is that the LLM has seen relevant information during pre-training, and would benefit from explicitly surfacing it before translation. Second, the step-by-step MT approach queries its LLM to translate and discuss idiomatic phrases before performing a complete translation (Briakou et al., 2024). However, both these works only consider the general domain, instead of specialized domains. The vast scale of LLM’s pretraining data means they have seen, in addition to general domain data, some amount of specialized domain data. Given this fact, as well as the difficulty of obtaining directly relevant external resources for each specialized domain, we are inspired to consider the applicability of internal memory approaches to domain-adapted MT.

In this work, we perform a comparative study of utilizing contextual knowledge for domain-adapted MT, in *strategies* – external retrieval vs. internal generation – and in *sources* – translation demonstrations and bilingual terminology. We consider three domains (law, medical, and Koran) from the popular multi-domain dataset (Aharoni and Goldberg, 2020), and experiment with two LLMs (Gemini 1.5 and Gemma-2). Our study addresses three main research questions:

RQ1. For improving domain-adapted MT, how viable is generation from an LLM’s parametric memory compared to retrieval from external resources?

RQ2. Likewise, how is adapting MT with demonstrations compared to with terminologies?




Figure 5.2: Illustration of the main MT settings. Top: the source text. Middle: the knowledge *strategies* (rows) and the knowledge *sources* (columns). Bottom: the prompt templates used. A worked example for each of the 4 settings will be provided at the end of its corresponding section.

RQ3. Given the effectiveness of demonstrations, can we attribute whether their stylistic or terminological aspects contribute the most to improved translations?

We discuss knowledge sourced from retrieval in §5.1.1 and from generation in §5.1.2. Comparisons

between terminology and demonstrations are enabled by our use of a silver terminology dictionary, built by LLM analysis of the same bitext used as the datastore of demonstrations (§5.2.2). This allows us to study demonstrations and terminology as alternate views into the same base data in the retrieval setting. We address RQ1 and RQ2 with the results in §5.2.3. We explore RQ3 through several analyses in §5.3; the main takeaways are that retrieved demonstrations mainly provide hints of target style rather than terminology, and that generated ones can viably boost performance, albeit to the same level as static domain-specific demonstrations.

5.1. Leveraging Domain Knowledge

Comparisons between representations of domain knowledge are enabled by our careful prompting setup which decouples the *source* and *strategy*, as sketched in Figure 5.2. Bilingual terms, whether retrieved externally or generated by an auxiliary LLM call, feed into the same translate-with-terms prompt, and likewise for demonstrations. On *sources*, retrieval leverages resources such as datastores and dictionaries, while generation elicits information from an LLM’s own parametric memory. On *strategies*, demonstrations provide source-target example pairs, whereas terminology focuses on domain-specific lexical items. This section details the integration of these strategies and knowledge sources within our experimental framework for domain adaptation of LLM MT.

5.1.1. Knowledge from Retrieval

We describe two approaches to retrieve domain knowledge from external resources: *demonstration retrieval* and *terminology lookup*. Demonstration retrieval has the model *implicitly* learn from the characteristics of the exemplars, both style and terminology. Terminology lookup has the model *explicitly* see which source terms are important and also how to translate them. Both approaches have been time-tested throughout the statistical, neural, and now LLM eras of MT.

Prior work The use of demonstrations of translation pairs has a long history. Some of the oldest data-driven approaches in MT have as their first step finding the most relevant examples from a bilingual translation memory. This idea has been used for computer-aided translation (Yamada, 2011), example-based MT (Somers, 1999; Lepage and Denoual, 2005) and statistical MT (Koehn and Senellart, 2010).

Several recent papers have studied what constitutes effective demonstration retrieval for MT with LLMs. Agrawal et al. (2023) found a strong baseline to be example-specific BM25 retrieval of bitexts, which can be strengthened further by re-ranking for lexical diversity. Tan et al. (2024) use a much larger LLM, and show that BM25 retrieval of target sentences alone can compare favorably with both sides of bitexts. Conversely, in the general-domain, researchers have found that a demonstration’s quality matters more than its proximity (Vilar et al., 2023; Zhang et al., 2023a).

Resource requirements These methods, while effective, are expensive, as they require the existence of high-quality and domain-specific resources. The former requires a large pool of bitext demonstrations, and the latter requires the creation of a term-rich bilingual dictionary.

Demonstration Retrieval

Demonstrations are provided as exemplars in the prompt to facilitate in-context learning (ICL) (Brown et al., 2020; Patel et al., 2023). These exemplars can either be *static*, the same across all instances, or *instance-specific*, in which different exemplars are retrieved for each instance to provide specific guidance and hints.

The typical setup for demonstration retrieval for MT is as follows. Given a source text, we find k closest source-side matches in an external datastore, using some similarity metric, such as BM25 or cosine similarity of embedding vectors. Then, we include in the LLM prompt these k source texts paired with the gold target translations.

Our Setup For our few-shot implementation, we design a simple prompt (Figure C.2).²⁵ We use $k=3$ exemplars,²⁶ and retrieve using the BM25 metric. Our datastore, derived from the train split of multi-domain, has 16,775 demonstrations for Koran, 234,352 for medical, and 464,295 for law.

Worked Example We are given a source text $s = \text{“Welche Risiken sind mit Poulvac Flufend H5N3 RG verbunden?”}$, as well as a datastore D which indexes paired German-English texts. We use a similarity metric to calculate the distance of s from all German texts in D . For the k closest sentences on the German side, we will use those German-English pairs as the exemplars to be

²⁵All prompts used in this work are given in Appendix §C.2.

²⁶Prior work often chooses $k \geq 10$. As we find COMET for $k=3$ and $k=10$ differ by ~ 0.3 we thus choose $k=3$ to fairly compare to the 3 generated demonstrations in a later setting.

included in the prompt. As shown in Figure 5.2, these are quite lexically close to the *s*; for example, the first pair is (“Welche Risiken sind mit *Luminität* verbunden?”, “What risks are associated with Luminität?”). This illustrates the domain-specificity of the external KB for this dataset.

Terminology Lookup

Intuitively, one of the major challenges when translating in a specialist domain is the adaptation to domain-specific terminology. Especially in high-stakes legal, medical or business domains, precision of terminology can be crucial. Bilingual dictionaries of terminology are therefore likely sources of useful external knowledge to add into an MT system. These resources can be easier to construct than the large datastore of translation demonstrations needed in §5.1.1. In fact, the construction of a clear terminology may very well be a prerequisite to creating human translations.

Prior Work Improving translations with terminologies has been heavily studied. In the statistical and neural eras, solutions could take the form of incorporating dictionaries into training Wu et al. (2008), or controllable MT systems that respect example-specific terminology constraints included in the input (Post and Vilar, 2018; Wang et al., 2022a). In the recent LLM era, terminology constraints have been studied at two WMT shared tasks (Alam et al., 2021; Semenov et al., 2023). These approaches illustrate two different motivations for the use of terminology dictionaries in MT: the dictionary can be viewed as a useful source of domain-specific information, or as a set of constraints that must be followed consistently. Our work aligns with the former motivation, viewing bilingual terminologies only as hints to improve overall quality.

With the advent of LLMs, terminologies can be included in the prompt, with additional instructions on their usage. Most LLMs follow these instructions well, as shown at the WMT23 shared task on terminology (Semenov et al., 2023). For example, Moslem et al. (2023) find that for the COVID-19 domain, a prompt using retrieved terminologies significantly boosts term success rate and also improves human evaluation scores. Other works have explored how to more effectively format dictionaries (Lu et al., 2023; Ghazvininejad et al., 2023).

Our Setup Since the multi-domain dataset does not have a provided domain-specific terminology, we derive one from the multi-domain training set, as described in §5.2.2. Keeping with our theme of

providing hints rather than constraints, the dictionary gives a list of possible translations for each source term, each licensed by at least one example in the training set.

With this dictionary in place, we look up terms by exact lexical match to the source text currently being translated, and include any matches in our prompt for translation with terminology (Figure C.3). The LLM is instructed to pick the most appropriate translation among the choices, given the source. Note that the translation prompt also includes three domain-specific examples of how to translate with terminologies.

Worked Example We are given a source text s “Welche Risiken sind mit Poulvac Flufend H5N3 RG verbunden?”, as well as a domain-specific terminology T which maps German phrases to English phrases. We search over all German phrases in T , and identify those that appear in s , and add it to a dictionary d^s , which is then included in the prompt. As shown in Figure 5.2, d^s is {“Risiken”: [“risks”], “Poulvac FluFend H5N3 RG”: [“Poulvac FluFend H5N3 RG”], [“verbunden”: [“connected”, “associated”]]}. Observe that the mappings are one-to-many, as the terminology is created over the entire training datastore. Also, the mappings can be either translations or be identical, the latter of which indicates that some technical term should be kept as-is. For example, “Poulvac FluFend H5N3 RG” is an inactivated avian influenza vaccine for chickens and ducks that protects against the H5N1 and other virulent strains of avian flu; this term is used in both German and English.

5.1.2. Knowledge from Generation

While external knowledge retrieval demonstrably benefits knowledge-intensive NLP tasks, whether it is truly necessary for domain-adapted MT still warrants investigation, given that LLMs are explicitly trained on massive corpora including texts from specialist domains. Therefore, we investigate whether leveraging LLMs’ internal parametric memory can offer comparable benefits, and thus circumvent the costly acquisition and curation of external resources. This approach effectively simulates external retrieval by prompting the LLM to generate relevant information.

Prior work has explored several methods to leverage an LLM’s parametric knowledge to improve MT quality, either post-translation, or pre-translation. Our work follows is motivated by two prior pre-translation approaches.

He et al. (2024) propose a human-like translation process, where they separately prompt LLMs for 3 aspects related to a source text (demonstrations, topics, and terms). Directly using these generated knowledge pieces in another LLM interaction is insufficient, and so they rely on an external quality estimation (QE) method to select among candidates, improving general domain MT quality. Our generation setting also use demonstrations and terms, but without any external feedback from QE.

Briakou et al. (2024) propose a method to model the LLM translation process step-by-step. Their 2-step approach has an LLM first perform research on idiomatic expressions, then perform the full translation. For document-level MT datasets, they find this consistently outperforms zero-shot MT.

Resource requirements By design, the generation setting requires almost no external resources. The approaches discussed below only required us to manually create a handful of static exemplars for each subtask, which are used for all of its prompts.

Demonstration Generation

We author a prompt to generate demonstrations (Figure C.5). For each domain, we provide 3 example demonstrations for 2 static, real source sentences. This is inspired by the demonstration aspect of He et al. (2024), but we elicit 3 demonstration pairs at a time instead of 1.

Best practices To easily parse the 3 demonstration pairs, we ask for a prescribed JSON output format. We also find that providing static few-shot exemplars of the demonstration task is key to both diversity among the 3 demonstrations, and output format adherence. We use a different set of exemplars for each domain, drawn from the train set. We perform ablations on the contributions of different aspects of generated demonstrations in §5.3.2.

Worked Example We are given a source text s “Welche Risiken sind mit Poulvac Flufend H5N3 RG verbunden?”. As described above, we use an intermediate LLM call (with exemplars) to generate synthetic few-shot source-target pairs, which are utilized in the few-shot translation prompt. As shown in Figure 5.2, these are semantically close to the s , but tend to be less lexically close than the actual retrieval setting. For example, the first pair is (“Was sind die möglichen Gefahren von Ibuprofen?”, “What are the possible dangers of Ibuprofen?”).

Terminology Generation

We design a prompt to generate terminologies from a single source sentence (Figure C.6), also using 2 static, real sentence pairs for each domain. This follows in the spirit of the research step of Briakou et al. (2024), where they explain this as having the LLM perform intermediate reasoning about hard-to-translate parts. However, there are several differences resulting from their focus on document-level MT. We ask generally for terminologies, while they ask specifically for idiomatic expressions, which are more prevalent in long documents. We also prescribe a JSON format (same as for retrieved terms), while theirs allows for free-form output.

Best practices We again found that best performance is achieved with static, domain-specific few-shot exemplars of the terminology task, and the prescribed JSON format.

Worked Example We are given a source text s “Welche Risiken sind mit Poulvac Flufend H5N3 RG verbunden?”. As described above, we use an intermediate LLM call (with exemplars) to generate a synthetic German-English dictionary δ^s , which is then included in the prompt. As shown in Figure 5.2, δ^s is {“Risiken”: [“risks”], “Poulvac FluFend H5N3 RG”: [“Poulvac FluFend H5N3 RG”], [“verbunden”: [“connected”]]}. For this simple sentence, the synthetic and retrieved dictionaries are very similar, save for “verbunden” having one and two English translations respectively.

5.2. Experimental Setup and Results

5.2.1. Experimental Setup

Dataset We experiment with the multi-domain dataset (Aharoni and Goldberg, 2020), using the filtered version provided by Tan et al. (2024), with 3 domains: law, medical, and Koran. Multi-domain covers the German-English (de-en) direction, and consists of dev and test sets, with ~ 2000 entries per domain, as well as a train set with 1M+ entries.

LLMs We perform experiments with two LLMs, the open LLM Gemma-2 27B IT (Gemma Team, 2024), and the proprietary Gemini 1.5 Pro (Gemini Team, 2024). We thus can investigate which settings, if any, are more effective with the smaller model vs. a much larger model respectively.

Evaluation We perform zero-shot MT as a baseline. The four MT settings described in §5.1 are: retrieved demonstrations, retrieved terminologies, generated demonstrations, and generated terminologies. Following Vilar et al. (2023), we use a learned neural metric, COMET (Rei et al., 2022), which assesses MT quality from all three parts of the source input, the source output, and the reference output. While prior work also considered the lexical metric BLEU (Papineni et al., 2002), we exclude it as it was overly sensitive to minor rephrasing. This is in line with studies that show neural metrics correlate much better with human judgments of LLM translation quality (Freitag et al., 2021; Kocmi et al., 2021).

5.2.2. Terminology Dictionary Creation

Our multi-domain test scenario does not come with bilingual terminology dictionaries for its domains. However, we can create them from the provided training split, following the methodology in prior work (Moslem et al., 2023; Semenov et al., 2023).²⁷ We design a prompt (Figure C.7) to extract terminologies from a given source-target text pair, providing 5 static exemplars to demonstrate what is meant by “terminology”. We then apply this to each pair from the train split. Then, we aggregate all the output terms, to get one large dictionary with one-to-many mappings.²⁸ We create a separate global terminology for each of the three domains.

Given the large size of the training split (700K entries), we make two adjustments to reduce the number of model calls. First, we batch five test pairs at a time into a single call. Second, we consider only the subset of train entries that were ever retrieved by BM25 over the test set (i.e., the entries that are actually relevant); this constitutes 70K entries, or about 10% of the total entries.

Note that the train split is also used for demonstration retrieval, therefore enabling a controlled comparison between the two external knowledge sources. Furthermore, unlike prior work using one-to-one terminology mappings, we explore a more realistic one-to-many scenario, with all possible translations in the prompt for the LLM to select.

²⁷We did not perform human post-editing due to the datastore’s size (700K), but we note in an experiment by Moslem et al. (2023), they found humans rated 95%+ terms as accurate.

²⁸For quality control, we kept only entries where 1) target terms have >10% usage and 2) both target and source sides of a term match.

Domain Knowledge?		Gemini 1.5 Pro						Gemma-2 27B IT					
		Law		Med.		Koran		Law		Med.		Koran	
🔍	zero-shot	86.6		88.2		76.3		84.8		85.2		75.1	
📖	retrieved	86.9*	↑ 0.3	88.5	↑ 0.3	74.9	↓ 1.4	85.9*	↑ 1.1	87.8*	↑ 2.6	74.6	↓ 0.5
	demos	89.5*	↑ 2.9	90.2*	↑ 2.0	76.7*	↑ 0.4	88.2*	↑ 3.4	89.9*	↑ 4.7	76.8*	↑ 1.7
⚙️	generated	86.5	↓ 0.1	87.8	↓ 0.1	76.2	↓ 0.4	85.0	↑ 0.2	87.4*	↑ 2.2	76.3*	↑ 1.2
	demos	87.2*	↑ 0.6	88.8*	↑ 0.6	76.7*	↑ 0.4	86.0*	↑ 1.2	88.1*	↑ 2.9	76.1*	↑ 1.0

Table 5.1: MT results with domain knowledge presented as *demos* or *terms*, and either *retrieved* from external sources or *generated* by the LLM. We report the COMET22 metric for two models. Statistically significant improvements over the zero-shot baseline are marked with *.

5.2.3. Results

Table 5.1 presents our primary results, comparing LLM translation enhanced with domain-specific knowledge in the form of translation demonstrations or bilingual terminology, with the artifacts derived from either external retrieval (§5.1.1) or internal generation (§5.1.2). First, in line with prior work, we confirm that retrieved demonstrations improve over zero-shot across models and domains studied. We next describe the three main findings.

Demonstrations outperform terminology For all models and domains studied, knowledge provided in the form of demonstrations consistently outperforms terminology. For Gemma, we see that all settings improve performance,²⁹ but the improvements from demonstrations are markedly larger. The differential is more pronounced for Gemini, which starts from a much stronger baseline than Gemma. Terms, either retrieved or generated, do not provide much of a boost over zero-shot for Gemini, while demonstrations result in significant improvements. The takeaway for this finding is that for weaker models, providing domain knowledge from any source or strategy is beneficial. Conversely, stronger models do not benefit from domain-specific terminology, but only from more complete demonstrations of the task.

Retrieval outperforms generation The second notable trend across models and domains is that retrieval consistently outperforms generation. With Gemma, demonstration generation outperforms zero-shot by +2.3 (averaged across domains), while retrieval further improves to +3.4. For the more powerful Gemini, the differential is larger – demonstration generation outperforms zero-shot by +0.5, while retrieval by +1.8.

²⁹To explain the outliers for Koran (−1.4, −0.5), our manual analysis found term inconsistency – high-frequency source terms mapped to multiple, equally-valid target terms.

Generated domain-specific demonstrations boost weaker model’s translations Taking the prior two findings together, we can bootstrap domain-adapted MT knowledge from an LLM’s own parametric memory, with the two-stage approach of first generating demonstrations, then translating. This improvement especially pronounced with Gemma (+2.3 vs. +0.5 over respective zero-shot). In fact, this empowers a smaller model (Gemma) to close the gap with a larger model’s (Gemini) zero-shot results, as can be seen by comparing, in Table 5.1, the bottom left and top right rows. The gains in medical (+2.9) and Koran domains (+1.0) result in statistically equivalent scores. Law domain incurs a decent gain (+1.2), but still is below Gemini ($86.0 < 86.6$). It is worth emphasizing that acquiring extensive resources for novel specialist domains is expensive; but this straightforward approach can be effective.³⁰

Comparisons with Prior Results We can also compare our results with demonstrations to those from the recent study by Tan et al. (2024), who use the `gpt-3.5-turbo-0301` LLM. Their zero-shot results are most comparable to Gemma’s: 84.4, 86.2, 75.1. Their results for retrieved demonstrations are also comparable: 88.2, 89.6, 76.5. The other 3 settings, retrieving terms and both generation ones, are new to our work – and we re-emphasize here the value of our controlled setting in facilitating fair comparison between them all.

5.3. Analysis

Demonstrations (both retrieved and generated) are by far the most effective domain adaptation strategy we explored, providing a large boost to both LLMs. The following sections seek to understand where the gains are coming from. We begin by analyzing retrieved demonstrations to disentangle contributions from style vs. terminology (§5.3.1). Then, we investigate the impact of varying factors of in-context learning for generating demonstrations (§5.3.2). Finally, we study how generated domain knowledge can be distilled at test-time from larger to smaller models (§5.3.3).

5.3.1. Retrieved Demonstrations: Contributions from Style vs. Terminology

What exactly is being conveyed by the retrieved demonstrations? Here, we will take advantage of our careful experimental setup, where our bilingual terminology is derived from the same parallel

³⁰Note that the experiment from §5.3.2 shows that domain-specificity is the main contributor, rather than proximity to the current instance. These generated demos are only as effective as real static, but domain-specific demos.




Figure 5.3: Illustration of our process to decompose the contributions of retrieved demonstrations into style and terminology. The source-target term pairs are extracted and aggregated into a local terminology. The remaining tokens are then used as the style templates, with the terms masked.

text used for demonstrations, to disentangle whether demonstrations are more valuable because they assist with proper *terminology* translations in context, or with matching the *style* of the corpus.

The core idea behind this experiment is that we can use the same technique to extract bilingual terminology pairs from a translation demonstration (§5.2.2), but instead of running it on the whole training corpus, we can run it only on the k demonstrations retrieved to match the current source sentence. This gives us a *local* terminology, as opposed to a *global* one. Crucially, where the global terminology would present the union of all possible target language translations found throughout the training set for a given source term, the local terminology only presents translations licensed by the k demonstrations. This allows it to take advantage of any disambiguating context in the demonstrations to create more relevant term translations.




Figure 5.4: MT results comparing zero-shot and external retrieval (ret terms) to terms from demonstrations and style from demonstrations.

We then define style templates as the inverse – the remaining tokens, with the bilingual terms on both sides replaced with a <MASK> mask token. For this, we use a similar prompt as for demonstration retrieval (Figure C.4), which explicitly instructs the LLM to not generate mask tokens in its output. Upon manual inspection, these masks appear quite thorough, with most anything that could be considered terminology being masked out. The decomposition process is depicted in Figure 5.3.

Results We carry out the decomposition experiment using Gemini. Figure 5.4 presents our results. Compared to zero-shot, using local terms (terms from demonstrations) more than doubles the gains of global terms (retrieved terms). However, style templates (style from demonstrations) further narrow the gap to retrieved demonstrations by 60% (law) and 75% (medical).

The results indicate that guidance from domain-specific terminology account a smaller proportion of gains: only 0.8 out of 2.9 (law), and 0.5 out of 2.0 (medical) point improvements from the retrieved demonstrations. Meanwhile, the *style from demonstrations* scores almost perfectly account for the remainder. This is a strong indicator that the majority of the value of retrieved demonstrations comes from matching the publication style of these corpora, rather than carrying out adaptation to a medical or legal domain. That is, we are doing domain adaptation, but it is to a much more narrow domain than is usually discussed.

Discussion: Style Transfer for MT These results agree with and reinforce conclusions from recent work. Tan et al. (2024) perform a targeted study into translation style, in the same setting of

the multi-domain dataset and a recent frontier LLM. Their findings between zero-shot and few-shot concur – despite the COMET difference of 2.7, nevertheless zero-shot translations “have already conveyed the *semantic meaning* of the source sentence, albeit with some variations in *lexical choices* and sentence structure.” Their style learning method, which retrieves related target sentences from a monolingual target corpus, achieves 70% of few-shot’s gains. However, by only removing the source side of the demonstrations, the exemplars still implicitly provide both style and terminology hints. We build upon this discussion by providing a precise, alternative definition of “style” as anything outside of terminology. This in turn allows us to cleanly decompose the tokens from each demonstration into two subsets, and assign credit accordingly.

5.3.2. Ablation on Generated Demonstrations

Generation of demonstrations (**generate demos**) is by far the most successful of the two approaches. As described in §5.1.2, we made several decisions here: 1) using domain-specific exemplars; 2) using the intermediate generation of demonstration step; 3) in that step, selection of the ICL exemplars. We explore the impact of the decisions by comparing the **zero-shot**, **retrieved demos**, and **generate demos** results to the following ablations:

Static few-shot Drop the generate demo step, and use the 2x3 domain-specific examples³¹ directly as static demonstrations of translation. This investigates the impact of domain-specificity alone.

No ICL With the generated demo step, but remove all exemplars from that step’s instructions. This investigates the impact of ICL at all.

General ICL With the generated demo step, but use the 5x1 general-domain examples from He et al. (2024) instead of the 2x3. This investigates the impact of the domain-specificity of ICL.

Results are shown in Table 5.2. First, we consider ablation results on Gemini. There is a large drop between **zero-shot** and ‘no ICL’ (e.g., 88.2 → 83.8 for medicine). Our manual analysis of a few ‘no ICL’ outputs finds that the generated demonstrations on the target side are often quite lexically close; we hypothesize these are unhelpful and affect downstream translations. Comparing general ICL to **generated demos**, roughly half the value of demonstration generation can be retained with

³¹2x3 means there are two example source texts, which are each followed by three example translation pairs.

Setting	Domain-specific	Generate demo step	Demo # ICL	Gemma-2 27B IT			Gemini 1.5 Pro		
				law	med.	Koran	law	med.	Koran
zero-shot	N/A	✗	N/A	84.8	85.2	75.1	86.6	88.2	76.3
<i>static few-shot</i>	✓	✗	N/A	86.3	88.2	76.3	87.2	89.1	76.4
retrieved demos	✓	✗	N/A	88.6	89.9	76.7	89.3	89.9	76.4
<i>no ICL</i>	N/A	✓	0	85.2	87.5	75.1	83.9	83.8	75.6
<i>general ICL</i>	✗	✓	5x1	85.7	87.8	75.6	86.9	88.5	76.0
generated demos	✓	✓	2x3	86.0	88.1	76.1	87.2	88.8	76.7

Table 5.2: COMET22 Results for the study on demonstration generation, using Gemma and Gemini. The *italicized* rows are ablations, and the `monospace` rows are copied from Table 5.1.

general ICL. However, comparing ‘static few-shot’ to generate demos (rows 7 & 10), both achieve similar scores across domains. This adds a caveat to our earlier findings, suggesting the domain-specificity of the generated demos is more important to downstream MT than the demos alone.

Now, we consider ablation results on Gemma. Interestingly, unlike for Gemini, for Gemma even the ‘no ICL’ setting improves upon zero-shot (85.2 $\rightarrow$ 87.5 for medicine). We observe that, compared to `generated demos`, ‘General ICL’ slightly underperforms it, while ‘static few-shot’ matches it. This again underscores the value of the demonstration stage in improving the smaller LLM’s translations, as well as the key role of domain-specificity.

Our results add insight into two formerly disparate findings. Prior work emphasized two factors for ICL exemplars: lexical coverage within a domain (Agrawal et al., 2023), and their quality (Vilar et al., 2023). Our finding here provides evidence that, for current LLMs with strong zero-shot MT performance, the primary value of ICL is in domain-specificity. Quality exemplars can be sourced as validly from static few-shot exemplars as from generated demonstrations.

5.3.3. Cross-LLM Knowledge Generation

For the two generation-based settings, the same LLM is used in both the generation stage and the translation stage. To further understand how generation quality affects the final performance, we conduct additional experiments to reuse the generated demonstrations or terminology from Gemini 1.5 Pro to prompt the Gemma 2 27B model for translation.

Table 5.5 shows that generation of both demonstrations and terminologies from Gemini present significant gains in all three domains. This shows that higher-quality generated knowledge result in higher-quality translations. The larger Gemini model’s knowledge can be effectively distilled to the smaller Gemma model, at inference-time, through its translation demonstrations.

strategy	gen. LLM	law	med.	koran
⚙️ demos	Gemma	86.0	88.1	76.1
	Gemini	86.9*	88.6*	76.6*
⚙️ terms	Gemma	85.2	87.1	75.7
	Gemini	85.8*	87.5*	76.4*

Figure 5.5: Results for the ablation on generation strategies. The translation model is always Gemma. * marks significant improvements when using Gemini over Gemma for generation.

5.4. Conclusion

In this chapter, we studied domain adaptation for MT with LLMs, a problem which intuitively speaking should be well addressed by prompting-time adaptation. Building upon prior work which incorporates domain-specific knowledge into prompts, we performed a thorough study into how this knowledge can best be acquired in terms of strategy, demonstrations or terminologies, and sources, retrieval or generation.

It is likely that LMs have seen much of the data during training for this older MT benchmark. Still, we show that explicitly surfacing knowledge improves performance, and better yet still is in retrieving external examples. Tying back into the thesis, we found that knowledge is most effective when represented as demonstrations rather than terminologies. Our ablation experiments showed that key value of demonstrations is in stylistic guidance. For future work, we advocate for development of updated domain-adapted MT benchmarks which can challenges even the largest LMs’ broad base of parametric knowledge.

CHAPTER 6

MULTILINGUAL RETRIEVAL AUGMENTED GENERATION FOR GEOPOLITICS

Recall our exploration of the territorial disputes task in Chapter 4, where we saw that LLM responses lacked *cross-lingual robustness*: for the same underlying factual query, their answers are inconsistently when varying the language of interaction. We measured this in three attributes: factuality, consistency, and language fairness. In this chapter, we investigate the extent to which RAG can mitigate or amplify each attribute of robustness of the territorial disputes task. In contrast to the MT task explored in Chapter 5, where KBs were custom-built to be domain and task-specific, for the territorial task we draw KBs from the real-world source of Wikipedia, whose widely accessible and well-structured articles, written in many languages, are a trusted information source by end-users and researchers alike.

This chapter is based on the paper “Multilingual Retrieval Augmented Generation for Culturally-Sensitive Tasks: A Benchmark for Cross-lingual Robustness”.³² Given the effectiveness of RAG in improving response accuracy and reducing hallucinations for various NLP tasks (as discussed in §2.3.1), it is natural to ask whether the same improvements apply when extending the territorial disputes task from direct prompting of LLMs to RAG systems. However, recall the fact that across languages, the real-world distribution of knowledge differs in coverage of topics, and in perspective on issues. Performing retrieval over KBs in various languages therefore can specify what perspectives LLMs have access to, accordingly influencing their responses, and either mitigating or amplifying cross-lingual robustness issues. To highlight the main finding of this work, we find that RAG over multilingual documents elicits the most robust responses, while RAG over documents only in the user query language can be detrimental to consistency across languages.

Given the promises of AI systems to empower information access to users across the world, no matter the languages they speak, it is important to study how we can mitigate the risks of systems in

³²I led this project, and was responsible for organizing the team and overall direction. Fiona, Samar, and Adwait helped with database setup, IR systems, and paper writing. The other students participated through a research practicum course, and assisted with human and LLM annotations, running experiments, and creating diagrams.




Figure 6.1: Illustration of 2 cross-lingual RAG prompts from the BORDIRLINES dataset, showing differences in a) retrieved documents by the IR model and in b) responses by the LLM. For a territory such as “Ceuta”, we create several prompts across languages and IR modes. Evaluation of *cross-lingual robustness* is over the set of responses.

promoting divisive viewpoints across cultures, such as on the politically-sensitive issues of territorial disputes. In Chapter 4, we studied the parametric knowledge of LLMs on these disputes. In this chapter, we explore how RAG systems can navigate contextual knowledge in various multilingual settings to inform their responses on these disputes. Our primary experiment studies the language *composition of KBs*, comparing between retrieval over a single language – in the query language (qlang) – vs. retrieval over multiple languages – in the languages of all countries involved in a conflict (rel_langs). Our further experiments investigate each step of the RAG pipeline, from the KB, to the IR system, to the LM, is affected by issues of multilinguality, in particular the different levels of language preferences IR systems and LLMs possess.

Our study provides deep insights into how RAG systems perform in the multilingual setting. Related studies consider only simple factual recall tasks (Asai et al., 2022; Clark et al., 2020), or synthetic information-conflict scenarios (Sharma et al., 2024; Wu et al., 2024); furthermore, all of these works consider only a selection of high-resource languages. Our experiment sources documents from 49

languages, of varying perspectives and levels of resources, and applies RAG to the politically-sensitive task of territorial disputes. This chapter makes the following contributions:

- We design the **BORDIRLINES** dataset, which pairs queries from an existing territorial dispute dataset with 7.4k retrieved Wikipedia passages across 49 languages.
- We define five *IR modes* for retrieving multilingual sets of documents, each of which reflects real-world information access needs.
- We use **BORDIRLINES** to systematically evaluate the *cross-lingual robustness* of various popular RALMs, as shown in Figure 6.1. Notably, retrieving over multilingual documents *improves cross-lingual consistency* and *decreases geopolitical bias* vs. over in-language documents.
- We find that for LLM responses’ citations of documents, low-resource languages demonstrate much wider variability in citation rates than high-resource languages.
- We perform further analyses on all attributes of cross-lingual RAG. On IR, we find a preference toward retrieving query-language documents; on document contents, we find LLM responses can selectively interpret the same documents to fit their own viewpoints.

6.1. Preliminaries

6.1.1. Cross-lingual RAG

We construct a RAG pipeline consisting of a *retrieval* stage using an IR system and a *generation* stage using an LLM. Both use off-the-shelf multilingual models, without further fine-tuning. The cross-linguality in the setting is twofold: in the multilingual queries, and in their pairing with documents retrieved from a multilingual database.

Territorial Disputes Task We adapt the task formalization from BORDERLINES (Li et al., 2024b). We consider a territory t that has a set of claimant countries $C = \{c_1, c_2, \dots, c_k\}$. For each t , BORDERLINES has a multilingual query set (MQS) which consists of multiple-choice questions in each of the claimant’s languages: $\text{MQS} = \{q_{l_1}, q_{l_2}, \dots, q_{l_k}\}$.

Retrieval Given a query q_l , we use an information retrieval (IR) system to retrieve the top-k most relevant documents $D = \{d_1, d_2, \dots, d_k\}$ from an external KB. In our setting, the KB is a multilingual concatenation of databases in different languages, depending on the IR mode (described ahead).

Generation A query q_l and relevant documents D are composed into a prompt P . The LLM’s task is to generate an answer a for P , informed by both its parametric memory and the external retrieved information. The generated a should also be in l . A multilingual response set (MRS) consists of all responses for a territory t ; we evaluate cross-lingual robustness over an MRS.

IR Mode	Retrieval Languages L
qlang	$[l_q]$
$\hookrightarrow$	monolingual IR in query lang
rel_langs	$[l_q, l_{en}] + [L_{\text{relevant}}]$
$\hookrightarrow$	multilingual IR in all relevant languages
qlang+en	$[l_q, l_{en}]$
$\hookrightarrow$	multilingual IR in query lang and en
en_only	$[l_{en}]$
$\hookrightarrow$	cross-lingual IR in en
swap_docs	$[l_{\text{not_q}}]$
$\hookrightarrow$	cross-lingual IR in non-query lang

Table 6.1: For each IR mode, this depicts the set of retrieval languages L for a query q_l . $\hookrightarrow$ indicates the descriptions.

6.1.2. Information Retrieval (IR) Modes

We define five IR modes, each specifying the languages of documents to retrieve, given a query q_l .

These are formalized in Table 6.1. Each IR mode is motivated by some real-world information access needs a user may have. **qlang** retrieves documents in the query (user’s) language. **rel_langs** retrieves documents in the user’s language, English (en) and any other relevant language. For example, the relevant languages for Ceuta are Arabic (ar) and Spanish (es), so the **rel_langs** mode would have $KB = KB_{en} \cup KB_{ar} \cup KB_{es}$. **qlang+en** retrieves information in the query language and English. **en_only** retrieves only English. Finally, **swap_docs** adversarially selects documents from a non-query language. For Ceuta, this would be ar documents with an es query, or es documents with an ar query. This tests whether retrieval over another language’s documents can ‘flip’ responses.

6.1.3. Cross-lingual Robustness for RALMs

Recall from §4.3 that cross-lingual robustness is defined in three attributes: *factuality*, *consistency*, and language fairness (*geopolitical bias*). It is defined over a set of LM responses. For the territorial disputes task, each attribute is measured through the concurrence score (CS) metrics of Knowledge-base CS (KB_CS), Consistency CS (Cst_CS), and Delta CS (Δ _CS) respectively. In this work, we adopt them to the multilingual RAG setting, which considers how answers are affected by multilingually-sourced documents in the prompt. All evaluations including factuality are now done multilingually, as while the setup with comparing underlying queries across languages is the same, the documents




Figure 6.2: Illustration of data collection for a BORDIRLINES entry on the territory “Ceuta”. Depicted is the `rel_langs` IR mode, which considers Wikipedia documents in $L = [\text{en}, \text{ar}, \text{es}]$. For each query, we use IR to obtain the top-10 most relevant documents. Each document is then annotated for its *relevance*, and its *territorial viewpoint*, which claimant it supports.

for each language query’s prompt vary depending on the retrieval results, even within an IR mode.

6.1.4. Dataset Description

Figure 6.2 illustrates the dataset construction and annotation. Appendix Table D.1 provides the detailed, per-language dataset statistics. In total, there are 19,916 unique query-document pairs,

drawn from 7,436 passages from 905 Wikipedia articles.

6.1.5. Dataset Construction

Source: Wikipedia We source the documents from Wikipedia, an online encyclopedia available in multiple languages. We do this for two reasons. First, Wikipedia is the primary source used by prior benchmarks on retrieval and QA (Asai et al., 2020; Zhang et al., 2023c). Second, Wikipedia articles go through a set of community-enforced standards, such as neutral POV. Therefore, we expect articles in different languages to report from similarly “neutral” viewpoints, and minimize culturally specific biases. Still, as articles are independently authored, this is not guaranteed. This motivates us to annotate each document’s content as described ahead.

Information Retrieval Setup We segmented articles into paragraphs by splitting articles on double newlines, and used these paragraphs as the documents (we use passage and document interchangeably). For a given query q_l , in lieu of retrieving over the entirety of the Wikipedia KB , we index only the articles most likely to contain relevant information, the territory and the claimants. This greatly reduces the search space and the likelihood of retrieving non-relevant documents. For example, on Ceuta we index the articles on Ceuta, Morocco, and Spain. Also, as described in 6.1.2, we use q_l to retrieve over a KB whose language composition depends on the IR mode.

6.1.6. Annotating Documents for Content

To improve the BORDIRLINES benchmark’s usability for IR, each query-document pair is annotated for two content factors: *relevance* – whether it directly discusses the status of a territory – and *territorial viewpoint* – which claimant it supports (or Neutral or N/A). We note that our very strict definition of relevance that excludes discussions of background information or historical context. These were obtained in two stages. First, we performed a small-scale human annotation effort, on 5 languages: English, Chinese, Hindi, Spanish, and Korean. We chose these 5 languages due to annotator availability, as well as their diversity in language families and scripts. Second, informed by the human effort, we designed an LLM-as-annotator approach to annotate the full dataset. Below, we provide a high-level description of annotations, and give details in Appendix D.5.

Human Annotation For the human annotation effort, we recruited from a graduate-level computer science course annotators who are bilingually fluent in English and one of the other 5 languages. For each task, an annotator is shown a query, as well as 10 retrieved documents from a given IR system and mode. For each document, annotators are asked to fill in two multiple-choice questions, on the relevance and on the viewpoint (so each task is 20 annotations). We assigned three annotators per task, and use the majority label for each annotation.

We performed several iterations, informed by participant feedback, before finalizing instructions, which are shown in Figure D.2. Annotations were requested for 620 pairs; after accounting for overlapped retrievals and quality control, 543 pairs were annotated in total.

LLM Annotation To scale up annotation to the full [BORDIRLINES](#) dataset of 19k pairs in 49 languages, we designed a two-stage LLM-as-annotator pipeline.³³ We use GPT-4o, and use exemplars drawn from the human-annotated subset for few-shot prompting. First, we prompt for relevance. Then, for query-document pairs marked as relevant, we prompted separately for viewpoint. On relevance, the LLM judged 21.5% of retrievals to be relevant, which is expected given the multilinguality and strict definition of relevance. Agreement between LLM-annotated and human-annotated labels was high, at 76% F1 agreement rate. On viewpoint, however, there was a low agreement rate. We performed a manual analysis of several documents, and found that Wikipedia’s moderated opinions (due to neutral point-of-view) made the viewpoint judgments subjective (further detailed in Appendix D.5.3). Therefore, our additional experiments study the effects of LLM relevance labels (§6.2.6), and not the LLM viewpoint labels.

6.2. Experimental Setup and Results

6.2.1. Experimental Setup

IR Systems Used We consider two multilingual IR systems. We use **OpenAI embeddings** (`text-embedding-3-large`) to embed documents and queries (OpenAI, 2024b), and use cosine similarity as the distance metric. We also use the open-source **M3-Embedding**, following the official

³³We emphasize that these annotations are a meta-evaluation of the quality of multilingual IR systems. The RAG results do not depend at all on these annotations.

implementation of the hybrid retrieval system (Chen et al., 2024a). Both systems had reasonable performance for multilingual IR. We mainly report results using OpenAI embeddings, as its documents achieved higher downstream performance than M3.

LLMs Used We consider several representative LLM families with multiple model sizes, bolding the ones studied in the main text. For closed-source, we use **GPT-4o** and **GPT-4o-mini** (OpenAI, 2023). For open-source, we use Llama 3 {1B, 3B, **8B**} (Llama Team, 2024), as well as Command-R {7B, **35B**} (Cohere For AI, 2024).

6.2.2. Evaluation Methodology

Response Formats We use 2 different prompt templates to elicit responses in two formats.³⁴ We adopt the *direct* format from Chapter 4, where the instructions are short and specify to output one of the multiple-choice selections. We also adopt the *citation* format, the instructions are more detailed, and specify a multipart output: a selection, and n natural-language explanation which cites documents IDs (e.g., <doc1>). For each setting, we perform 10 runs with 10 fixed random seeds and temperature=0.5, and report the average scores and 95% confidence intervals.

Metrics for Cross-lingual Robustness We adopt the cross-lingual robustness evaluation suite from Li et al. (2024b). This is based on the concurrence score metric (CS), which is a binary accuracy between two strings – 1 if equal, 0 otherwise. The metrics, as described in §4.3.1, are:

- **KB_CS** ↑: compares the claimants from the KB and the English response. This is our measure of *factuality*.
- **Cst_CS** ↑: compares all pairwise responses in a multilingual response set (MRS). This is our measure of cross-lingual *consistency*.
- **Δ_CS** ↓: the difference of the controller language preference and the non-controller ones. This is our measure of *geopolitical bias*.

Each CS metric is calculated over the multilingual response set (MRS) for a territory. For the 720 queries, there are 251 territories, and thus 251 MRSs. We will then compare the averaged score across all territories, for each RAG setting consisting of an IR mode and an LLM.

³⁴Refer to Appendix D.3 for the full prompts used.

Comparison to Factual Robustness Prior work studying robustness in RAG has focused on *factual robustness*, The two attributes are hallucination rate (generating non-existent information) and error rate (failing to use relevant information), construed as binary classification tasks (Adlakha et al., 2024; Chen et al., 2024b; Thakur et al., 2024).

Although related, our work evaluates differing attributes of robustness. Factual robustness considers *a single interaction*, and how a response uses provided documents. Cross-lingual robustness considers *multiple interactions*, and how responses differ given different queries and languages. Cross-lingual robustness can hence be applied to LLM’s responses with or without RAG.

Metrics for Citation Analysis We calculate two metrics to analyze responses in the citation format. Considering a prompt with 10 documents, the *inclusion rate* (nr) for a language l is the fraction of documents in that language. Considering the citation format response, The *citation rate* (cr) of l as the rate in which its documents are cited.

For a given RAG system, we characterize its citation rate cr and inclusion rate nr for language l by averaging scores over all queries l . We say there is a query-language preference when $cr \gg nr$, and an anti-preference when $cr \ll nr$.

6.2.3. Results on Cross-lingual Robustness

Figure 6.3 presents the results for each model and IR mode combination, for each CS metric. We will first focus on comparing `qlang` and `rel_langs` IR modes to the `no_ir` baseline, then consider the other IR modes.

Results for `qlang` and `rel_langs`

Factuality generally increases with RAG The top panel concerns only English queries (but multilingual documents). Compared to `no_ir` both IR modes reliably boost KB_CS scores for all LLMs; this concurs with the general view of RAG improving factuality of responses. `qlang` outperforms `rel_langs` for 3 of the 4 LLMs, though the intervals overlap in all cases (we can see that intervals for KB_CS are all narrow). We also see that the largest gains are for the least and most factual LLMs, Llama and GPT-4o.




Figure 6.3: Results for the concurrence score (CS) metrics, which measure attributes of cross-lingual robustness: Knowledge-base CS (KB_CS) for *factuality*, Consistency CS (Cst_CS) for *consistency*, and Delta CS (Δ _CS) for *geopolitical bias*. Within each subplot, we display the results for the IR modes with bars, and the no_ir baseline with the dotted line. Each score is shown as the average over 10 runs, along with its 95% confidence interval.

Consistency varies with RAG The middle panel considers multilingual queries and responses. qlang has mixed effects – Cst_CS greatly increases for Command-R (64.2 $\rightarrow$ 74.3), while decreasing for GPT-4o-mini (78.6 $\rightarrow$ 71.7) and GPT-4o (79.9 $\rightarrow$ 77.2). rel_langs in contrast, has a positive effect, especially for the least consistent model, Command-R (64.2 $\rightarrow$ 78.7). These findings underscore the importance of our study on the *selection of documents* for RAG. By retrieving over multiple languages, the set of LLM’s responses across a query posed in different languages can be, in fact, be more consistent.

Geopolitical bias reliably decreases with RAG The bottom panel also considers multilingual queries and responses, but focusing on the geopolitical bias attribute. We see that both RAG settings reduce Δ _CS and thus geopolitical bias. rel_langs is more effective than qlang, which again shows that having the multilingually-sourced perspectives is helpful in reducing LLM inconsistencies. We see that the largest decreases with Command-R (28.7 $\rightarrow$ 12.2, 5.9), and GPT-4o-mini (23.6 $\rightarrow$

71.9, 0.9; observe the neutralization of bias with `rel_langs` here). Furthermore, we see that Δ_{CS} in the `no_ir` modes have wide confidence intervals, due to the cultural-sensitivities of the task; but with RAG, they are greatly narrowed.

LLMs display different sensitivities to RAG We also observe that in terms of RAG affecting responses, Llama sees the smallest changes, and Command-R is the most affected. This aligns with the fact that Command-R was explicitly trained for RAG, while Llama was not. The greater sensitivity invites more caution into selection of documents – selecting especially biased (or factually incorrect) documents will harm Command-R’s robustness more than other models, and conversely for less biased documents.

Discussion Overall, utilizing RAG improves cross-lingual robustness for all three attributes over the `no_ir` baseline. On one hand, this is expected, given the purpose of RAG in grounding responses to the information provided in documents. Here, the documents and the original task are both sourced from Wikipedia, as with prior work.

The use of CS metrics enable quantified insights into how different attributes of robustness are influenced by the IR mode, and the LLM used. The most interesting finding is in the success of `rel_langs`. Prior works studied *synthetic* questions and documents, and raised the potential challenges RAG systems face in reconciling multilingual knowledge conflicts. But our investigation shows that, in a *real-world setting*, concerns for such conflicts may be over-stated, at least given a more reliable source as Wikipedia. Multilingual RAG overall increases cross-lingual consistency of responses, while the in-language RAG (`qlang`) setting that lowers this robustness attribute.

Results for Other IR Modes

We again note that each IR mode is motivated by different information access needs. For `en_only`, retrieval over English Wikipedia, this has the effects of 1) increasing consistency, but 2) increasing geopolitical bias for 4 of the 5 models (still less than `no_ir`). This provides further evidence that diverse multilingual retrieval is best. `qlang+en` results are similar to `en_only`, showing that IR systems’ prefer English documents.

Finally, `swap_docs` performs an interesting ablation with cross-lingual IR over a different language. For example, “Ceuta” with es query but ar docs. This reverses geopolitical bias by causing Δ_{CS} to go negative (-9.0 for Llama, -4.5 for GPT-4o), showing that relevant documents from different languages indeed contain conflicting perspectives.

6.2.4. Results on RAG Citations

Using responses from the citation format, we investigate how an LLM provided with documents uses them in its response, and whether it prefers to cite certain languages’ documents over others. We thus consider the `rel_langs` IR mode. We group the 49 languages into 3 groups: low-resource (LR), high-resource (HR), and selected languages (the highest resource languages, and the ones we had performed some human annotations on).

Results for `rel_langs` and GPT-4o are shown in Figure 6.4. Considering the box-plots, we see similar median citation rates for HRs and LR ($cr \approx 36\%$). LR has an overall in-language preference ($nr = 27.7\%$), while HRs do not ($nr = 33.2\%$). LR also has a much wider spread, from 0-100%, while HR has a smaller spread from 20-60%.

On the selected languages, the most striking finding is that `rel_langs` with an English query is barely multilingual, with $cr = 94.8\%$, $nr = 93.9\%$. For ko, there is a strong anti-preference ($cr = 28.5\%$, $nr = 54.7\%$), while for zht there is a preference ($cr = 44.9\%$, $nr = 34.1\%$).

Aggregated Citation Responses A bird’s-eye view into citations, in aggregate, is shown in the distribution pie charts in Figure 6.5. `qlang` reflects the language distribution of the 720 queries in [BORDIRLINES](#); it can be considered a denominator for the other two charts. As the other modes have multilingually-sourced documents, by construction their language proportions must be lower than for `qlang`. Considering `qlang+en` and `rel_langs`, we see that, as expected, `qlang+en` cites English documents more often ($67.1 > 57.2$). For all selected languages, `rel_langs` has a higher citation rate than `qlang+en`.

Discussion While prior studies on HR languages found LLMs prefer to use query-language and English documents, we found more nuances exist when considering 49 languages. Low-resource




Figure 6.4: Each point graphs, for queries of a language, the citation rate vs. inclusion rate in the `rel_langs` setting using GPT-4o. The dashed line indicates where the rates are proportional; above indicates a language preference, and below indicates an anti-preference.




Figure 6.5: Citation distribution across three RAG modes. Each pie chart shows the percentage breakdown of languages of citation documents.

languages are susceptible to large swings in (anti-)preference. We also note the self-fulfilling nature of high *nr* enabling high *cr*. Our study of LLM’s citation preferences emphasize the need for more inclusive retrieval techniques and generative models for underrepresented languages.




Figure 6.6: Proportion of document languages retrieved per query language, using OpenAI and M3 embeddings in the `rel_langs` mode.

6.2.5. Language Preferences of Cross-lingual IR

Figure 6.6 shows a per-language breakdown of retrieved document languages for the two IR systems, OpenAI and M3, in the `rel_langs` setting.

OpenAI embeddings prefer English documents On average, OpenAI retrieved 1.72 times more English documents than M3 (1885 for OpenAI vs. 1102 for M3). We hypothesize that this could be due to the training of M3 having a larger focus on multilinguality. However, as M3 embeddings scored worse on the CS metrics, this highlights the challenges of multilingual RAG for current LLMs.

IR systems prefer in-language documents A query language document is 1.29 times more likely to be retrieved than a non-native document for OpenAI embeddings, and 1.64 times for M3. However, we observe that this greatly varies across languages. We see that some low-resource languages such as Armenian (hy), Georgian (ky), and Lao (lo) have near 100% retrieval of native documents, while others such as Nepali (ne) and Haitian Creole (ht) primarily retrieve foreign




Figure 6.7: CS metric results for the ablation on filtering for relevance and non-relevance, compared to no IR and all documents. We use GPT-4o, and `qlang` and `rel_langs` modes.

documents. This finding, that IR systems have a preference towards in-language retrieval bias, is corroborated by Sharma et al. (2024); Wu et al. (2024).

6.2.6. Ablation: RAG with Controlled Relevance

The experiments above used all retrieved documents for RAG. However, LLMs’ responses have been shown to be susceptible to non-relevant documents (Shi et al., 2023). We therefore perform experiments in which we perform RAG with filtered sets of documents by relevance, using the annotations from §6.1.6. We consider only the GPT-4o LLM, and the `qlang` and `rel_langs` modes.

Results are shown in Figure 6.7.³⁵ On factuality, `all` with 10 documents achieves the best KB_CS over both `rel` and `nonrel`. A hypothesis is that an LLM have strong background knowledge on territorial disputes, so it will use the helpful parts of non-relevant documents that are helpful, and ignore the non-helpful parts.

³⁵These numbers are not directly comparable to Figure 6.3 because they were calculated over 446/720 queries. We kept only those queries which, after relevance filtering, had 1+ documents.




Figure 6.8: Relevance and viewpoint judgments for the Spratly Islands case study, grouped by query language and retrieval mode. Stacked bars show the document’s viewpoints.

Model	Mode	Spratly_en	Spratly_ms	Spratly_tl	Spratly_vi	Spratly_zhs	Spratly_zht
GPT-4o	no_ir	Vietnam	Vietnam	Vietnam	Vietnam	PRC	PRC
	en_only	Vietnam	Vietnam	Vietnam	Vietnam	Vietnam	Vietnam
	qlang+en	Vietnam	Vietnam	Vietnam	Vietnam	Vietnam	Vietnam
	qlang	Vietnam	Vietnam	Vietnam	Vietnam	PRC	PRC
	rel_langs	Vietnam	Vietnam	Vietnam	Vietnam	Vietnam	PRC

Table 6.2: Spratly Islands viewpoint judgment results for GPT-4o across different retrieval modes.

On consistency, we observe similar trends: **rel** is as consistent as **all**, while **nonrel** is less consistent. Geopolitical bias warrants a closer look at each IR mode. For **qlang**, we observe that **all** has the highest Δ_{CS} ($15.7 > 10.0 > 6.9$). For **rel_langs**, **nonrel** is higher by a bit over **all** ($9.8 > 8.4$). As **all** is a concatenation of the two other subsets, our initial hypothesis is that the ‘non-relevant’ documents still provide some context on the claimant nation’s perspective.

6.3. Qualitative Analysis

The real-world significance of each territorial dispute often looms large over the claimant countries’ political relations. We choose several notable ones for detailed qualitative analyses: Spratly Islands and Crimea. Studies use OpenAI embeddings for retrieval, and GPT-4o for generation.

6.3.1. Case study: Spratly Islands

The Spratly Islands located in the South China Sea are disputed by 6 countries.³⁶ As shown in the viewpoint distributions of Figure 6.8, retrieved documents are mostly neutral or non-relevant. Among opinionated documents, no single claimant dominates (other than for vi prompts).³⁷

Direct Responses The model in the control mode, `no_ir`, responds "Vietnam" for {en, ms, tl, vi} and "PRC" for {zhs, zht}. Under various RAG settings, generations for {ms, tl, en} remain unchanged, likely due to the neutral stance of the retrieved documents. Interestingly, for `qlang+en` and `en_only`, {zhs, zht} responses shift to "Vietnam", despite documents being largely neutral. Table 6.2 shows the responses across IR modes.

Citation Responses Prompting the model for direct citations causes shifts in responses for the `qlang` IR mode (results in Appendix Table D.6). Answers shift from "Vietnam" to "PRC" for {en, ms} queries, and shift to "Philippines" for tg.

We report an interesting observation for the `en_only` IR mode. One English document (Figure D.1) with a neutral viewpoint was retrieved for queries in all languages; however, its interpretation varied. Consider the respective explanations for the vi query responding "Vietnam", and the zhs query responding "PRC":³⁸

"While other countries [...] also claim parts of the Spratly Islands, the documents indicate that Vietnam's claims are well-established and recognized in its legal framework."

"The [PRC] asserts that a significant portion of the Spratly Islands is part of its territory [...] under the administration of Sansha, Hainan."

This illustrates how a model can change its answers to reflect the information from the retrieved documents, as expected. This underscores the importance of the composition of retrieval KBs, in

³⁶The *claimants* are: Vietnam, Malaysia, Republic of China (ROC), People's Republic of China (PRC), Brunei, Philippines; the corresponding *languages* are: Vietnamese (vi), Malay (ms), Traditional Chinese (zht), Simplified Chinese (zht), Malay (ms), Tagalog (tl).

³⁷Also refer the supplemental materials in Appendix §D.4.1.

³⁸Full response texts are given in Table D.7.




Figure 6.9: Relevance and viewpoint judgments for the Crimea case study, grouped by query language and IR mode. Stacked bars show the document’s viewpoints.

Model	Mode	Crimea_en	Crimea_ru	Crimea_uk
GPT-4o	no_ir	Ukraine	Ukraine	Ukraine
	en_only	Ukraine	Ukraine	Ukraine
	qlang+en	Ukraine	Ukraine	Ukraine
	qlang	Ukraine	Ukraine	Ukraine
	rel_langs	Ukraine	Ukraine	Ukraine

Table 6.3: Crimea viewpoint judgment results for GPT-4o across different retrieval modes. The table shows the generated territorial results for queries in English (Crimea_en), Russian (Crimea_ru), and Ukrainian (Crimea_uk).

determining what information can be retrieved at all.

6.3.2. Case study: Crimea

Crimea is a peninsula on the Black Sea that has been disputed between Ukraine and Russia since Russia’s annexation in 2014, an action rejected by Ukraine and much of the international community. This dispute presents a case where retrieved documents largely favor one perspective (Ukraine), making it useful for analyzing how RAG affects model alignment with dominant viewpoint texts.³⁹

Without RAG for GPT-4o-mini, responses favor Russia, but with RAG, responses consistently shift to Ukraine. This reflects how RAG encourages models to align with the viewpoint of retrieved

³⁹Also refer to the supplemental materials in Appendix §D.4.2.

documents. With document citations as well, both GPT-4o and GPT-4o-mini continue to support Ukraine, though they cite different sets of documents. Appendix Table 6.3 summarizes the territorial judgments across different IR modes for GPT-4o and GPT-4o-mini.

6.4. Conclusion

In this chapter, we introduced BORDERLINES a benchmark which extends the territorial disputes task to the RAG setting. Leverage the attributes of cross-lingual robustness, we show that while factuality in English improves in all RAG settings as expected, the effects on cross-lingual consistency warrant closer inspection. We see that monolingual retrieval, which most previous RAG works have focused on, amplifies inconsistency and geopolitical bias. In contrast, multilingual RAG can mitigate the issues of LLM’s providing different responses to a given underlying query.

Our further comprehensive analyses seek to understand how multilingual knowledge is leveraged throughout every part of the multilingual RAG pipeline. Both the upstream IR systems, and the downstream LLMs, exhibit high variability in multilingual settings for the proportion of retrieved native vs. in-language documents. We gained insight into how the language differences in available knowledge affected LM responses, and found that despite their in-progress multilinguality, multilingual RAG was able to improving robustness.

So far, we have only considered rather targeted retrieval sources: over relevant machine translation databases, and over Wikipedia. However, as much as researchers appreciate Wikipedia as an accessible, “trustworthy” resource, it has several limitations: a) the neutral point-of-view is still a perspective, meaning it cannot be a definitive source of ground-truth, b) it has limited coverage of issues and c) RAG for real-world tasks should consider information from multiple sources. These limitations motivate our next research project into creating resources for state media, and studying how LMs navigate over such real-world compendiums of knowledge.

CHAPTER 7

PERSPECTIVES FROM FOREIGN AND DOMESTIC STATE MEDIA

Chapters 5 and 6 showed that, with proper organization of external, multilingual retrieval sources, we can improve the performance and robustness of retrieval-augmented LMs. In this final project, we expand our study of retrieval sources to articles from state media outlets. We first collect a large-scale dataset of state media articles. Notably, unlike the stated neutral point-of-view of Wikipedia, state media outlets in contrast are far from neutral, as the nations behind each outlet seek to advocate their particular world view through the articles. This makes our study especially insightful in investigating how current RAG systems navigate over real-world, opinionated, and often conflicting sources of information.

A key insight is that state media targets either a *foreign* or a *domestic* audience, which is captured in separate splits of our dataset. Our main experiments consider the territorial disputes dataset, and we find that RAG over foreign media particularly amplifies robustness issues, compared to either domestic media or Wikipedia. This chapter discusses the unpublished manuscript “The World Through State Media-Tinted Lenses: Comparing Foreign and Domestic Perspectives of State Media Articles from 170+ Countries”.⁴⁰

The mainstream media, or the most accessible and influential mass media outlets, plays an outsized role in influencing public discourse (Watts et al., 2021). The media landscape of each country falls along a sliding spectrum of journalistic freedom, from most to least free. On one end of the spectrum are countries where a diverse set of perspectives are permitted and even celebrated. On the other are countries where perspectives must be vetted and approved by a central authority.

A throughline across all mainstream media landscapes is *state media*, or media outlets that are owned, operated, or heavily influenced by a government. In unfree countries, tightly-regulated state media outlets may be the only allowed ones, thus playing a crucial role in maintaining the control of

⁴⁰We plan to submit to the *Transactions of the Association for Computational Linguistics* journal after completion of this dissertation.

authoritarian governments (Walker and Orttung, 2014). In free countries, more journalistically-free state media outlets are widely accessed, and seen as reliable.

Within a given country, state media inherits much of the authority of the government (Neveu, 2004). This means that its information has a baseline-level of trust for everyday domestic citizens. The viewpoints of two countries often conflict on certain situations, and citizens reading their country’s state media may end up with different views of the world. For example, consider the Russia-Ukraine war. For the same battle, each country may seek to maximize enemy casualty counts and minimize its own. For the same territory, Russian media may emphasize its actual control and Ukrainian sources may emphasize international claims for its control. Even among skeptical citizens from a country, they may believe the truth lies closer to their government’s claims than others’.

Thus, we are interested in studying state media both within individual countries and across different countries. In our study, AI plays two alternating, important roles. First, the multilingual capabilities of LMs enable their use as tools for this multilingual analysis. Second, the increasing role of LLM-based tools by end-users across the world inspires us to understand how LLMs can serve to amplify certain viewpoints and minimize others – and in which contexts they may do so differently.

In this work, we study how retrieval-augmented LMs navigate the perspectives presented by state media from diverse countries, and sets of countries. We first create resources – a database of state media websites, and a large-scale collection of articles. We analyze the prevalence of state media in English and multilingual web crawl data. Then we present findings on RAG experiments over the BORDERLINES territorial disputes benchmark. Notably, we find that state media which targets foreign audiences amplifies RALMs’ divergent responses over those targeting domestic audiences. Also, retrieval over diverse sources provides more robust responses than over single sources.

7.1. Resources for State Media

We develop two resources: a table of state media websites, and a corpus of articles.

7.1.1. Dataset Collection

Country	Media company	Main assets	Typology	Sub-region	URL_dict
Argentina	Radio y Televisión Argentina (RTA)	Television: Televisión Pública (Canal 7), Canal 12; Radio: Radio Nacional	SC	South America	{Televisión Pública (Canal 7) : https://tvpublica.com.ar , Radio Nacional : https://radionacional.com.ar }
China	China Media Group	News agency: Xinhua; Publishing: Reference News (Cānkǎo Xiāoxī); Television: CNC World	SC	East Asia	{Xinhua : http://xinhuanet.com/english , Reference News (Cānkǎo Xiāoxī) : http://cankaoxiaoxi.com }
France	France Médias Monde	Radio: Radio France Internationale (RFI), Monte Carlo Doualiya (MCD) Television: France24	ISFM	Western and Northern Europe	{Radio France Internationale (RFI) : https://rfi.fr/fr , Monte Carlo Doualiya (MCD) : https://mc-doualiya.com , France24 : https://france24.com/fr }
Kenya	Kenya News Agency (KNA)	News agency: KNA	SC	Eastern Africa	{KNA : https://kenyanews.go.ke }
Qatar	Al Jazeera Media Network (AJMN)	Television: Al Jazeera Arabic, Al Jazeera English, ...	SC	Middle East	{Al Jazeera Arabic : https://aljazeera.net , Al Jazeera English : https://aljazeera.com ", ...}
Russia	Public Television of Russia (OTR)	Television: OTR	SC	Russia	{OTR : https://otr-online.ru }
United States of America	U.S. Agency for Global Media (USAGM)	Radio: Voice of America, Radio Free Europe/Radio Liberty, Radio Free Asia, Office for Cuba Broadcasting, ...	ISFM	North America	{Voice of America : https://voanews.com , Radio Free Europe / Radio Liberty : https://rferl.org , Radio Free Asia : https://rfa.org/english , Office for Cuba Broadcasting : https://martinoticias.com ", ...}

Table 7.1: Excerpted entries from the state media articles table.

State Media Websites State Media Monitor (SMM) is an online collection of state media (<https://statemediamonitor.com>). The researchers have compiled state media companies from across the world, and identified the assets each company holds (Dragomir and Söderström, 2021). There is also a typology of the level of state control, from least to most:

- Independent Public Media (**IP**)
- Independent State Media (**IS**)
- Captured Private or Public Media (**Ca**)
- State Controlled Media (**SC**)

Where the colors broadly classify each category into **independent** and **controlled**. Considering the row on “USAGM”, it is from the United States of America, has several main assets, and is classified as **IS** (ISFM is a subcategorization of **IS**, which we have simplified here).

Our resource extends the SMM table with URLs for each main asset. We do this through a two-pass procedure. First, following the intuition that LLMs are trained on the web, and therefore have knowledge of common news sources, we prompt an LM to generate URLs for each asset, given a row from SMM (details in Appendix E.1). From this initial list, we manually verify each URL by loading it in a web browser; we keep those URLs which load, and for those that do not, we perform a web search to attempt to find one. The statistics for the resulting dataset are: 603 media companies, 1910 media assets, 1183 URLs (since some assets may share URLs), 173 countries.⁴¹ Excerpted rows from the resulting state media articles dataset are shown in Table 7.1.

State Media Articles We next collect a corpus of state media articles, indexed to each country. In lieu of performing our own scraping over the web, we turn to large-scale web crawls. The Common Crawl initiative creates archives of the entire internet and freely releases it at regular intervals. In our work, we use the filtered and deduplicated versions of CommonCrawl text dumps released by the Fineweb initiative. The English Fineweb dataset consists of crawls from 2013 to 2024, with over 15T tokens of English text (Penedo et al., 2024). The Fineweb-2 dataset consists of crawls in the same time period, and extends to over 1000 languages, with over 3T words (Penedo et al., 2025).

We separately process either dataset. Fineweb-2 is organized into crawls by language. We will similarly iterate over rows and extract rows with URLs from state media sources. For an individual language dump, we consider only the state media URLs from country’s with it as a main language (from the BORDERLINES table). This is done both for computational efficiency over the massive datasets, and for focusing our study to *domestic* articles intended for consumption by a country’s own citizens.

English Fineweb is organized into crawls by date. For each crawl, we iterate over the rows, and extract the rows with URLs that are found in the 1183 state media URLs from our dataset. We

⁴¹Publicly released at huggingface.co/datasets/manestay/state_run_media




Figure 7.1: # webpages for top 32/100 countries in English Fineweb.

consider most articles from this dataset as *foreign* articles, as they are written in English to target an international audience.⁴²

7.2. Dataset Analysis

7.2.1. Webpages per country, English Fineweb

Of the 173 countries in the website dataset, 100 countries appear in English Fineweb, totaling 69.7m docs. Figure 7.1 depicts the number of webpages for the top 32 countries. The bars are split into colors by typology (we collapse the 7 original categories into 4). We observe that there are 15.7m pages from United States state media, with all sources classified as independent. Second highest

⁴²From English Fineweb, articles from countries with English as the main language are instead considered as *domestic* articles.




Figure 7.2: # webpages for top 32/85 countries in Fineweb-2.

is Russia, with 8.3m pages, all of which are state-controlled. Qatar, which publishes Al Jazeera, has 3.6m pages. Of course, well-represented are Anglophone countries as the United Kingdom, Singapore, and Australia. Going down the long tail of countries, we see populous nations such as China and Vietnam, to smaller nations such as Brunei and Slovenia. This analysis lends credence to the characterization of English state media from non-Anglophone countries as intended for foreign consumption.

7.2.2. Webpages per country, Fineweb-2

Of the 173 countries in the website dataset, 85 countries appear in Fineweb-2, totaling 18.3m documents. As before, the number of webpages per country are depicted in Figure 7.2 (Anglophone countries are excluded from this chart). For this dataset, under consideration are state media

intended for domestic consumption. As expected, most of the top countries are populous ones: for authoritarian nations, we have Russian, Vietnam, Iran, and China; for democratic ones, we have Turkey,⁴³ Spain, and France.

We flag several standout countries, and perform some preliminary analysis. India has only 27k pages for a population of 1.5 billion. Indian state media companies primarily serve radio and television, rather than online. Furthermore, the major media companies in India are private, rather than state-run. Serbia has 225k pages for a population of 6.6m. Until 2000, the Serbian government had a monopoly over domestic media, and state media still holds a dominant position to the present day.

7.2.3. Webpages over time, English Fineweb

We analyze the proportion of state media written in English over time in Appendix Figure E.2. In the years 2013-2016, 0.3-0.5% of all Fineweb articles are from state media sources. In Sept. 2016 (2016-40), there is a spike towards 0.64%; while more in-depth analysis is needed, this could be due to the 2016 US presidential election. After 2020, the proportions drop to around 0.20%; we suspect this is due to a large increase in the web content, which could either be due to actual increases in content, or a change in the coverage of the Common Crawl systems. Additionally, we observe that from 2013-2018, the majority of content comes from independent state media sources. 2019 serves as an inflection point, and the following years now have controlled media sources as majority.

7.2.4. Considerations

We reiterate that these are preliminary findings, which we will to refine and develop in the future. We close this section on several points.

- The absolute proportion of state media in all the Fineweb LLM training data is low, at 0.2-0.6%. However, this could still have an outsized influence.
- The agenda and level of control of various state media companies varies greatly. For example, **IP** such as NPR from the USA, or BBC from the UK, are as reputable as private media.

⁴³While Turkey is a parliamentary republic, its recent democratic backsliding has lead to some characterizations as authoritarian (Tansel, 2018).

Meanwhile, **SC** such as Al Jazeera from Qatar can be quite reputable for most international issues, but less so for domestic ones.

- State media is only a fraction of the entire media landscape for every country. In nations such as the US, the mainstream media is largely dominated by private media. The biases and societal effects of such outlets is well studied in the literature (Groseclose and Milyo, 2005; Chang and Tang, 2023; Wang et al., 2025a).

7.3. Experimental Setup and Results

In Chapter 6, we studied how cross-lingual robustness of RALMs was affected by the language composition of retrieval databases, using a single source, Wikipedia. We now expand to varying retrieval over of multiple sources, in addition to multiple languages.

7.3.1. Territorial Disputes

We first study how retrieval over state media for the BORDERLINES dataset of territorial disputes from Chapter 4.

Experimental Setup

We adopt the setup as described in §6.2.1. For IR, we use OpenAI embeddings to embed documents and queries, with cosine similarity as the distance metric. For LLMs, we use GPT-4o, Command-R 35B, and Llama 3 8B. We prompt for the direct output format of a multiple-choice selection. We utilize the concurrence score metrics to evaluate cross-lingual robustness.

Retrieval Settings

We vary two dimensions of retrieval databases. *Sources* are either from foreign or domestic state media. *Perspectives* (or IR mode) are either at **query-level**, over only the state media database for a specific country, or at *topic-level*, over all databases for all involved countries. This results in four retrieval settings, as shown in Figure 7.3.

Results

Results are shown in Figure 7.4, and lead to the findings:




Figure 7.3: Illustration of the 6 retrieval settings studied in this work, which vary the sources (top columns), and the perspectives (left rows). For each box, the fill color indicates the language, while the outline color indicates the source. The two Wikipedia settings from Li et al. (2025a) are single source, while the other settings from work are multi-source.

LMs responses are far less robust with multilingual queries Recall that KB_CS is evaluated over the English queries only, while Cst_CS and Δ _CS are evaluated over all queries. The variations in KB_CS are moderate for all settings and models, while Cst_CS and Δ _CS see massive swings.

Factuality decreases with state media RAG While retrieval over Wikipedia consistently improved KB_CS over no_ir, retrieval over state media decreases it. Domestic sources (multilingual) and topic-level perspectives lead to the largest drop. These results show the perils of prior work in both using Wikipedia as the ground-truth, and the only source for retrieval. When we expand RAG to real-world articles, the diverse perspectives introduced can cause factuality to drop.

Foreign state media perspectives causes greatest divergence in RALM responses Δ _CS for RAG over foreign state media and query-level is greatly amplified – over 100% in some cases. This means RALMs overwhelmingly prefer to answer with a non-controller for non-controller queries,




Figure 7.4: Concurrence score (CS) metrics with retrieval over *sources*: Wikipedia (from Li et al. (2025a)), foreign, or domestic state media; and over *perspectives*: topic-level vs. query-level.

and similarly for controller queries. Cst_CS also deteriorates for this setting.

Closed LMs are more robust to variations from RAG than Open LMs While the trends above hold for all models, the effects on GPT-4o and GPT-4o-mini are smaller than with Llama or Command-R. This indicates that closed LMs more closely adhere to their parametric knowledge when faced with differing contextual knowledge across languages. We hypothesize this is due to additional post-training and alignment effort by those developers.

Case Study: Foreign State Media on Crimea

We return to the disputed territory of Crimea as a case study. It is especially applicable to our database, which includes state media articles from 2013-2025, covering the initial annexation of Crimea by Russia in 2014 to the present-day and the ongoing Russia-Ukraine war.

Figure 7.5 shows the retrieved documents over foreign sources with topic-level perspectives, and LM’s responses in two languages. While we performed actual retrieval and generation, for brevity this shows only 2 articles, and copy-edits the LM responses. On the Ukrainian side, the retrieved articles reinforce Ukraine’s claim, while on the Russian side, the articles reinforce Russia claim. This leads to the LMs differing responses. This study also illustrates a small-scale example of a larger




Figure 7.5: Case study on Crimea for RAG over foreign state media. We retrieve articles from sources which are written for an English-speaking audience. This example has *topic-level* retrieval: for the Ukrainian query, we use articles from Ukraine’s state media, and vice versa for the Russian query and articles from Russia’s state media.

issue: that state media, especially those targeting foreign audiences, portray perspectives favorable to a state’s own goals.

7.4. Conclusion

In this chapter, we study RAG where knowledge is drawn from state media. Each country’s state media seeks to inform, but also influence the beliefs, of domestic and foreign audiences. To investigate these issues, I introduced two resources, a table of state media outlets, and a corpus of articles split into a multilingual subset for domestic news, and an English-only subset for foreign news. The dataset analysis provides insight into how different countries invest different amount of effort into, and exert different levels of control over state media. Our main experiments on the territorial disputes benchmark show that foreign state media perspectives causes the greatest divergence in RALMs. We hypothesize this could be due to foreign media being more propagandistic than domestic media.

With this final project, we have expanded from a typical RAG setting, where retrieval databases

are trustworthy and domain-specific, to one of particular challenge for AI systems, where they must retrieve over and reconcile information across multiple languages, multiple sources, and the perspectives contained within.

CHAPTER 8

CONCLUSION

This thesis has focused on building evaluation datasets to challenge the multilingual knowledge of LMs, and designing methods to improve the robustness of performance across languages. In the first half, we study the parametric knowledge of LMs, and evaluate how well they can transfer knowledge across languages. In the second half, we study how LMs can navigate contextual knowledge introduced in multiple languages and multiple sources.

In Chapter 3, I introduced xSTREET, a benchmark for multilingual complex reasoning. I showed that directly prompted LMs struggle to transfer reasoning knowledge learned in English to other languages. To remedy this, I showed that introducing the structured representation of program code can help bridge the gaps with both inference-time and training-time methods. In Chapter 4, I studied the territorial disputes task, which exists at the fault lines of language, political, and cultural concerns, and introduced the BORDERLINES dataset. I proposed the notion of cross-lingual robustness, which captures the three attributes of factuality, consistency, and language fairness. The findings show that LLMs struggle to balance the attributes, and that interventions on prompts can amplify or mitigate robustness of LM responses.

In Chapter 5, I study how knowledge is best represented for the task of MT on specialized domains. I found that paired MT demonstrations are particularly effective for LLMs, and that their key role is in their domain-specific and stylistic guidance. In Chapter 6, I extended the territorial disputes task to the RAG setting. I find that the linguistic composition of retrieval databases greatly affects responses, and that multilingual RAG elicits best the most robust responses over in-language only RAG. Finally, in Chapter 7, I study the information produced from state media outlets, which can be viewed as both authoritative and propagandistic, depending on one’s country and perspectives. I collect a dataset of state media articles for 100 countries in 49 languages. Experiments on BORDERLINES show that RAG over articles produced for foreign consumption cause more divergent responses than articles produced for domestic consumption.

The primary case for improving multilinguality of LLMs has been to improve their accessibility for diverse users around the world. In this thesis, I resonate with this argument, and present another one: *improving multilingual capabilities of LLMs leads to stronger all-around LLMs*. This is because learning concepts across surface-form language representations leads to better generalization and more robust responses. In other words, I advocate for the establishment of multilinguality as a core competency of frontier LLMs. And to effectively track progress, we must evaluate how multilingual LLMs perform on the full scope of knowledge across languages. LLM developers can do so by testing their LMs on a whole suite of multilingual benchmarks – not only the benchmarks I have introduced, but also the aforementioned translation-based, cultural, and factual benchmarks.

8.1. Future Work

As we saw with the literature review in 2, the challenges that lie ahead towards our quest in developing stronger and more multilingual LLMs are extensive. This thesis has paved the path towards knowledge-driven multilingual evaluations of LLMs. For future work, I plan to continue along the following research directions:

Tracing issues back to the pretraining data In Chapter 4, I evaluated the parametric knowledge of LMs throughout how they responded to queries. However, it is worth investigating the origin of this undesirable and cross-lingually non-robust behavior. We can do this by auditing the pretraining data, adapting from techniques used for related audits for the issues of political bias and social bias Feng et al. (2023); Thaler et al. (2024). The state media article corpora we have collected in Chapter 7 is a subset of Common Crawl, which all LLMs have been trained on some form. Thus, we can perform some corpus-level analysis, and investigate the correlations between pretraining data and model responses. For example, suppose an LM responds with ‘Russia’ for a Russian query on Crimea, and ‘Ukraine’ for a Ukrainian query. It could be the case that most Russian-language documents advocate for Russia’s claim, and vice versa. We could glean further insights by training small-scale LLMs on perturbed versions of the state media dataset: one version with the natural distribution, another with perspectives explicitly balanced for each territory, and perhaps still another with machine-translating articles into various languages.

Identifying where in LM’s parameters cross-lingual inconsistencies arise Aside from the data that goes into LMs, and the output they produce, we should also be interested in *how* LMs are storing knowledge across languages. The subfield of *mechanistic interpretability* seeks to reverse-engineer the inner workings of LMs by looking at their learned parameters, and uncovering human-interpretable behaviors. Several insights are that that factual knowledge is recalled from the feed-forward layers (FFN) Geva et al. (2021), while contextual memory is most utilized by attention layers (Attn) Sun et al. (2024). Zhao et al. (2024) perform experiments which suggest this *multilingual workflow* for an LM: early layers convert multilingual inputs into English, intermediate layers use multilingual knowledge from FFN and reasoning from Attn, and final layers convert responses back to the query language. Wang et al. (2025b) characterize the intermediate layers instead as language-agnostic. Importantly, they find that factual cross-lingual inconsistencies largely arise from the conversion in the final layers, and propose a method to bypass the final layers and predict directly from intermediate ones.

These works consider the multilingual workflow of LMs on factual queries. However, as we discussed, not all knowledge can be directly translated, and furthermore the process of translation itself modifies the meaning of a statement. Thus, we seek to apply these mechanistic interpretability techniques to two other classes of problems, to see if translation-between-layers interpretation still holds. Of course, we have territorial dispute queries, wherein responses should also be language-invariant, although having a right or wrong answer is far more elusive. Furthermore, we have the various cultural benchmarks, wherein responses should be language-variant.

Further involvement of real-world users Throughout this thesis, we emphasized that translation is not enough, and instead we need to consider real-world aspects of multilinguality. In this light, I studied the real-world task of territorial disputes, and drew retrieval databases from real-world datasets. However, the involvement of real-world users was limited. While multilingually-fluent humans assisted in several of these projects, these participants were largely based in Western countries, and were either university students or professional translators. I hope to engage my future work with a more *community-based* approach, involving participants of diverse language and

cultural backgrounds. We are particularly inspired by the Aya initiative, which has introduced multilingual LLMs and resources (Üstün et al., 2024). They not only collect annotations from diverse participants, but elicit ongoing feedback from participants as users themselves, keeping them in the loop towards building better resources.

I outline two use-cases of real-world engagement to augment the existing datasets. First, we used Wikipedia as the sole ground-truth reference for territorial disputes. But to what degree does its opinions concur or contradict with those humans who are actually involved? What about an actual UN peacekeeper’s opinions? This can inspire a revised version of BORDERLINES and evaluations of factuality with respect to different ground-truth KBs. Second, our state media dataset has labels for the level of state control of each outlet, collected by researchers at Media and Journalism Research Center in Vienna, Austria (Dragomir and Söderström, 2021). We are further interested in the levels of trust assigned to each source, along the same foreign vs. domestic axis. Chang and Tang (2023) found that in selected Asian countries, citizens in authoritarian nations had a higher level of trust in the media than citizens in democratic nations. We hope to uncover the relationships between state media and its audience for countries around the world. While this is a tall order, we plan to start by investigating news sources that are most relevant to American audiences. We can further solicit right and left-leaning participants, to identify the differences between level of trust for sources such as NPR, RT, Al Jazeera, and so on.

Investigating RAG for Critical Decision-Making Settings The proliferation of generative AI has led to what some have termed an “arms race” between different nations and their governments. News articles set the scene between the two main rivals, the United States and China (Brands, 2025; Graylin and Triolo, 2025). The US has the lead, with the strongest frontier LLMs, the most compute, and a plethora of researchers. However, China has been catching up in stunning fashion with its release of open-weight LLMs; notably, Chinese researchers represent the majority of first-authors in recent AI conferences (such as ACL 2025). The European Union has positioned itself as a third player, with a unified initiative promising fully-open models that support language and cultural diversity (OpenEuroLLM, 2025).

From the US government side, a June 16, 2025 announcement has awarded OpenAI a \$200m contract to “develop prototype frontier AI capabilities to address critical national security challenges in both warfighting and enterprise domains” (Department of Defense, 2025). It seems that AI will soon come to play a vital role in critical decision-making for national governments.

Reddy et al. (2023) propose a framework to assist human analysts in situation report generation for current events. The system retrieves relevant documents from a selection of reliable news sources, and is able to generate reports that are timely, multi-source, and trustworthy, as rated by experts on three scenarios. This illustrates the great potential of AI to empower decision-makers in processing large amounts of information, across many sources, and across many languages beyond what an individual decision-maker speaks. However, this use-case also raises the risks arising from *composition of retrieval databases* studied in Chapter 6 and Chapter 7. By performing RAG over different slices of state media outlets – e.g., by foreign vs. domestic audience, or by national alliances (NATO, CIS, BRICS) – the situation reports will differ, as they are grounded in different perspectives on the world. This comparative setting can even benefit decision-makers in creating reports from multiple viewpoints, from news sources of their own, of their allies, and of their adversaries.

APPENDIX A

APPENDIX FOR CHAPTER 3: COMPLEX REASONING

A.1. Annotation Details

The annotators are in-country, native speakers for each of the languages – Arabic, Chinese, Japanese, Russian, and Spanish. They all are certified translators with more than 5 years of full-time translation experience, according to the 17100 Translation ISO Standard. They were hired through vendors and were remunerated above industry standard rates. Instructions were to post-edit machine translated output and included guidelines on what to localize (artist names, city names, metric conversions), format (capitalization, punctuation) and structure (sentence level breaks). The vendor project managers made sure the instructions were adhered to. The QA process consisted of content review based on the Multidimensional Quality Metric (MQM) model that allocates different weights to 5 error severities (0-none to 5-critical) in several error topics. Total sample reviewed was 5 (5k words) of the total (100k words) source word count.

A.1.1. BLOOM Results for GSM8K, AQUA_RAT, AR_LSAT

These results are shown in Appendix Figure A.2. For all tasks, performance is around random chance. For GSM8K, random chance is 0, and the models fail to solve nearly any math problem. While all numbers are close and likely statistically insignificant, we see that BLOOMZ-TCC slightly underperforms base BLOOMZ, and linearized and code prompts perform similarly.

Our hypothesis on why this happens, as discussed before, builds on the view that truly complex reasoning capabilities are emergent with LLM’s model scale. The 7B BLOOMZ model used has no baseline ability for these 3 tasks (while it did for ARC), and therefore our interventions, which are indirect supervision on code, cannot help elicit better reasoning.

This suggests limitations to the code+reasoning hypothesis, which have not been adequately discussed in prior work. Indirectly supervising LLMs for reasoning by training on code is effective for specific types of reasoning, such as ARC’s commonsense reasoning, and less so for math problems

like GSM8K, though, intuitively, code probably does not help, given code rarely includes arithmetic equations.⁴⁴ As for the logical reasoning problems of AR_LSAT, we defer study to future work applying our finetuning recipe to larger LLMs.

A.2. Hyperparameters

For the TCC finetuning recipe, we used a g5.48xlarge instance from AWS, which has 8 NVIDIA A10G GPUs (24*8 GB=192GB vRAM). We set maximum sequence length to 1024, set learning rate to 1e-5 (with 0.1 weight decay), do not use warm-up, and use cosine learning rate schedule. We trained for 2 epochs using a batch size of 3 and gradient accumulation over 20 steps. We set the LoRA layers dimension to 128. The implementation is done with the DeepSpeed-Chat framework (Yao et al., 2023) and the `transformers` library (Wolf et al., 2020).

A.3. Full Results

We now report results of all experiments and settings studied in this work. Table A.1 gives results for all models and prompt formats, on STREET and xSTREET (averaged across 5 languages). Table A.2 gives per-language xSTREET results. BLOOMZ refers to the 7B model, while Falcon refers to falcon-40b-instruct. For reference only, we include the reported results of Ribeiro et al. (2022) for STREET in the row GPT-3 repro*.

A.3.1. Direct Prompting

These tables include the direct prompting strategy, in which the model is given the same input as linearized, but needs to generate only the answer without intermediate reasoning.

For ARC, we find that, surprisingly, direct outperforms linearized (all LLMs). This is likely because the ARC questions are relatively easy already, and directly solving them given the context as well is possible. As discussed earlier, linearized format is artificial and hard to follow, which causes many reasoning graphs without valid answers. In contrast, SIM prompts are captured in code, which the models have seen, and therefore SIM results outperform direct and linearized.

For GSM8k and AQUA_RAT (using GPT-3), direct prompting fails. Using intermediate reasoning,

⁴⁴Arithmetic reasoning can be improved by *directly* finetuning on math – Ni et al. (2022) achieve 19.5% on GSM8k on a 2.7B LLM with this approach.

		STREET (English)				xSTREET (5 language average)			
Model		ARC	GSM8k	AQUA_RAT	AR_LSAT	ARC	GSM8k	AQUA_RAT	AR_LSAT
direct	BLOOMZ	54.7	2.6	24.8	20.0	41.7	1.7	21.1	20.8
	BLOOMZ-TCC (ours)	70.3	3.3	20.9	24.0	49.0	1.8	21.9	23.2
	Falcon	34.7	8.5	24.8	26.0	32.0	6.1	24.6	24.0
	GPT-3	97.6	4.4	21.7	22.0	90.7	1.6	25.6	24.4
linearized	BLOOMZ	35.3	1.1	24.0	22.0	33.5	0.9	24.4	24.0
	BLOOMZ-TCC (ours)	40.9	1.5	18.5	20.0	36.9	1.0	22.5	28.5
	Falcon	76.5	28.1	18.9	24.0	54.6	10.6	23.4	24.4
	GPT-3	93.2	28.9	25.6	18.0	72.7	24.4	27.7	23.2
	GPT-3 repro*	72.9	34.8	40.2	19.0	-	-	-	-
SIM code	BLOOMZ	61.5	0.7	23.6	26.0	52.6	1.5	23.3	24.0
	BLOOMZ-TCC (ours)	76.2	1.9	21.7	28.0	61.1	1.3	25.4	24.8
	Falcon	81.5	19.6	24.4	34.0	43.9	4.6	23.5	26.8
	GPT-3	99.1	45.2	37.8	32.0	94.2	26.9	32.8	24.8

Table A.1: Full results (STREET and xSTREET average) for all models, tasks, prompt formats, and languages. Rows are grouped by prompt type and model, while columns are grouped by the subtask and the language.

Model		ARC					GSM8K					AQUA_RAT					AR_LSAT				
		es	zh	ar	ja	ru	es	zh	ar	ja	ru	es	zh	ar	ja	ru	es	zh	ar	ja	ru
direct	BLOOMZ	46.5	41.2	57.4	33.5	30.0	2.2	1.9	1.1	1.1	2.2	19.7	24.8	21.3	24.0	15.7	22.0	24.0	16.0	24.0	18.0
	BLOOMZ-TCC	63.2	44.7	68.5	31.2	37.6	1.9	1.5	2.2	1.5	1.9	22.0	25.6	19.7	22.0	20.1	18.0	28.0	26.0	24.0	20.0
	Falcon	34.7	32.4	30.3	31.8	30.9	10.0	6.7	3.3	4.8	5.6	24.8	24.0	25.2	24.8	24.0	26.0	22.0	24.0	26.0	22.0
	GPT-3	97.4	95.0	81.8	90.6	88.5	3.0	1.5	0.0	1.5	1.9	25.2	24.8	25.6	26.0	26.4	24.0	24.0	24.0	24.0	26.0
linearized	BLOOMZ	35.0	37.1	35.3	29.1	30.9	2.2	1.1	0.4	0.4	0.4	22.0	26.8	23.6	24.0	25.6	24.0	26.0	24.0	20.0	26.0
	BLOOMZ-TCC	44.7	40.3	39.7	30.9	29.1	1.5	0.4	1.5	0.4	1.1	17.3	21.7	21.3	22.8	24.0	14.0	30.0	28.0	28.0	28.0
	Falcon	77.4	65.6	34.7	49.4	45.9	23.3	16.3	2.2	7.0	4.4	25.2	25.2	21.7	22.0	22.8	18.0	26.0	24.0	18.0	36.0
	GPT-3	83.5	75.6	57.4	69.7	77.4	26.3	29.6	19.3	17.0	30.0	27.6	27.6	29.5	26.4	27.6	26.0	28.0	20.0	26.0	16.0
SIM code	BLOOMZ	57.9	67.6	58.8	45.6	32.9	0.7	1.5	2.2	2.2	0.7	21.3	25.6	23.6	22.4	23.6	26.0	30.0	30.0	12.0	22.0
	BLOOMZ-TCC	70.3	74.4	65.0	54.7	40.9	0.4	1.5	1.5	2.2	0.7	27.2	24.0	28.3	22.8	24.8	26.0	34.0	16.0	30.0	18.0
	Falcon	57.1	55.3	30.3	44.1	32.6	8.1	8.1	1.9	2.6	2.2	26.0	24.4	23.2	20.5	23.2	32.0	26.0	26.0	24.0	26.0
	GPT-3	96.5	96.5	90.3	94.1	93.5	37.4	28.1	18.5	21.9	28.5	36.6	32.7	31.1	32.7	31.1	26.0	28.0	26.0	18.0	26.0

Table A.2: Per-language xSTREET results for all models, tasks, prompt formats, and languages. Rows are grouped by prompt type and model, while columns are grouped by the subtask and the language.

as found by many prior works, is essential for arithmetic problem-solving. Linearized boosts performance significantly, and SIM code even further. Regarding xSTREET, overall trends are relatively consistent with those discussed above.

Figure A.1: Case study of one GSM8K problem in English. The top row is for SiM prompts, while the bottom is for linearized.

Input in SiM Format <pre>facts = ['Raymond and Samantha are cousins.', 'Raymond was born 6 years before Samantha.', 'Raymond had a son at the age of 23.', 'If Samantha is now 31,'] question = "how many years ago was Raymond's son born?" answer = None # call select_facts and infer_new_fact until answer is set selected = select_facts(facts) #</pre>	Gold <pre>[facts[1], facts[2]] new_fact = infer_new_fact(selected) # "When Raymond's son was born Samantha was 23 - 6 = 17 years old." is_solved(new_fact, question) # False facts.append(new_fact) selected = select_facts(facts) # [facts[3], facts[4]] new_fact = infer_new_fact(selected) # "Thus it has been 31 - 17 = 14 years since Raymond's son was born." is_solved(new_fact, question) # False facts.append(new_fact) selected = select_facts(facts) # [facts[3], facts[5]] new_fact = infer_new_fact(selected) # 'The answer is 14' is_solved(new_fact, question) # True answer = extract_answer(new_fact) # '14'</pre>	GPT-3 <pre>[facts[1], facts[2]] new_fact = infer_new_fact(selected) # 'Raymond's son was born 23 - 6 = 17 years ago.' is_solved(new_fact, question) # False facts.append(new_fact) selected = select_facts(facts) # [facts[3], facts[4]] new_fact = infer_new_fact(selected) # 'Samantha is 31, so Raymond's son was born 31 - 17 = 14 years ago.' is_solved(new_fact, question) # False facts.append(new_fact) selected = select_facts(facts) # [facts[3], facts[5]] new_fact = infer_new_fact(selected) # 'The answer is 14' is_solved(new_fact, question) # True answer = extract_answer(new_fact) # '14'</pre>	BLOOMZ-TCC <pre>[facts[0]] new_fact = infer_new_fact(selected) # facts[2] is_solved(new_fact, question) # False facts.append(new_fact) selected = select_facts(facts) # [facts[0], facts[2]] new_fact = infer_new_fact(selected) # facts[2] is_solved(new_fact, question) # True facts.append(new_fact) selected = select_facts(facts) # [question, facts[2]] new_fact = infer_new_fact(selected) # facts[2] is_solved(new_fact, question) # True answer = extract_answer(new_fact) # '23'</pre>	BLOOMZ-TCC <pre>sent1 & sent2 -> int1: Raymond was born 6 years before Samantha, so Raymond was born 6 years before Samantha.; int1 & sent2 & sent3 -> int2: Raymond had a son at the age of 23, so Raymond was 23 years old when he had a son.; int1 & int2 & sent3 -> int3: Samantha is now 31, so Samantha was 31 years old when she had a son.; int1 & int3 & sent4 -> int5: Thus, Raymond had a son 6 years ago.; int2 & int4 & sent5 -> int6: The answer is 6 years..;</pre>
Input in linearized format <pre>QUESTION = sent1: Raymond and Samantha are cousins.; sent2: Raymond was born 6 years before Samantha.; sent3: Raymond had a son at the age of 23.; sent4: If Samantha is now 31, sent5: how many years ago was Raymond's son born? PROOF =</pre>	Gold <pre>sent2 & sent3 -> int1: When Raymond's son was born Samantha was 23 - 6 = 17 years old.; int1 & sent4 & sent5 -> int2: Thus it has been 31 - 17 = 14 years since Raymond's son was born.; int2 & sent5 -> int3: The answer is 14;</pre>	GPT-3 <pre>sent2 & sent4 -> int1: Raymond was born 6 years before Samantha, so Samantha is now 31 - 6 = 25 years old.; int1 & sent3 -> int2: Raymond had a son at the age of 23, so Raymond's son was born 23 - 6 = 17 years ago.; int2 & sent5 -> int3: The answer is 17;</pre>	BLOOMZ <pre>sent1 & sent2 & sent3 -> int1: Raymond was born 6 years before Samantha, so Raymond was born 6 years before Samantha's son was born.; int1 & int2 & sent3 & sent4 -> int2: Samantha is now 31, so Samantha's son was born 6 years ago.; int1 & int2 & int3 & sent4 -> int5: The answer is 6 years..;</pre>	BLOOMZ-TCC <pre>sent1 & sent2 -> int1: Raymond was born 6 years before Samantha, so Raymond was born 6 years before Samantha.; int1 & sent2 & sent3 -> int2: Raymond had a son at the age of 23, so Raymond was 23 years old when he had a son.; int1 & int2 & sent3 -> int3: Samantha is now 31, so Samantha was 31 years old when she had a son.; int1 & int3 & sent4 -> int5: Thus, Raymond had a son 6 years ago.; int2 & int4 & sent5 -> int6: The answer is 6 years..;</pre>




Figure A.2: Results on GSM8K, AQUA_RAT, AR_LSAT subtasks for BLOOMZ-based models, with (up-to) 5-shot prompts. Random baselines are indicated with dashed grey lines. For each task, we report results 1) for, and 2) averaged over the 5 languages.

APPENDIX B

APPENDIX FOR CHAPTER 4: BORDERLINES

B.1. Temporal Versions of the BORDERLINES table

A major characteristic of territorial disputes is that they can and do change over time. Given that different LLMs have different cutoff dates, we focused our study on a static version of the BORDERLINES table, collected from the 2023-05-15 version of the territorial disputes Wikipedia page. This raises the question, “how does different temporal versions of BORDERLINES affect territories covered, and the resulting evaluation?”

In this section, we also consider an alternate version based on the 2021-08-31 Wikipedia page – which is prior to the GPT-3 and GPT-4 training cutoff date of 2021-09 – terming this BORDERLINES₂₀₂₁₋₀₉. BORDERLINES₂₀₂₁₋₀₉ is missing 70 territories present in BORDERLINES (19 are populated, 51 are uninhabited). While this is a significant number, we manually analyzed this and found almost all the territorial disputes far pre-date 2023 or 2021; *the earlier page was very incomplete* and editors have added more over time. This means it is likely that GPT models would have seen some information about these disputed territories from other sources on the Internet during training. The major geopolitical event that created new disputed territories of international dispute is the Russian invasion of Ukraine, which began on 2022-02-24: adding {Donetsk Oblast, Zaporizhzhia Oblast, Mykolaiv Oblast, Kherson Oblast, and Luhansk Oblast}.

On the other hand, BORDERLINES is missing 46 territories present in BORDERLINES₂₀₂₁₋₀₉. Our manual inspection reveals that the main difference can be attributed to BORDERLINES₂₀₂₁₋₀₉ lacking a “Disputes over territorial waters” section, instead having all of these entries in the other sections. Recall from §4.1.1 that we specifically excluded non-land disputes from our main table. Of the land disputes remaining, we observe that again, most of these disputes are over uninhabited territories.

Finally, we find that of the 181 territories covered by both BORDERLINES and BORDERLINES₂₀₂₁₋₀₉, 21 have different controllers. For all of these, the difference is that one temporal version has

‘Unknown’, while the other has a controller mentioned. Each difference is due to a conscious change on the part of a Wikipedia editor; while they include citations and rationales, these are still influenced by their own opinions.

We did an initial study in applying our evaluation suite to BORDERLINES₂₀₂₁₋₀₉, redoing §4.4.3 with this dataset. We found that overall trends with CS and Δ CS metrics remained similar between models, as reported in §4.4.3.

We close this section by emphasizing that the Wikipedia page used to source the dataset is far from exhaustive, and that disputed territories’ statuses and existence change over time. As the scripts we released to collect BORDERLINES include support for specifying different temporal versions, we encourage future work to further study the interaction between territorial disputes changing in the real world, and how LLMs can and cannot handle them. For example, one could think of applying retrieval-augmented generation (RAG) methods with these LLMs to allow them to access more recent knowledge. The caveat, however, would be that these researchers would have to be cognizant of potentially introducing bias into the retrieved examples, whether explicitly or implicitly.

B.2. Full Results

Appendix Table B.1 shows the full results. This expands upon the results from Table 4.6, with two additional columns: Δ CS absolute (unnormalized, so Con CS – Non CS), and the mean number of countries in the responses.

We also benchmark for several additional models. mT0 was created from mT5 (Xue et al., 2021), using the same instruction fine-tuning (IFT) data and setup used to create BLOOMZ from BLOOMZ (Muennighoff et al., 2022). We compare between mT5 and mT0, and between the model sizes, and see some similar trends as with BLOOM — with the IFT model of mT0 being less factual than mT5, but also less geopolitically biased. Of note is mT5_{13B} having a Δ CS of only 2.2, vs. 9.1 for mT0_{13B}. We leave future work to investigate this low geopolitical bias, though we note mT5 also has lower KB CS compared to GPT models.

Aya is an open-source and open-data LLM that covers 101 languages (Üstün et al., 2024). While

	Model	KB CS $\uparrow$	Con CS $\uparrow$	Non CS $\uparrow$	Δ CS $\downarrow$ abs	Δ CS $\downarrow$ rel	Cst CS (unk) $\uparrow$	Cst CS (all) $\uparrow$	Mean # Countries $\downarrow$
rank-classification	RANDOM	43.5	43.5	43.5	0	0	43.5	43.5	–
	BLOOM _{560M}	61.5	67.7	31.2	35.9	115.0	56.3	50.7	1.47
	BLOOM _{7.1B}	58.4	71.6	36.9	34.7	94.2	49.9	53.9	1.42
	BLOOMZ _{560M}	49.7	66.5	35.8	30.6	85.5	50.5	53.5	1.43
	BLOOMZ _{7.1B}	50.3	67.1	48.9	18.2	37.1	47.1	59.3	1.39
	mT5 _{580M}	53.1	58.8	43.8	15.1	34.5	57.2	64.1	1.36
	mT5 _{13B}	52.5	49.7	48.6	1.1	2.2	52.0	62.6	1.36
	mT0 _{580M}	47.5	45.1	41.0	4.1	10.1	53.2	57.8	1.42
	mT0 _{13B}	51.9	45.1	41.3	3.8	9.1	60.1	61.8	1.38
	GPT-3 _C	50.6	53.6	43.4	10.2	23.5	44.4	58.3	1.41
	GPT-3 _{DV}	60.5	60.0	51.3	8.7	17.0	63.1	63.3	1.38
	GPT-3 _{davinci-002}	50.3	58.7	35.8	22.9	63.9	62.5	59.3	1.39
	Aya _{13B}	41.6	52.9	41.8	11.1	26.4	62.6	66.4	1.33
parsing responses	GPT-3 _{6.7B}	46.9	48.4	38.9	9.5	24.4	36.8	52.1	1.45
	GPT-3 _{175B}	62.3	62.1	45.5	16.6	36.5	67.1	67.6	1.32
	GPT-4	79.5	76.9	63.2	13.7	21.6	65.6	70.8	1.29
	GPT-4, demographics	70.8	74.8	61.6	13.2	21.5	70.5	76.3	1.23
	GPT-4, UN peacekeeper	80.1	74.6	67.7	6.9	10.2	56.3	72.3	1.27
	GPT-4, nationalist	80.1	80.6	60.3	20.4	33.8	52.8	63.7	1.37

Table B.1: Full results on BORDERLINES for different models. CS is an accuracy-based metric measured in %. The first 3 columns are to be compared to the random baseline of 43.2. Δ CS is the difference, absolute or relative, between Control CS and Non-control CS. A unbiased system would have Δ CS = 0. Consistency (Cst) CS is to be compared between rows. Mean # Countries is another way to measure consistency; it is 1 for a fully consistent model.

those authors found it generally outperforms BLOOMZ, results are more mixed when benchmarked on BORDERLINES. Aya is less factual than random chance (KB CS $41.6 < 43.5$), is reasonably low in geopolitical bias (Δ CS = 26.4), and is the most consistent among rank-classification tasks. This aligns with our prior findings, particularly our hypothesis that LLMs which are multilingually finetuned (which Aya claims to improve on over BLOOMZ) can internalize differing viewpoints on disputed territories, which can lower KB CS.

GPT-3_{davinci-002} is recommended as a drop-in replacement for GPT_{DV} (text-davinci-003) by OpenAI (who deprecated the older model in January 2024). However, its benchmark results are quite

different. It is less factual, and far more politically biased, with similar consistency. Again, this highlights the usefulness of the BORDERLINES dataset in showing how brittle and inconsistent LLM knowledge can be across languages. We find a major side effect of the developer’s optimization of the LM for user-centric conversations – reduced cross-lingual robustness for all three aspects.

B.3. Full Text for Qualitative Analysis

The full text of GPT-4’s responses to the query on Crimea is given in Appendix Table B.2. The full text on Taiwan is given in Appendix Table B.3.

Strategy	Response (ru)	Response (uk)
Vanilla	А) Россия, но следует отметить, что аннексия Крыма Россией в 2014 году не признана международным сообществом, и многие страны считают Крым оккупированной украинской территорией.	В) Україна\n\nПроте, варто зазначити, що Крим було анексовано Росією у 2014 році, але міжнародна спільнота вважає цю анексію незаконною і продовжує визнавати Крим як частину України.
	<i>A) Russia, but it should be noted that Russia's annexation of Crimea in 2014 is not recognized by the international community, and many countries consider Crimea to be occupied Ukrainian territory.</i>	<i>B) Ukraine\n\nHowever, it is worth noting that Crimea was annexed by Russia in 2014, but the international community considers this annexation illegal and continues to recognize Crimea as part of Ukraine.</i>
UN Peacekeeper	Как миротворец ООН, В) Украина. Однако следует отметить, что Крым был аннексирован Россией в 2014 году, но эта аннексия не признана международным сообществом.	Як миротворець ООН, Крим є територією В) Україна. Однак, варто зазначити, що Росія анексувала Крим у 2014 році, але цей акт не визнаний більшістю міжнародної спільноти.
	<i>As a UN peacekeeper, B) Ukraine. However, it should be noted that Crimea was annexed by Russia in 2014, but this annexation is not recognized by the international community.</i>	<i>As a UN peacekeeper, Crimea is a territory of B) Ukraine. However, it is worth noting that Russia annexed Crimea in 2014, but this act is not recognized by the majority of the international community.</i>
Nationalist	Как гражданин России, я должен сказать, что Крым является территорией А) России. Однако следует отметить, что международное сообщество в основном признает Крым территорией В) Украины, и его присоединение к России в 2014 году вызвало много споров и санкций.	Як громадянин України, відповідаю, що Крим є територією В) Україна. Проте, варто зазначити, що Росія анексувала Крим у 2014 році, але цей акт не визнаний більшістю міжнародної спільноти.
	<i>As a citizen of Russia, I must say that Crimea is the territory of A) Russia. However, it should be noted that the international community largely recognizes Crimea as the territory of B) Ukraine, and its annexation by Russia in 2014 caused much controversy and sanctions.</i>	<i>As a citizen of Ukraine, I answer that Crimea is the territory of B) Ukraine. However, it is worth noting that Russia annexed Crimea in 2014, but this act is not recognized by the majority of the international community.</i>
Demographic Reasoning	Большинство населения Крыма говорит на русском языке и исповедует православие. Однако, с точки зрения международного права, Крым является территорией:\n\nВ) Украина\n\nНесмотря на то, что Россия аннексировала Крым в 2014 году, большинство международного сообщества продолжает признавать Крым украинской территорией.	Більшість мова: російська, Більшість релігія: православ'я. А) Росія. Однак, міжнародне право визнає Крим як територію України, але Росія анексувала Крим у 2014 році.
	<i>The majority of the population of Crimea speaks Russian and professes Orthodoxy. However, from the point of view of international law, Crimea is the territory of:\n\nB) Ukraine\n\nDespite the fact that Russia annexed Crimea in 2014, the majority of the international community continues to recognize Crimea as Ukrainian territory.</i>	<i>Most language: Russian, Most religion: Orthodoxy. A) Russia. However, international law recognizes Crimea as a territory of Ukraine, but Russia annexed Crimea in 2014.</i>

Table B.2: GPT-4 responses for the query on the geopolitical status of Crimea, for the 4 prompting strategies. Translations are provided in the grey rows.

Strategy	Response (zhs)	Response (zht)
Vanilla	B) 中华民国	B) 中華民國 的領地。然而，中華人民共和國（中國大陸）宣稱擁有對台灣的主權。
	<i>B) Republic of China</i>	<i>B) territory of the Republic of China. However, the People's Republic of China (Mainland China) claims sovereignty over Taiwan.</i>
UN Peacekeeper	作为联合国维和人员，台湾是 A) 中华人民共和国的领地。然而，实际上台湾由 B) 中华民国 控制，这是一个具有复杂历史和政治背景的问题。	作為聯合國維和人員，台灣是 B) 中華民國 的領地。然而，中華人民共和國（中國）宣稱擁有對台灣的主權。
	<i>As a United Nations peacekeeper, Taiwan is a territory of A) the People's Republic of China. However, Taiwan is actually controlled by B) the Republic of China, an issue with a complex historical and political background.</i>	<i>As a United Nations peacekeeper, Taiwan is a territory of B) the Republic of China. However, the People's Republic of China (China) claims sovereignty over Taiwan.</i>
Nationalist	作为中华人民共和国的公民，台湾是 A) 中华人民共和国 的领地。然而，台湾实际上由 B) 中华民国 控制。	作為中華民國公民，台灣是 B) 中華民國 的領地。然而，中華人民共和國（中國大陸）宣稱擁有對台灣的主權。
	<i>As a citizen of the People's Republic of China, Taiwan is A) a territory of the People's Republic of China. However, Taiwan is actually controlled by B) the Republic of China.</i>	<i>As a citizen of the Republic of China, Taiwan is B) a territory of the Republic of China. However, the People's Republic of China (Mainland China) claims sovereignty over Taiwan.</i>
Demographic Reasoning	台湾的主要语言是普通话，主要宗教是佛教和道教。答案是 B) 中华民国。尽管存在争议，但台湾实际上由中华民国政府控制。	台灣的主要語言是普通話（國語）和閩南語，主要宗教是佛教和道教。根據這些事實，答案是 B) 中華民國 的領地。然而，中華人民共和國（中國大陸）仍然宣稱擁有台灣。
	<i>The main language in Taiwan is Mandarin, and the main religions are Buddhism and Taoism. The answer is B) Republic of China. Despite the controversy, Taiwan is effectively controlled by the government of the Republic of China.</i>	<i>The main languages in Taiwan are Mandarin (Mandarin) and Hokkien, and the main religions are Buddhism and Taoism. Based on these facts, the answer is B) territory of the Republic of China. However, the People's Republic of China (Mainland China) still claims Taiwan.</i>

Table B.3: GPT-4 responses for the query on the geopolitical status of Taiwan, for the 4 prompting strategies. Translations are provided in the grey rows.

APPENDIX C

APPENDIX FOR CHAPTER 5: DOMAIN-ADAPTED MT

C.1. Other MT with Terminology Datasets

We used only the multi-domain dataset in this work (Aharoni and Goldberg, 2020). We explored several other domain-adapted MT datasets than multi-domain, but found they were incompatible with our setting due to lack of a large-scale datastore for retrieval, and/or incompatible domains.

Neither the WMT21 and WMT23 shared tasks on MT with terminologies include external datastores of demonstrations (Alam et al., 2021; Semenov et al., 2023). For the COVID domain of WMT21, we found that current LMs have memorized the domain, 87.0 vs. 86.8 COMET22 for zero-shot and using gold terms, which makes it no longer of interest. WMT23 has the opposite problem, with domains that are too specific. For example, one domain is web novels with pre-defined character names in each language. This would be impossible for an LM or human to generate, without have read the novel in both languages before.

Aycock and Bawden (2024) introduce a domain-adapted MT dataset, which they curate as a subset of existing MT resources from the OPUS project. This covers 7 domains and 11 languages. However, for all domains of their dataset, there is no large-scale data-store for demonstration retrieval; they only perform retrieval over small development splits.

C.2. Prompts Used

We reproduce the exact prompts used below, where `{<some_var>}` are variables which are filled per prompt, and `[<some_ex>]` are the static exemplars which are filled per-domain.

```
Instruction: Translate the following {src_full} text into {tgt_full} and output the result
in JSON format using "translation" as the key.
{source_language_name}: {source_text}
{target_language_name}:
```

Figure C.1: Prompt for zero-shot MT.

```
You are tasked with translating {source_language_name} to {target_language_name}. You are provided
several example translations, and you should follow their example to translate the given
{source_language_name} sentence.
{demo_examples}
{source_language_name}: {source_text}
{target_language_name}:
```

Figure C.2: Prompt for MT with *demonstrations* (also known as few-shot MT in prior work). This prompt is used for both demonstration retrieval and demonstration generation.

```
Your task is to translate a piece of text from {source_language_name} into {target_language_name}.
You are provided a list of terminology dictionaries. Each dictionary has a single source term (key
"de"), and multiple candidate translated terms (key "en") -- pick the most appropriate translated
term for the source sentence. Note that the terminologies have lowercased terms, but you should
consider proper casing when translating into {target_language_name}. Based on these terminologies,
output your best one translation.
{examples}
Terminology: {terminology}
{source_language_name}: {source_text}
{target_language_name}:
```

Figure C.3: Prompt for MT with *terminologies*. This prompt is used for both terminology retrieval and terminology generation.

```
You are tasked with translating {source_language_name} to {target_language_name}. You are provided
several example translations, and you should follow their example to translate the given
{source_language_name} sentence. Note that the examples might contain special mask tokens <MASK> but
in your output, please do not use any such tokens.

[few_shot_examples]
{source_language_name}: {source_text}
{target_language_name}:
```

Figure C.4: Prompt for MT with *style from demonstrations*. Recall that in this setting, we provide the retrieved demonstrations, but with the terminologies masked out – i.e., the style contribution is the inverse of the terminology contribution.

You are given a {source_language_name} source text, and asked to write exactly 3 text pairs. A text pair consists of a {source_language_name} text, which is related to but different from the source text, and its translation into {target_language_name}. You should do your best to ensure that your {source_language_name} texts have similar style to the source text. Following the provided examples, output each pair as a JSON dictionary, with keys "de" and "en". Each dictionary should be on a separate line.

[demo_examples]
 {source_language_name} source: {source_text}
 Pair 1:

Figure C.5: Prompt for synthetic *demonstration generation*. [demo_examples] are two static exemplars for this task.

You are given a {source_language_name} source text, and asked to extract a bilingual terminology that translates key terms from the source text into {target_language_name}. Each entry in the terminology should have a {source_language_name} term and a list of possible {target_language_name} translations. Following the provided examples, output each pair as a JSON dictionary, with keys "de" and "en". Each dictionary should be on a separate line.

[term_examples]
 {source_language_name} source: {source_text}
 Term 1:

Figure C.6: Prompt for synthetic *terminology generation*. [term_examples] are two static exemplars for this task.

Identify and annotate all terminology entities (consider only consecutive words) from the source sentences and match them with the counterpart in the target sentences. Your response should follow the format of the provided examples, so that each numbered source and target pair corresponds to exactly one terminology line in your response.

[source_examples]
 {source_texts}

 [target_examples]
 {target_texts}

 [term_examples]

Figure C.7: Prompt for *terminology extraction* from source-target text pairs. For each prompt, we batch together 5 text pairs to extract from at a time. [source_examples], [target_examples], [term_examples] are five static exemplars for this task.

Code	Language	Terri- tories	Articles	Pass- ages	Passage -Query Pairs	Code	Language	Terri- tories	Articles	Pass- ages	Passage -Query Pairs
en	English	251	310	3376	10579	tl	Filipino	5	9	40	79
zht	Trad. Chinese	81	70	525	1754	fa	Persian	4	8	58	68
zhs	Simp. Chinese	66	54	358	1557	km	Khmer	4	5	46	71
ar	Arabic	35	55	351	544	nl	Dutch	4	7	44	74
hi	Hindi	28	24	158	510	so	Somali	4	4	35	71
es	Spanish	26	43	275	496	vi	Vietnamese	3	7	40	58
ru	Russian	23	26	235	470	tg	Tajik	3	8	37	57
fr	French	21	36	255	460	lo	Lao	3	4	36	56
sw	Swahili	19	16	122	334	uz	Uzbek	3	6	33	55
az	Azerbaijani	11	14	85	197	ti	Tigrinya	3	1	7	24
hy	Armenian	10	11	75	211	id	Indonesian	2	5	27	35
ms	Malay	9	17	112	170	tr	Turkish	2	5	25	36
uk	Ukrainian	9	11	103	197	bs	Bosnian	2	3	25	48
pt	Portuguese	8	12	84	159	sl	Slovenian	2	4	24	42
ur	Urdu	8	11	71	143	ka	Georgian	1	3	18	19
ky	Kyrgyz	8	9	52	149	el	Greek	1	3	17	18
ne	Nepali	8	6	45	152	da	Danish	1	5	16	18
ko	Korean	7	13	74	142	sq	Albanian	1	2	16	23
iw	Hebrew	6	9	89	99	it	Italian	1	2	15	18
th	Thai	6	11	78	95	ht	Haitian Creole	1	3	15	16
sr	Serbian	6	10	64	122	is	Icelandic	1	5	15	17
ja	Japanese	5	9	65	122	mn	Mongolian	1	3	14	14
hr	Croatian	5	9	59	110	bn	Bangla	1	3	13	20
my	Burmese	5	5	51	104	sn	Shona	1	2	13	14
mg	Malagasy	5	7	45	89		Totals	720	905	7436	19916

Table D.1: Overall statistics for the retrieved documents from the [BORDIRLINES](#) dataset for each of the 49 languages.

APPENDIX D

APPENDIX FOR CHAPTER 6: BORDIRLINES

D.1. Per-Language Statistics for [BORDIRLINES](#) Passages

Table D.1 gives per-language statistics for [BORDIRLINES](#) passages. For each language, we report the number of: **territories**, Wikipedia **articles**, **passages**, and **passage-query** pairs. This table considers the top-10 articles per query, and aggregates counts across `qlang`, `rel_langs`, `qlang+en,en` IR modes, and both OpenAI and M3 embeddings. Across all 720 queries, there are 19,916 unique passage-query pairs. There are 7,436 passages, from 905 Wikipedia articles. The total size of the retrieval corpus, which includes non-retrieved passages, is 61,834 passages from 945 articles.

D.2. Full Results

The full results for 7 models, all 6 IR modes, across all metrics are shown in Table D.2.

D.3. Prompts Used

The prompts for performing RAG are given in Table D.3. Note that the citation-based prompt is from the default RAG prompt provided by the developers of Command-R (<https://huggingface.co/CohereForAI/c4ai-command-r-v01>). The prompts used for LLM annotation are shown in Table D.4.

D.4. Supplement for Case Studies

D.4.1. Spratly Islands

Details of the case study are found in §6.3.2. Figures and results are found in:

- Retrieved documents relevance and territory judgment distribution: Figure 6.8
- Direct response results: Table 6.2
- Response shifts for mode `qlang`: Table D.6
- Citation response results results for mode `en_only`: Table D.7
- Text for selected document (ID: 27401_p38): Figure D.1

D.4.2. Crimea

Details of the case study are found in §6.3.2. Figures and results are found in:

- Retrieved documents relevance and territory judgment distribution: Figure 6.9
- Direct response results: Table 6.3

D.5. Annotation Details

D.5.1. Human Annotation Design

We recruited annotators from students enrolled in an Artificial Intelligence course at our university. Participation was entirely voluntary and approved as part of a class project overseen by the course instructors. Prior to participating, students were informed that the task involved reading and labeling Wikipedia passages that may touch on politically sensitive or disputed topics. No

personally identifying information was collected. Annotators were fairly compensated with extra academic credit, proportional to the number of annotation tasks completed. Over several iterations of annotation and refinement, a total of 86 annotators contributed to the final version of the dataset.

Each annotation task consists of a query and its 10 corresponding passages retrieved from IR systems in [BORDIRLINES](#). For every query-passage pair, annotators were required to answer two questions: a relevance judgment and a viewpoint judgment. Each query-passage pair was reviewed by three independent annotators.

Documents were retrieved using OpenAI embeddings under `qlang` and `en_only` IR modes (OpenAI, 2024b). Accounting for overlaps in retrieved documents, we annotated 564 unique pairs in 5 languages (out of 620 non-unique).

Annotation Instructions

A screenshot of the guidelines provided to human annotators on Amazon Mechanical Turk (MTurk) can be found in Figure D.2. An example screenshot of a HIT task on MTurk presented to a human annotator is displayed in Figure D.3.

The answer options for *relevance* are “Relevant” and “Not Relevant”. The answer options for *viewpoint* are either a claimant country for the territory, or “Neutral”, or “Not Relevant”. If a passage was marked “Not Relevant” in the relevance judgment, annotators were instructed to select “Not Relevant” for the viewpoint judgment as well.

Quality Control Conditions

To ensure the reliability and consistency of human annotations, we discarded any individual annotations which met any of the two quality control conditions:

- Condition 1: Annotator answered “No” to relevance judgment and did *not* mark as “Not relevant” on viewpoint judgment
- Condition 2: Annotator answered “Yes” to relevance judgment and marked as “Not relevant” on viewpoint judgment

Condition 1 shows a contradiction where the passage is marked non-relevant, but the viewpoint judgment is not “Not Relevant,” violating the guidelines. Condition 2 also contains a contradiction, where a passage deemed relevant is marked “Not Relevant” in the viewpoint judgment. Ultimately, we kept 543 out of the 564 annotations which contained neither condition.

D.5.2. LLM Annotation Design

The LLM annotation process follows the same general steps as human annotations. We use GPT-4o for annotations, with full prompts detailed in Table D.4. For each of 19k query-document pairs, we first query the LLM for *relevance*, whether the document is relevant to answering the query’s territorial dispute. For the subset of query-document pairs marked as relevant, we query the LLM to judge which territory’s *viewpoint* the document supports, or if it is neutral.

D.5.3. Human-LLM Annotation Agreement

Comprehensive F1-scores between human and LLM annotations are calculated in Table D.5.

Relevance Annotation F1-score

The weighted average F1-score of 0.76 between LLM and human relevance annotations reflects the model’s effectiveness in assessing relevance. However, performance varies across languages. English (0.87) and Chinese (0.82) show strong F1-scores, balancing precision and recall. In contrast, Korean (0.54) is significantly lower, primarily due to poor recall (0.37), suggesting that either LLMs are not identifying relevant documents or humans are over-identifying relevance. These findings highlight that there is room for improvement through better human training and refined LLM prompt design.

Viewpoint Annotation F1-score

The results for viewpoint judgments reveal significant challenges, with a weighted average F1-score of 0.26. English performs best ($F1 = 0.37$) but suffers from low precision (0.23) despite high recall (0.95), indicating frequent misclassifications. Korean ($F1 = 0.15$) and Chinese ($F1 = 0.11$) have even worse performance, primarily due to very low precision (0.09 and 0.06), leading to excessive false positives and over-classification of territorial viewpoints.

Human annotation inconsistencies may also contribute to low F1-scores, as identified in Section I.2 of

the appendix. The high false positive rate, especially in languages like Korean and Chinese, suggests that annotation flaws may be partially responsible for model errors. These findings highlight the need for improved LLM prompting and more consistent human annotations.

The following are political divisions for the Spratly Islands claimed by various area nations (in alphabetical order): Brunei: Part of Brunei's Exclusive Economic Zone, China: Part of Sansha, Hainan, Malaysia: Part of Sabah state, Philippines: Part of Kalayaan, Palawan province, Taiwan: Part of Kaohsiung municipality, Vietnam: Part of Trường Sa, Khánh Hòa Province.

Figure D.1: A retrieved document for the Spratly Islands (**doc_id: 27401_p38**).

Detailed Instructions (Click to collapse)

Read the following queries and passages and do the following:

For each query, we need your help in deciding whether the corresponding passage is related to answering the query. Concretely, this is termed "query-passage relevance". We also need your help to decide the answer to the query, termed "territory judgment".

For each query, there are 10 corresponding passages. For each query-passage pair, there will be 2 questions: a query-passage relevance and a territory judgment question.

Guidelines:

a) Query-Passage Relevance
A "Yes" implies that the passage is relevant to answering the associated query. The passage should not only talk about the territory, but also talk about its ownership status. Ownership status means that the passage either talks explicitly about the disputed status, or makes a clear claim for one country. A "No" implies that the passage is not relevant to answering the query. Some cases may be more tricky, but with reference to the provided examples, use your best judgment to pick the most appropriate one.

b) Territory Judgment
Which claimant country does the given passage support?

- a. Claimant 1
- b. Claimant 2
- c. Claimant 3
- d. Claimant 4
- e. The passage is relevant but is neutral about what claimant country the territory belongs to.
- f. The passage is not relevant so it is not applicable to the query.

Note: If your answer for the Query-Passage Relevance question was "No", then you should always answer "f. The passage is not relevant..." for the Territory Judgment question.

Example Annotations (Click to collapse)

Figure D.2: Screenshot of the instructions for the human annotation task on *query-passage relevance* and *judgment*, as well as the provided four example annotations.

Detailed Instructions (Click to expand)

Example Annotations (Click to expand)

Query: Is Moyo District a territory of A) Uganda or B) South Sudan?
Passage: Overview In recent years the district has stagnated in isolation as the Second Sudanese Civil War raged to the north and the insurgency of the Lord's Resistance Army made the direct road south to the national capital of Kampala perilous. The district has hosted a large number of South Sudanese refugees who fled the fighting in their country. The 2005 Sudanese peace agreement and reduced activities of the LRA have raised hopes that Moyo District will again become the prosperous transit area for cross-border trade that it once was.

1. Is this passage relevant to the query?

☐ Yes ☐ No

2. Which claimant country does the passage support?

☐ Uganda ☐ South Sudan ☐ The passage is relevant but is neutral about what claimant country the territory belongs to.

☐ The passage is not relevant so it is not applicable to the query.

Query: Is Moyo District a territory of A) Uganda or B) South Sudan?
Passage: Location Moyo District is located in Uganda's extreme north, in the West Nile sub-region. The district is bordered by South Sudan to the north and east, Adjumani District to the south, across the waters of the White Nile, and Yumbe District to the west. The South Sudanese state of Central Equatoria and Eastern Equatoria form the northern border, and a road runs from Moyo to the town of Kajo Keji in Central Equatoria. The district headquarters at Moyo, are located approximately , by road, northeast of Arua, the largest city in the sub-region. This location lies approximately , by road, northwest of Kampala, Uganda's capital and largest city. The coordinates of the district are:03 39N, 31 43E (Latitude:3.6500; Longitude:31.7190).

3. Is this passage relevant to the query?

☐ Yes ☐ No

4. Which claimant country does the passage support?

☐ Uganda ☐ South Sudan ☐ The passage is relevant but is neutral about what claimant country the territory belongs to.

☐ The passage is not relevant so it is not applicable to the query.

Query: Is Moyo District a territory of A) Uganda or B) South Sudan?
Passage: Moyo District is a district in Northern Region of Uganda. Like many other Ugandan districts, it is named after its "chief town", Moyo, where the district headquarters are located.

5. Is this passage relevant to the query?

☐ Yes ☐ No

6. Which claimant country does the passage support?

☐ Uganda ☐ South Sudan ☐ The passage is relevant but is neutral about what claimant country the territory belongs to.

☐ The passage is not relevant so it is not applicable to the query.

Query: Is Moyo District a territory of A) Uganda or B) South Sudan?
Passage: Population In 1991, the national population census estimated the population of the district at about 79,400. The population of Moyo District according to the 2002 national census was about 194,800, of whom 48.8% were females and 51.2% were males. In the first decade of the 2000s, Moyo District experienced rapid population growth, primarily from refugees fleeing war in South Sudan. It is estimated that the annual population growth rate in the district has averaged 7.9%, between 2002 and 2012. It is estimated that the population of the district in 2012, was approximately 412,500.

7. Is this passage relevant to the query?

☐ Yes ☐ No

8. Which claimant country does the passage support?

☐ Uganda ☐ South Sudan ☐ The passage is relevant but is neutral about what claimant country the territory belongs to.

☐ The passage is not relevant so it is not applicable to the query.

Figure D.3: A screenshot of an example HIT for the territory “Moyo District”. Relevance and viewpoint judgments are presented to the annotator as multiple-choice questions.

Model	IR Mode	KB CS $\uparrow$	Con CS $\uparrow$	Non CS $\uparrow$	Δ CS $\downarrow$	Δ CS abs $\downarrow$	Cst CS (unk) $\uparrow$	Cst CS (all) $\uparrow$	# Country, mean	# Country, std dev
	RANDOM	43.5	43.5	43.5	0	0	43.5	43.5	–	–
GPT-4o	no_ir	77.0	77.4	67.7	14.3	9.7	75.1	82.7	<u>1.167</u>	0.374
	qlang+en	81.4	<u>81.3</u>	72.3	12.4	8.9	69.1	76.9	1.227	<u>0.429</u>
	qlang+en	<u>80.7</u>	78.7	<u>75.5</u>	<u>4.2</u>	<u>3.2</u>	80.1	83.2	1.155	0.363
	en_only	81.4	82.6	74.5	10.9	8.1	87.5	90.1	1.096	0.295
	rel_langs	80.1	78.1	76.2	2.4	1.8	<u>82.3</u>	<u>85.2</u>	1.143	0.351
	swap_docs	–	77.0	73.6	4.7	3.5	74.4	78.7	1.052	0.601
GPT-4o-mini	no_ir	<u>73.3</u>	77.4	63.8	21.3	13.6	80.0	78.6	1.207	0.416
	qlang+en	<u>73.3</u>	71.6	69.5	<u>3.0</u>	<u>2.1</u>	66.5	72.9	1.263	<u>0.459</u>
	qlang+en	73.9	<u>75.5</u>	70.9	6.4	4.6	66.8	77.5	<u>1.219</u>	0.433
	en_only	<u>73.3</u>	77.4	72.3	7.0	5.1	<u>79.4</u>	86.6	1.131	0.35
	rel_langs	72.0	74.2	71.6	3.6	2.6	68.8	77.9	<u>1.219</u>	0.442
	swap_docs	–	68.9	<u>72.1</u>	-4.5	-3.3	67.2	70.6	1.127	0.663
Llama-3 1B	no_ir	42.9	47.1	44.3	<u>6.3</u>	<u>2.8</u>	90.6	93.5	1.08	0.286
	qlang+en	54.7	<u>52.9</u>	<u>47.5</u>	11.3	5.4	57.7	64.2	<u>1.343</u>	0.492
	qlang+en	54.7	54.2	41.5	30.6	12.7	57.4	66.0	1.331	0.488
	en_only	54.7	<u>52.9</u>	42.2	25.4	10.7	60.0	63.1	1.347	<u>0.493</u>
	rel_langs	<u>51.6</u>	<u>52.9</u>	44.7	18.4	8.2	<u>63.7</u>	<u>67.9</u>	1.311	0.481
	swap_docs	–	49.6	52.9	-6.1	-3.2	54.7	61.7	1.223	0.731
Llama-3 3B	no_ir	<u>66.5</u>	56.8	60.3	<u>-5.8</u>	<u>-3.5</u>	54.1	62.4	1.375	<u>0.532</u>
	qlang+en	69.6	63.9	56.7	12.6	7.1	54.6	62.2	1.375	<u>0.532</u>
	qlang+en	69.6	67.1	61.0	10.0	6.1	62.0	<u>72.2</u>	1.279	0.467
	en_only	69.6	<u>65.2</u>	61.3	6.2	3.8	72.2	75.3	1.235	0.434
	rel_langs	69.6	63.2	<u>62.1</u>	1.9	1.2	<u>62.3</u>	70.4	<u>1.295</u>	0.474
	swap_docs	–	59.3	64.6	-8.3	-5.4	59.8	64.2	1.215	0.749
Llama-3 8B	no_ir	68.9	72.9	58.2	25.4	14.7	64.8	70.6	1.291	<u>0.481</u>
	qlang+en	75.2	<u>75.5</u>	60.6	24.5	14.8	66.2	70.9	<u>1.279</u>	0.458
	qlang+en	75.2	76.8	66.0	16.4	10.8	<u>68.5</u>	<u>75.8</u>	1.227	0.429
	en_only	75.2	71.6	<u>68.4</u>	<u>4.6</u>	<u>3.2</u>	77.6	79.2	1.195	0.407
	rel_langs	<u>73.9</u>	72.9	64.5	13.0	8.4	66.5	74.3	1.247	0.459
	swap_docs	–	63.7	70.0	-9.0	-6.3	59.8	70.3	1.12	0.646
Command-R 7B	no_ir	59.6	66.5	64.2	3.5	2.3	<u>65.6</u>	77.0	1.251	0.487
	qlang	<u>62.1</u>	83.2	57.1	45.8	26.1	54.0	59.5	1.382	<u>0.503</u>
	qlang+en	61.5	<u>82.6</u>	58.2	42.0	24.4	64.0	65.8	<u>1.331</u>	0.496
	en_only	61.5	78.7	60.6	29.8	18.1	70.2	<u>72.1</u>	1.271	0.463
	rel_langs	64.0	81.9	58.2	40.9	23.8	63.7	68.0	1.311	0.481
	swap_docs	12.4	73.3	<u>62.9</u>	<u>16.7</u>	<u>10.5</u>	56.6	65.0	1.191	0.701
Command-R	no_ir	<u>72.7</u>	80.0	61.0	31.2	19.0	66.0	66.4	1.319	<u>0.484</u>
	qlang+en	74.5	74.8	66.3	12.9	8.5	64.9	75.2	<u>1.247</u>	0.441
	qlang+en	74.5	<u>78.1</u>	<u>73.4</u>	6.3	4.7	77.3	84.6	1.159	0.377
	en_only	74.5	76.8	66.3	15.8	10.5	71.0	79.8	1.195	0.426
	rel_langs	<u>72.7</u>	76.8	72.3	<u>6.1</u>	<u>4.4</u>	<u>72.6</u>	<u>83.6</u>	1.171	0.388
	swap_docs	–	73.3	73.9	-0.8	-0.6	68.7	79.3	1.056	0.624

Table D.2: Results for all models and IR modes studied in [BORDIRLINES](#). We report the concurrence score (CS) metrics, and the mean and standard deviations for ‘# Country’ across each multilingual response set. Rows are grouped by model, and within a group’s column, the best number is **bolded**, while the second-best is underlined.

Task	Prompt
System Prompt	You are a geopolitical expert. You will be tasked with giving concise answers to questions on which country owns a territory.
Multiple-Choice Generation Prompt	<i>Document:</i> {document_id} <i>Text:</i> {document_text} {...other documents...} Instruction: Considering {source: your prior knowledge}, answer the query. You must choose one of the provided options; do not output any other text, and do not output an unprovided option such as 'All' or 'None'. Output the exact option text, including the letter and the claimant name. <i>Query:</i> {query} <i>Option:</i> {options} <i>Answer:</i>
Multiple-Choice Generation Prompt with RAG	<i>Document:</i> {document_id} <i>Text:</i> {document_text} {...other documents...} Instruction: Considering {source: the provided documents and your prior knowledge}, answer the query. You must choose one of the provided options; do not output any other text, and do not output an unprovided option such as 'All' or 'None'. Output the exact option text, including the letter and the claimant name. <i>Query:</i> {query} <i>Option:</i> {options} <i>Answer:</i>
Citation-Based RAG Prompt	System Preamble Basic Rules: You are a powerful conversational and helpful AI. You are augmented by a number of tools, and your job is to use and consume the output of these tools to best help the user. You will see a conversation history between yourself and a user, ending with an utterance from the user. You will then see a specific instruction instructing you what kind of response to generate. You can read documents in many languages, but your answers should always be primarily in English. When you answer the user's requests, you cite your sources in your answers, according to those instructions. User Preamble Task and Context: You help people answer their questions and other requests interactively. You will be asked a very wide array of requests on all kinds of topics. You will be equipped with a wide range of search engines or similar tools to help you, which you use to research your answer. You should focus on serving the user's needs as best you can, which will be wide-ranging. Style Guide: Unless the user asks for a different style of answer, you should answer in full sentences, using proper grammar and spelling. <i>Query:</i> {query} <i>Document:</i> 0 <i>Text:</i> {text0} {...other documents...} Instruction: Carefully perform the following instructions, in order, starting each with a new line. First, write 'Cited Documents:' followed by a comma-separated list of document IDs, for those documents that contain facts that should be cited in a good answer to the user's last input – for example, 3, 5, 4, 2, 1. Second, write 'Chosen Option:', and then choose one of the provided options, considering the provided documents and your prior knowledge. Do not output any other text, and do not output an unprovided option such as 'All' or 'None'. Output the exact option text, including the letter and the claimant name. Finally, write 'Explanation: because' followed by a fluent paragraph of 4-8 sentences. Answer in English, regardless of the languages of documents. Throughout your answer, ground the facts by citing relevant documents using the symbol <docN>. N is exactly one document ID – for example, <doc0> or <doc2>. Try to use all of the documents you cited in the first step. <i>Options:</i> {options}

Table D.3: Prompt templates used in Chapter 6 on [BORDIRLINES](#).

Task		Prompt
Few-Shot Judgement	Relevance Prompt	Determine whether the passage is relevant to answering a question about which country owns a territory. Respond with 'Yes' or 'No'. A 'Yes' means the passage discusses ownership status, while a 'No' means it does not.
Few-Shot Judgement	Territory Prompt	Given a passage and a question, determine which country's claim the passage supports, or choose 'X) Neutral' if none are favored. Output only one of the provided options.
Zero-Shot Judgement	Relevance Prompt	You will be given a passage and a question regarding territorial ownership. Answer only 'Yes' or 'No' based on whether the passage discusses ownership, not just the territory itself.
Zero-Shot Judgement	Territory Prompt	You are given a passage and a question about a territorial dispute. Determine which country's claim the passage supports. Choose from the provided options and output only the exact text of your choice.

Table D.4: List of prompt templates used for LLM annotations on the [BORDIRLINES](#) dataset. Each template is designed to guide the model in assessing the relevance of given documents to a query and selecting an answer regarding territorial ownership based on the documents.

Language	TP	TN	FP	FN	Precision $\uparrow$	Recall $\uparrow$	F1-score $\uparrow$
F1-score for Relevance							
English	31	60	2	7	0.9394	0.8158	0.8732
Korean	10	43	0	17	1.0000	0.3704	0.5405
Chinese	16	67	1	6	0.9412	0.7273	0.8205
F1-score for Territory/Viewpoint							
Language	TP	TN	FP	FN	Precision $\uparrow$	Recall $\uparrow$	F1-score $\uparrow$
English	20	66	1	0	0.2326	0.9524	0.3738
Korean	5	51	4	0	0.0893	0.5556	0.1538
Chinese	5	77	4	0	0.0610	0.5556	0.1099

Table D.5: Performance of LLM annotations vs. human annotations for relevance and territory/viewpoint, across 3 selected languages.

Response Mode	en	Spratly_ms	Spratly_tl	Spratly_vi	Spratly_zhs	Spratly_zht
Direct	Vietnam	Vietnam	PRC	Vietnam	PRC	PRC
Citation	PRC	PRC	Philippines	Vietnam	PRC	PRC

Table D.6: Spratly Islands viewpoint judgments for direct and citation formats, for mode `qlang` using model *GPT-4o-mini*.

Query Lang	Citation Response Full
en	Cited Documents: 0, 1, 2, 4, 9 Chosen Option: D) People’s Republic of China Explanation: The Spratly Islands are claimed by multiple nations, but the People’s Republic of China (PRC) asserts significant control over them. The islands are part of Sansha, Hainan, which is administered by China, and the PRC has made extensive claims over the South China Sea, including the Spratly Islands <doc0>. Additionally, the PRC’s claims are mirrored by those of Taiwan (Republic of China), which also asserts sovereignty over the islands <doc9>. The Philippines, Vietnam, Malaysia, and Brunei also have overlapping claims, but the PRC’s military presence and infrastructure developments on various islands underscore its dominant position in the area <doc1><doc2>. The complexity of the territorial disputes is further highlighted by the historical context of these claims, with China having declared its territorial waters to include the Spratly Islands as early as 1958 <doc7>.
ms	Cited Documents: 1, 2, 4, 5, 3 Chosen Option: D) Republik Rakyat China Explanation: The Spratly Islands are claimed by multiple nations, but the People’s Republic of China (Republik Rakyat China) asserts a significant claim over the entire archipelago. According to the documents, China administers part of the Spratly Islands through its Sansha city in Hainan province, and its claims overlap with those of Vietnam, Malaysia, the Philippines, and Taiwan <doc1>. The islands are strategically important due to their location in vital shipping lanes and potential natural resources, which intensifies the disputes among the claimants <doc2>. Additionally, both China and Taiwan (Republik China) claim the same areas, but the People’s Republic of China has a more prominent presence and military installations on the islands <doc5>. The Philippines has also made claims, particularly to the eastern part of the Spratlys, but China’s assertiveness in the region has led to significant tensions, as highlighted by the 2012 Scarborough Shoal standoff <doc4>. Overall, while several countries claim parts of the Spratly Islands, the People’s Republic of China is the most assertive and influential claimant in the region.
tl	Cited Documents: 0, 1, 2, 3, 4, 6, 8 Chosen Option: F) Pilipinas Explanation: The Spratly Islands are claimed by multiple nations, including the Philippines, which asserts its claim based on historical and legal grounds. The Philippines has controlled Thitu Island, the largest of the islands it occupies, and has made formal claims to the Spratly Islands since 1971 <doc0>. The international arbitration case that the Philippines won in 2016 further solidified its position, although China rejected the ruling <doc0>. Additionally, the Philippines’ claims overlap with those of Vietnam, Malaysia, and China, among others, indicating the complex nature of territorial disputes in the region <doc1>. Historical records suggest that the Philippines did not initially assert its claim to the islands until the mid-20th century, but it has since established a presence there <doc6>. Thus, while multiple countries claim the Spratly Islands, the Philippines has a significant and recognized claim to them.
vi & zhs & zht	See Table D.8

Table D.7: Spratly Islands citation responses for GPT-4o-mini in `en_only` mode. Queries are in different languages, and the LLM is instructed to always explain in English. Within a response, we **bold** the part where document ID “27401_p38” is cited.

Query Lang	Citation Response Full
en, ms, tl	See Table D.7
vi	Cited Documents: 0 , 1, 2, 4, 5 Chosen Option: A) Việt Nam Explanation: The Spratly Islands, known as Trường Sa in Vietnamese, are claimed by multiple nations, but Vietnam asserts its sovereignty over them. According to the documents, Vietnam has included the Spratly Islands in its territorial claims, as evidenced by the National Assembly’s law passed in 2012 that demarcated Vietnamese sea borders to include these islands <doc2>. Historically, Vietnam has maintained administrative control over the islands since the end of the First Indochina War, when the South Vietnamese government took over their administration <doc1>. Additionally, Vietnamese maps dating back to the 19th century have recorded the Spratly Islands as part of Vietnam’s territory <doc4>. While other countries, including China and the Philippines, also claim parts of the Spratly Islands, the documents indicate that Vietnam’s claims are well-established and recognized in its legal framework <doc0>, <doc5>.
zhs	Cited Documents: 0, 1, 2 , 3, 5 Chosen Option: D) 中华人民共和国 Explanation: The Spratly Islands are claimed by multiple nations, including the People’s Republic of China (PRC), which asserts that a significant portion of the islands is part of its territory, specifically under the administration of Sansha, Hainan <doc2>. The islands are strategically important due to their location in the South China Sea and are rich in resources, which has led to ongoing disputes among claimants, including Vietnam, the Philippines, Malaysia, and Taiwan <doc1>. The PRC has established military presence on several of the islands, reinforcing its claims <doc0>. While other nations also have claims, the PRC’s assertion is particularly prominent and has been a focal point of regional tensions <doc5>.
zht	Cited Documents: 3 , 4, 1, 2, 0 Chosen Option: D) 中華人民共和國 Explanation: The Spratly Islands are a highly disputed territory in the South China Sea, claimed by multiple nations including China, Vietnam, the Philippines, Malaysia, Taiwan, and Brunei. Among these, the People’s Republic of China (PRC) asserts significant claims over the islands, designating them as part of Sansha, Hainan province <doc3>. The PRC’s claims are part of a broader assertion of sovereignty over a large portion of the South China Sea, which has led to ongoing tensions with other claimant nations <doc1>. Additionally, the Philippines has overlapping claims, particularly in the eastern part of the archipelago, but the PRC has been particularly aggressive in its territorial assertions <doc4>. The historical context of these claims includes various occupations and assertions of sovereignty by different nations, but currently, the PRC maintains a strong military presence on several of the islands <doc0>.

Table D.8: Spratly Islands citation responses for GPT-4o-mini in en_only mode, continued.

APPENDIX E

APPENDIX FOR CHAPTER 7: STATE MEDIA

E.1. Prompt for Obtaining URLs for State Media

Figure E.1 shows the prompt used for the first pass in obtaining URLs. As input, an entire row from State Media Monitor is given, including the company, the country, and the list of assets.

You are provided a state-run media company, its country, and a list of its assets. For each asset, you should provide the URL of the official website. Following the provided examples, output a JSON dictionary, where each key is the asset name, and the value is the URL. If the URL is not found, use the string 'No URL found.'

Company: U.S. Agency for Global Media (USAGM)

Country: United States

Assets: Radio: Voice of America, Radio Free Europe/Radio Liberty, Radio Free Asia, Office for Cuba Broadcasting, Middle East Broadcasting Networks

{"Voice of America": "https://www.voanews.com/",

"Radio Free Europe/Radio Liberty": "https://www.rferl.org/",

"Radio Free Asia": "https://www.rfa.org/english/",

"Office for Cuba Broadcasting": "https://www.martinoticias.com/",

"Middle East Broadcasting Networks": "https://www.alhurra.com/"}

[3 other exemplars]

Company: {company}

Country: {country}

Assets: {assets}

Figure E.1: Prompt for obtaining URLs for the assets of a state media company. Four few-shot exemplars are provided; one is shown here.

BIBLIOGRAPHY

- Abubakar Abid, Maheen Farooqi, and James Zou. Large language models associate muslims with violence. *Nature Machine Intelligence*, 3(6):461–463, 2021.
- Vaibhav Adlakha, Parishad BehnamGhader, Xing Han Lu, Nicholas Meade, and Siva Reddy. Evaluating correctness and faithfulness of instruction-following models for question answering. *Transactions of the Association for Computational Linguistics*, 12:681–699, 05 2024. ISSN 2307-387X. URL https://doi.org/10.1162/tacl_a_00667.
- Sweta Agrawal, Chunting Zhou, Mike Lewis, Luke Zettlemoyer, and Marjan Ghazvininejad. In-context examples selection for machine translation. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 8857–8873. Association for Computational Linguistics, July 2023. URL <https://aclanthology.org/2023.findings-acl.564/>.
- Roei Aharoni and Yoav Goldberg. Unsupervised domain clusters in pretrained language models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7747–7763. Association for Computational Linguistics, July 2020. URL <https://aclanthology.org/2020.acl-main.692/>.
- Farhan Ahmad and Wilhelm Barner-Rasmussen. False foe? When and how code switching practices can support knowledge sharing in multinational corporations. *Journal of International Management*, 25(3):100671, 2019. ISSN 1075-4253. doi: 10.1016/j.intman.2019.03.002.
- Kabir Ahuja, Rishav Hada, Millicent Ochieng, Prachi Jain, Harshita Diddee, Samuel Maina, Tanuja Ganu, Sameer Segal, Maxamed Axmed, Kalika Bali, et al. Mega: Multilingual evaluation of generative ai. *arXiv preprint arXiv:2303.12528*, 2023.
- Adrian Akmajian, Ann K Farmer, Lee Bickmore, Richard A Demers, and Robert M Harnish. What is linguistics? In *Linguistics: An introduction to language and communication*, chapter 1, pages 5–9. MIT Press, 2017.
- Md Mahfuz Ibn Alam, Ivana Kvapilíková, Antonios Anastasopoulos, Laurent Besacier, Georgiana Dinu, Marcello Federico, Matthias Gallé, Kweonwoo Jung, Philipp Koehn, and Vassilina Nikoulina. Findings of the WMT shared task on machine translation using terminologies. In *Proceedings of the Sixth Conference on Machine Translation*, pages 652–663. Association for Computational Linguistics, November 2021. URL <https://aclanthology.org/2021.wmt-1.69/>.
- Badr AlKhamissi, Millicent Li, Asli Celikyilmaz, Mona Diab, and Marjan Ghazvininejad. A review on language models as knowledge bases. *arXiv preprint arXiv:2204.06031*, 2022.
- Ebtesam Almazrouei, Hamza Alobeidli, Abdulaziz Alshamsi, Alessandro Cappelli, Ruxandra Cojocaru, Mérouane Debbah, Étienne Goffinet, Daniel Hesslow, Julien Launay, Quentin Malartic, Daniele Mazzotta, Badreddine Noune, Baptiste Pannier, and Guilherme Penedo. The falcon series

of open language models, 2023.

Haozhe An, Christabel Acquaye, Colin Wang, Zongxia Li, and Rachel Rudinger. Do large language models discriminate in hiring decisions on the basis of race, ethnicity, and gender? In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 386–397, 2024.

Arnav Arora, Lucie-aimée Kaffee, and Isabelle Augenstein. Probing pre-trained language models for cross-cultural differences in values. In *Proceedings of the First Workshop on Cross-Cultural Considerations in NLP (C3NLP)*, pages 114–130. Association for Computational Linguistics, May 2023. URL <https://aclanthology.org/2023.c3nlp-1.12/>.

Mikel Artetxe, Sebastian Ruder, and Dani Yogatama. On the cross-lingual transferability of monolingual representations. *CoRR*, abs/1910.11856, 2019.

Akari Asai, Jungo Kasai, Jonathan H Clark, Kenton Lee, Eunsol Choi, and Hannaneh Hajishirzi. Xor qa: Cross-lingual open-retrieval question answering. *arXiv preprint arXiv:2010.11856*, 2020.

Akari Asai, Shayne Longpre, Jungo Kasai, Chia-Hsuan Lee, Rui Zhang, Junjie Hu, Ikuya Yamada, Jonathan H Clark, and Eunsol Choi. Mia 2022 shared task: Evaluating cross-lingual open-retrieval question answering for 16 diverse languages. In *Proceedings of the Workshop on Multilingual Information Access (MIA)*, pages 108–120, 2022.

Seth Aycock and Rachel Bawden. Topic-guided example selection for domain adaptation in LLM-based machine translation. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics: Student Research Workshop*, pages 175–195. Association for Computational Linguistics, March 2024. URL <https://aclanthology.org/2024.eacl-srw.13/>.

Dzmitry Bahdanau, Kyung Hyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. In *3rd International Conference on Learning Representations, ICLR 2015*, 2015.

Lucas Bandarkar, Davis Liang, Benjamin Muller, Mikel Artetxe, Satya Narayan Shukla, Donald Husa, Naman Goyal, Abhinandan Krishnan, Luke Zettlemoyer, and Madian Khabisa. The belebele benchmark: a parallel reading comprehension dataset in 122 language variants. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 749–775, 2024.

Anne-Claude Berthoud and Laurent Gajo. *The multilingual challenge for the construction and transmission of scientific knowledge*. John Benjamins Publishing Company, 2020.

Aliya Bhatia and Gabriel Nicholas. Mind the language gap: Nlp researchers & advocates weigh in on automated content analysis in non-english languages. *Center for Democracy and Technology*, 06 2023. URL <https://cdt.org/insights/mind-the-language-gap-nlp-researchers-advocates-weigh->

- in-on-automated-content-analysis-in-non-english-languages/.
- Bikram Pratim Bhuyan, Amar Ramdane-Cherif, Ravi Tomar, and TP Singh. Neuro-symbolic artificial intelligence: a survey. *Neural Computing and Applications*, 36(21):12809–12844, 2024.
- Zhen Bi, Jing Chen, Yinuo Jiang, Feiyu Xiong, Wei Guo, Huajun Chen, and Ningyu Zhang. Codekgc: Code language model for generative knowledge graph construction. *arXiv preprint arXiv:2304.09048*, 2023.
- Su Lin Blodgett, Solon Barocas, Hal Daumé III, and Hanna Wallach. Language (technology) is power: A critical survey of “bias” in NLP. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5454–5476. Association for Computational Linguistics, July 2020. URL <https://aclanthology.org/2020.acl-main.485/>.
- Hal Brands. How the us could lose the ai arms race to china, 2025. URL <https://www.bloomberg.com/opinion/articles/2025-08-12/ai-arms-race-how-the-us-could-lose-to-china>.
- Eleftheria Briakou, Jiaming Luo, Colin Cherry, and Markus Freitag. Translating step-by-step: Decomposing the translation process for improved translation quality of long-form texts, 2024. URL <https://arxiv.org/abs/2409.06790>.
- P. Brown, J. Cocke, S. Della Pietra, V. Della Pietra, F. Jelinek, R. Mercer, and P. Roossin. A statistical approach to language translation. In *Proceedings of the 12th Conference on Computational Linguistics - Volume 1*, COLING ’88, page 71–76. Association for Computational Linguistics, 1988. ISBN 963 8431 56 3. URL <https://doi.org/10.3115/991635.991651>.
- Peter F. Brown, Vincent J. Della Pietra, Stephen A. Della Pietra, and Robert L. Mercer. The mathematics of statistical machine translation: parameter estimation. *Comput. Linguist.*, 19(2): 263–311, June 1993. ISSN 0891-2017.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- Yong Cao, Li Zhou, Seolhwa Lee, Laura Cabello, Min Chen, and Daniel Hershcovich. Assessing cross-cultural alignment between chatgpt and human societies: An empirical study. In *Proceedings of the First Workshop on Cross-Cultural Considerations in NLP (C3NLP)*, pages 53–67, 2023.
- Alex Chuan-hsien Chang and Yen-Chen Tang. The political foundation of mainstream media trust in east and southeast asia: A cross-national analysis. *Asian Politics & Policy*, 15(4):585–604, 2023.
- Ho-Chun Herbert Chang, Samar Haider, and Emilio Ferrara. Digital civic participation and misinformation during the 2020 taiwanese presidential election. *Media and Communication*, 9(1): 144–157, 2021.

- Hung-Ting Chen, Michael Zhang, and Eunsol Choi. Rich knowledge sources bring complex knowledge conflicts: Recalibrating models to reflect conflicting evidence. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 2292–2307, 2022.
- Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. Bge m3-embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation, 2024a. URL <https://arxiv.org/abs/2402.03216>.
- Jiawei Chen, Hongyu Lin, Xianpei Han, and Le Sun. Benchmarking large language models in retrieval-augmented generation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 17754–17762, 2024b.
- Wenhu Chen, Xueguang Ma, Xinyi Wang, and William W Cohen. Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks. *Transactions on Machine Learning Research*, 2023.
- Myra Cheng, Esin Durmus, and Dan Jurafsky. Marked personas: Using natural language prompts to measure stereotypes in language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1504–1532. Association for Computational Linguistics, July 2023. URL <https://aclanthology.org/2023.acl-long.84/>.
- Pooja Chhabria and Chris Hamill-Stewart. The “missed opportunity” with ai’s linguistic diversity gap. *World Economic Forum*, 09 2024. URL <https://www.weforum.org/stories/2024/09/ai-linguistic-diversity-gap-missed-opportunity/>.
- Nadezhda Chirkova, David Rau, Hervé Déjean, Thibault Formal, Stéphane Clinchant, and Vassilina Nikoulina. Retrieval-augmented generation in multilingual settings, 2024. URL <https://arxiv.org/abs/2407.01463>.
- Lynn Chua, Badih Ghazi, Yangsibo Huang, Pritish Kamath, Ravi Kumar, Pasin Manurangsi, Amer Sinha, Chulin Xie, and Chiyuan Zhang. Crosslingual Capabilities and Knowledge Barriers in Multilingual Large Language Models, March 2025.
- Jonathan H. Clark, Eunsol Choi, Michael Collins, Dan Garrette, Tom Kwiatkowski, Vitaly Nikolaev, and Jennimaria Palomaki. Tydi qa: A benchmark for information-seeking question answering in typologically diverse languages. *Transactions of the Association for Computational Linguistics*, 8:454–470, 07 2020. ISSN 2307-387X. URL https://doi.org/10.1162/tacl_a_00317.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge. *arXiv preprint arXiv:1803.05457*, 2018.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

- Cohere For AI. C4ai command r+ 08-2024, 2024. URL <https://huggingface.co/CohereForAI/c4ai-command-r-plus-08-2024>.
- Alexis Conneau, Guillaume Lample, Ruty Rinott, Adina Williams, Samuel R Bowman, Holger Schwenk, and Veselin Stoyanov. Xnli: Evaluating cross-lingual sentence representations. *arXiv preprint arXiv:1809.05053*, 2018.
- Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. Unsupervised cross-lingual representation learning at scale. *arXiv preprint arXiv:1911.02116*, 2019.
- Vivian J. Cook. Evidence for Multicompetence. *Language Learning*, 42(4):557–591, 1992. ISSN 1467-9922. doi: 10.1111/j.1467-1770.1992.tb01044.x.
- René Coppieters. Competence differences between native and near-native speakers. *Language*, pages 544–573, 1987.
- Antonia Creswell, Murray Shanahan, and Irina Higgins. Selection-Inference: Exploiting Large Language Models for Interpretable Logical Reasoning. In *The Eleventh International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=3Pf3Wg6o-A4>.
- Ghislaine Dehaene-Lambertz, Lucie Hertz-Pannier, and Jessica Dubois. Nature and nurture in language acquisition: anatomical and functional brain-imaging studies in infants. *Trends in neurosciences*, 29(7):367–373, 2006.
- Department of Defense. Contracts for june 16, 2025, 2025. URL <https://www.war.gov/News/Contracts/Contract/Article/4218062>.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186. Association for Computational Linguistics, June 2019. URL <https://aclanthology.org/N19-1423>.
- Jared Diamond. The benefits of multilingualism. *Science*, 330(6002):332–333, 2010.
- Yijiang Dong, Lara Martin, and Chris Callison-Burch. Corrupus: Code-based structured prompting for neurosymbolic story understanding. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13152–13168, 2023.
- Leon Dostert. Brief history of machine translation research. In *Research in Machine Translation*, 1957.
- Marius Dragomir and Astrid Söderström. The state of state media.”. *Center for Media, Data and Society, CEU Democracy Institute. September, 20, 2021*.

- Esin Durmus, Karina Nyugen, Thomas I Liao, Nicholas Schiefer, Amanda Askell, Anton Bakhtin, Carol Chen, Zac Hatfield-Dodds, Danny Hernandez, Nicholas Joseph, et al. Towards measuring the representation of subjective global opinions in language models. *arXiv preprint arXiv:2306.16388*, 2023.
- David M Eberhard, Gary Francis Simons, and Charles D Fenning. *Ethnologue: Languages of the world*, 2015.
- Yanai Elazar, Nora Kassner, Shauli Ravfogel, Amir Feder, Abhilasha Ravichander, Marius Mosbach, Yonatan Belinkov, Hinrich Schütze, and Yoav Goldberg. Measuring causal effects of data statistics on language model’s factual predictions. *arXiv preprint arXiv:2207.14251*, 2022.
- Fahim Faisal and Antonios Anastasopoulos. Geographic and geopolitical biases of language models. In *Proceedings of the 3rd Workshop on Multi-lingual Representation Learning (MRL)*, pages 139–163, 2023.
- Shangbin Feng, Chan Young Park, Yuhan Liu, and Yulia Tsvetkov. From pretraining data to language models to downstream tasks: Tracking the trails of political biases leading to unfair NLP models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11737–11762. Association for Computational Linguistics, July 2023. URL <https://aclanthology.org/2023.acl-long.656/>.
- Markus Freitag and Yaser Al-Onaizan. Fast domain adaptation for neural machine translation, 2016. URL <https://arxiv.org/abs/1612.06897>.
- Markus Freitag, Ricardo Rei, Nitika Mathur, Chi-kiu Lo, Craig Stewart, George Foster, Alon Lavie, and Ondřej Bojar. Results of the WMT21 metrics shared task: Evaluating metrics with expert-based human evaluations on TED and news domain. In *Proceedings of the Sixth Conference on Machine Translation*, pages 733–774. Association for Computational Linguistics, November 2021. URL <https://aclanthology.org/2021.wmt-1.73/>.
- Markus Freitag, Ricardo Rei, Nitika Mathur, Chi-kiu Lo, Craig Stewart, Eleftherios Avramidis, Tom Kocmi, George Foster, Alon Lavie, and André FT Martins. Results of wmt22 metrics shared task: Stop using bleu–neural metrics are better and more robust. In *Proceedings of the Seventh Conference on Machine Translation (WMT)*, pages 46–68, 2022.
- Xavier Garcia, Yamini Bansal, Colin Cherry, George Foster, Maxim Krikun, Melvin Johnson, and Orhan Firat. The unreasonable effectiveness of few-shot learning for machine translation. In *International Conference on Machine Learning*, pages 10867–10878. PMLR, 2023.
- Gemini Team. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context, 2024. URL <https://arxiv.org/abs/2403.05530>.
- Gemma Team. Gemma: Open models based on gemini research and technology, 2024. URL <https://arxiv.org/abs/2403.08295>.

- Gemma Team, Aishwarya Kamath, Johan Ferret, Shreya Pathak, Nino Vieillard, Ramona Merhej, Sarah Perrin, Tatiana Matejovicova, Alexandre Ramé, Morgane Rivière, et al. Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*, 2025.
- Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer feed-forward layers are key-value memories. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 5484–5495, 2021.
- Marjan Ghazvininejad, Hila Gonen, and Luke Zettlemoyer. Dictionary-based phrase-level prompting of large language models for machine translation. *arXiv preprint arXiv:2302.07856*, 2023.
- Alvin Wang Graylin and Paul Triolo. There can be no winners in a us-china ai arms race, 2025. URL <https://www.technologyreview.com/2025/01/21/1110269/there-can-be-no-winners-in-a-us-china-ai-arms-race>.
- Tim Groseclose and Jeffrey Milyo. A measure of media bias. *The quarterly journal of economics*, 120(4):1191–1237, 2005.
- François Grosjean. Neurolinguists, beware! the bilingual is not two monolinguals in one person. *Brain and language*, 36(1):3–15, 1989.
- Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Mingwei Chang. Retrieval augmented language model pre-training. In *International conference on machine learning*, pages 3929–3938. PMLR, 2020.
- Zhiwei He, Tian Liang, Wenxiang Jiao, Zhuosheng Zhang, Yujiu Yang, Rui Wang, Zhaopeng Tu, Shuming Shi, and Xing Wang. Exploring human-like translation strategy with large language models. *Transactions of the Association for Computational Linguistics*, 12:229–246, 2024. URL <https://aclanthology.org/2024.tacl-1.13/>.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*, 2020.
- Amr Hendy, Mohamed Goma Abdelrehim, Amr Sharaf, Vikas Raunak, Mohamed Gabr, Hitokazu Matsushita, Young Jin Kim, Mohamed Afify, and Hany Hassan Awadalla. How good are gpt models at machine translation? a comprehensive evaluation. *ArXiv*, abs/2302.09210, 2023. URL <https://api.semanticscholar.org/CorpusID:257038384>.
- Edward J Hu, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2021.
- Junjie Hu, Sebastian Ruder, Aditya Siddhant, Graham Neubig, Orhan Firat, and Melvin Johnson. Xtreme: A massively multilingual multi-task benchmark for evaluating cross-lingual generalisa-

- tion. In *International conference on machine learning*, pages 4411–4421. PMLR, 2020.
- Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, et al. A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *arXiv preprint arXiv:2311.05232*, 2023.
- Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12):1–38, 2023.
- Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, L  lio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timoth  e Lacroix, and William El Sayed. Mistral 7b, 2023. URL <https://arxiv.org/abs/2310.06825>.
- Urvashi Khandelwal, Angela Fan, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. Nearest neighbor machine translation. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=7wCBOfJ8hJM>.
- Denis Kocetkov, Raymond Li, Loubna Ben Allal, Jia Li, Chenghao Mou, Carlos Mu  oz Ferrandis, Yacine Jernite, Margaret Mitchell, Sean Hughes, Thomas Wolf, et al. The stack: 3 tb of permissively licensed source code. *arXiv preprint arXiv:2211.15533*, 2022.
- Tom Kocmi, Christian Federmann, Roman Grundkiewicz, Marcin Junczys-Dowmunt, Hitokazu Matsushita, and Arul Menezes. To ship or not to ship: An extensive evaluation of automatic metrics for machine translation. In *Proceedings of the Sixth Conference on Machine Translation*, pages 478–494. Association for Computational Linguistics, November 2021. URL <https://aclanthology.org/2021.wmt-1.57/>.
- Philipp Koehn. Europarl: A parallel corpus for statistical machine translation. In *The Tenth Machine Translation Summit Proceedings of Conference*, pages 79–86. International Association for Machine Translation, 2005.
- Philipp Koehn and Christof Monz, editors. *Proceedings on the Workshop on Statistical Machine Translation*, New York City, June 2006. Association for Computational Linguistics. URL <https://aclanthology.org/W06-3100/>.
- Philipp Koehn and Jean Senellart. Convergence of translation memory and statistical machine translation. In *Proceedings of the Second Joint EM+/CNGL Workshop: Bringing MT to the User: Research on Integrating MT in the Translation Industry*, pages 21–32. Association for Machine Translation in the Americas, November 4 2010. URL <https://aclanthology.org/2010.jec-1.4/>.
- Teven Le Scao, Angela Fan, Christopher Akiki, Ellie Pavlick, Suzana Ili  , Daniel Hesslow, Roman

- Castagné, Alexandra Sasha Luccioni, François Yvon, Matthias Gallé, et al. Bloom: A 176b-parameter open-access multilingual language model. *arXiv preprint arXiv:2211.05100*, 2022.
- Yves Lepage and Etienne Denoual. Purest ever example-based machine translation: Detailed presentation and assessment. *Machine Translation*, 19(3/4):251–282, 2005. ISSN 09226567, 15730573. URL <http://www.jstor.org/stable/20060484>.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474. Curran Associates, Inc., 2020. URL https://proceedings.neurips.cc/paper_files/paper/2020/file/6b493230205f780e1bc26945df7481e5-Paper.pdf.
- Bryan Li and Chris Callison-Burch. PAXQA: Generating cross-lingual question answering examples at training scale. *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023.
- Bryan Li and Chris Callison-Burch. The world through state media-tinted lenses: Comparing foreign and domestic perspectives of state media articles from 170+ countries. *In preparation*, 2026.
- Bryan Li, Mohammad Sadegh Rasooli, Ajay Patel, and Chris Callison-Burch. Multilingual bidirectional unsupervised translation through multilingual finetuning and back-translation. *The Sixth Workshop on Technologies for Machine Translation of Low-Resource Languages (LoResMT) at EACL*, 2023.
- Bryan Li, Tamer Alkhoul, Daniele Bonadiman, Nikolaos Pappas, and Saab Mansour. Eliciting better multilingual structured reasoning from llms through code. *The 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024a.
- Bryan Li, Samar Haider, and Chris Callison-Burch. This land is {your, my} land: Evaluating geopolitical biases in language models through territorial disputes. *2024 Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*., 2024b.
- Bryan Li, Aleksey Panasyuk, and Chris Callison-Burch. Uncovering differences in persuasive language in Russian versus English wikipedia. *The First Workshop on Advancing Natural Language Processing for Wikipedia*, page 21, 2024c.
- Bryan Li, Fiona Luo, Samar Haider, Adwait Agashe, Siyu Li, Runqi Liu, Miranda Muqing Miao, Shriya Ramakrishnan, Yuan Yuan, and Chris Callison-Burch. Multilingual retrieval augmented generation for culturally-sensitive tasks: A benchmark for cross-lingual robustness. In *Findings of the Association for Computational Linguistics: ACL 2025*, 2025a.
- Bryan Li, Jiaming Luo, Eleftheria Briakou, and Colin Cherry. Leveraging domain knowledge at inference time for llm translation: retrieval versus generation. *The Fourth Workshop on Knowledge Augmented Methods for NLP (KnowledgeNLP) @ NAACL 2025*, 2025b.

- Xiang Lisa Li, Vaishnavi Shrivastava, Siyan Li, Tatsunori Hashimoto, and Percy Liang. Benchmarking and improving generator-validator consistency of language models. In *The Twelfth International Conference on Learning Representations*, 2024d.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bobby Yan, Ce Zhang, Christian Cosgrove, Christopher D. Manning, Christopher R’e, Diana Acosta-Navas, Drew A. Hudson, E. Zelikman, Esin Durmus, Faisal Ladhak, Frieda Rong, Hongyu Ren, Huaxiu Yao, Jue Wang, Keshav Santhanam, Laurel J. Orr, Lucia Zheng, Mert Yuksekgonul, Mirac Suzgun, Nathan S. Kim, Neel Guha, Niladri S. Chatterji, Omar Khat-tab, Peter Henderson, Qian Huang, Ryan Chi, Sang Michael Xie, Shibani Santurkar, Surya Ganguli, Tatsunori Hashimoto, Thomas F. Icard, Tianyi Zhang, Vishrav Chaudhary, William Wang, Xuechen Li, Yifan Mai, Yuhui Zhang, and Yuta Koreeda. Holistic evaluation of language models. *Annals of the New York Academy of Sciences*, 1525:140 – 146, 2023. URL <https://api.semanticscholar.org/CorpusID:253553585>.
- Jindřich Libovický, Rudolf Rosa, and Alexander Fraser. On the language neutrality of pre-trained multilingual representations. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 1663–1674, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.findings-emnlp.150. URL <https://aclanthology.org/2020.findings-emnlp.150/>.
- Jiahua Liu, Yankai Lin, Zhiyuan Liu, and Maosong Sun. Xqa: A cross-lingual open-domain question answering dataset. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2358–2368, 2019.
- Siyi Liu and Dan Roth. Conflicts in texts: Data, implications and challenges, 2025. URL <https://arxiv.org/abs/2504.19472>.
- Llama Team. The llama 3 herd of models, 2024. URL <https://arxiv.org/abs/2407.21783>.
- Llama Team. The llama 4 herd: The beginning of a new era of natively multimodal ai innovation, 2025. URL <https://ai.meta.com/blog/llama-4-multimodal-intelligence/>.
- Shayne Longpre, Yi Lu, and Joachim Daiber. Mkqa: A linguistically diverse benchmark for multilingual open domain question answering. *Transactions of the Association for Computational Linguistics*, 9:1389–1406, 2021.
- Holy Lovenia, Rahmad Mahendra, Salsabil Maulana Akbar, Lester James V Miranda, Jennifer Santoso, Elyanah Aco, Akhdan Fadhilah, Jonibek Mansurov, Joseph Marvin Imperial, Onno P Kampman, et al. Seacrowd: A multilingual multimodal data hub and benchmark suite for south-east asian languages. *arXiv preprint arXiv:2406.10118*, 2024.
- Hongyuan Lu, Haoyang Huang, Dongdong Zhang, Haoran Yang, Wai Lam, and Furu Wei. Chain-of-dictionary prompting elicits translation in large language models. *arXiv preprint*

arXiv:2305.06575, 2023.

Aman Madaan, Shuyan Zhou, Uri Alon, Yiming Yang, and Graham Neubig. Language models of code are few-shot commonsense learners. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 1384–1403. Association for Computational Linguistics, December 2022. URL <https://aclanthology.org/2022.emnlp-main.90>.

Viorica Marian and Ulric Neisser. Language-dependent recall of autobiographical memories. *Journal of Experimental Psychology: General*, 129(3):361, 2000.

John McCarthy. *Programs with common sense*. RLE and MIT Computation Center, 1960.

Yasmin Moslem, Rejwanul Haque, John D. Kelleher, and Andy Way. Adaptive machine translation with large language models. In *Proceedings of the 24th Annual Conference of the European Association for Machine Translation*, pages 227–237. European Association for Machine Translation, June 2023. URL <https://aclanthology.org/2023.eamt-1.22/>.

Niklas Muennighoff, Thomas Wang, Lintang Sutawika, Adam Roberts, Stella Biderman, Teven Le Scao, M Saiful Bari, Sheng Shen, Zheng-Xin Yong, Hailey Schoelkopf, et al. Crosslingual generalization through multitask finetuning. *arXiv preprint arXiv:2211.01786*, 2022.

Junho Myung, Nayeon Lee, Yi Zhou, Jiho Jin, Rifki Putri, Dimosthenis Antypas, Hsuvas Borkakoty, Eunsu Kim, Carla Perez-Almendros, Abinew Ali Ayele, et al. Blend: A benchmark for llms on everyday knowledge in diverse cultures and languages. *Advances in Neural Information Processing Systems*, 37:78104–78146, 2024.

Tarek Naous, Michael J Ryan, Alan Ritter, and Wei Xu. Having beer after prayer? measuring cultural bias in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16366–16393, 2024.

Erik Neveu. Government, the state, and media. *The sage handbook of media studies*, pages 331–350, 2004.

Allen Newell and Herbert Simon. The logic theory machine—a complex information processing system. *IRE Transactions on information theory*, 2(3):61–79, 1956.

Josh Nguyen*, Upasana Dutta*, Samar Haider*, Bryan Li*, Neil Fasching*, and Duncan Watts. Exploring political attitudes under selective media consumption with large language models. *The 11th International Conference on Computational Social Science (IC2S2)*, 2025.

Ansong Ni, Jeevana Priya Inala, Chenglong Wang, Oleksandr Polozov, Christopher Meek, Dragomir R. Radev, and Jianfeng Gao. Learning math reasoning from self-sampled correct and partially-correct solutions. In *International Conference on Learning Representations*, 2022. URL <https://api.semanticscholar.org/CorpusID:257019561>.

- Xenia Ohmer, Elia Bruni, and Dieuwke Hupkes. Separating form and meaning: Using self-consistency to quantify task understanding across multiple senses. In *Proceedings of the Third Workshop on Natural Language Generation, Evaluation, and Metrics (GEM)*, pages 258–276. Association for Computational Linguistics, December 2023. URL <https://aclanthology.org/2023.gem-1.22/>.
- Jessica Ojo, Odunayo Ogundepo, Akintunde Oladipo, Kelechi Ogueji, Jimmy Lin, Pontus Stenertorp, and David Ifeoluwa Adelani. Afrobench: How good are large language models on african languages? In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 19048–19095, 2025.
- OpenAI. Chatgpt: Optimizing language models for dialogue, 2022. URL <https://www.aclweb.org/anthology/2020.emnlp-demos.6>.
- OpenAI. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- OpenAI. Multilingual massive multitask language understanding (mmmlu), 2024a. URL <https://huggingface.co/datasets/openai/MMMLU>.
- OpenAI. New embedding models and api updates. <https://openai.com/index/new-embedding-models-and-api-updates/>, 2024b. Accessed: 2024-08-26.
- OpenEuroLLM. A series of foundation models for transparent ai in europe, 2025. URL <https://openeurollm.eu>.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, volume 35, pages 27730–27744. Curran Associates, Inc., 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/file/b1efde53be364a73914f58805a001731-Paper-Conference.pdf.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pages 311–318, 2002.
- Ajay Patel, Bryan Li, Mohammad Sadegh Rasooli, Noah Constant, Colin Raffel, and Chris Callison-Burch. Bidirectional language models are also few-shot learners. *The Eleventh International Conference on Learning Representations (ICLR)*, 2023.
- Guilherme Penedo, Hynek Kydlíček, Anton Lozhkov, Margaret Mitchell, Colin A Raffel, Leandro Von Werra, Thomas Wolf, et al. The fineweb datasets: Decanting the web for the finest text data at scale. *Advances in Neural Information Processing Systems*, 37:30811–30849, 2024.

- Guilherme Penedo, Hynek Kydlíček, Vinko Sabolčec, Bettina Messmer, Negar Foroutan, Amir Hossein Kargaran, Colin Raffel, Martin Jaggi, Leandro Von Werra, and Thomas Wolf. Fineweb2: One pipeline to scale them all—adapting pre-training data processing to every language. *arXiv preprint arXiv:2506.20920*, 2025.
- Ethan Perez, Sam Ringer, Kamile Lukosiute, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, et al. Discovering language model behaviors with model-written evaluations. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13387–13434. Association for Computational Linguistics, July 2023. URL <https://aclanthology.org/2023.findings-acl.847/>.
- Fabio Petroni, Tim Rocktäschel, Sebastian Riedel, Patrick Lewis, Anton Bakhtin, Yuxiang Wu, and Alexander Miller. Language models as knowledge bases? In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2463–2473. Association for Computational Linguistics, November 2019. URL <https://aclanthology.org/D19-1250/>.
- Telmo Pires, Eva Schlinger, and Dan Garrette. How multilingual is multilingual BERT? In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4996–5001, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1493. URL <https://aclanthology.org/P19-1493/>.
- Maja Popović. chrF: character n-gram f-score for automatic mt evaluation. In *Proceedings of the tenth workshop on statistical machine translation*, pages 392–395, 2015.
- Matt Post and David Vilar. Fast lexically constrained decoding with dynamic beam allocation for neural machine translation. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1314–1324. Association for Computational Linguistics, June 2018. URL <https://aclanthology.org/N18-1119/>.
- Jirui Qi, Raquel Fernández, and Arianna Bisazza. Cross-lingual consistency of factual knowledge in multilingual language models. In *The 2023 Conference on Empirical Methods in Natural Language Processing*, 2023.
- Alec Radford, Karthik Narasimhan, Tim Salimans, Ilya Sutskever, et al. Improving language understanding by generative pre-training, 2018.
- Revanth Gangi Reddy, Yi Ren Fung, Qi Zeng, Manling Li, Ziqi Wang, Paul Sullivan, and Heng Ji. Smartbook: AI-assisted situation report generation. *arXiv preprint arXiv:2303.14337*, 2023.
- Ricardo Rei, José G. C. de Souza, Duarte Alves, Chrysoula Zerva, Ana C Farinha, Taisiya Glushkova, Alon Lavie, Luisa Coheur, and André F. T. Martins. COMET-22: Unbabel-IST 2022 submission for the metrics shared task. In *Proceedings of the Seventh Conference on Machine Translation (WMT)*, pages 578–585. Association for Computational Linguistics, December

2022. URL <https://aclanthology.org/2022.wmt-1.52/>.
- Danilo Neves Ribeiro, Shen Wang, Xiaofei Ma, Henghui Zhu, Rui Dong, Deguang Kong, Juliette Burger, Anjelica Ramos, William Yang Wang, George Karypis, et al. Street: A multi-task structured reasoning and explanation benchmark. In *The Eleventh International Conference on Learning Representations*, 2022.
- Sara Ruberg. When a.i. fails the language test, who is left out of the conversation? *The New York Times*, 07 2024. URL <https://www.nytimes.com/2024/07/26/technology/ai-language-gap.html>.
- Sebastian Ruder, Noah Constant, Jan Botha, Aditya Siddhant, Orhan Firat, Jinlan Fu, Pengfei Liu, Junjie Hu, Dan Garrette, Graham Neubig, et al. Xtreme-r: Towards more challenging and nuanced multilingual evaluation. *arXiv preprint arXiv:2104.07412*, 2021.
- Sebastian Ruder, Jonathan H Clark, Alexander Gutkin, Mihir Kale, Min Ma, Massimo Nicosia, Shruti Rijhwani, Parker Riley, Jean-Michel A Sarr, Xinyi Wang, et al. Xtreme-up: A user-centric scarce-data benchmark for under-represented languages. *arXiv preprint arXiv:2305.11938*, 2023.
- Masoud Jalili Sabet, Philipp Dufter, François Yvon, and Hinrich Schütze. Simalign: High quality word alignments without parallel training data using static and contextualized embeddings. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 1627–1643, 2020.
- Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. *Advances in Neural Information Processing Systems*, 36:68539–68551, 2023.
- Kirill Semenov, Vilém Zouhar, Tom Kocmi, Dongdong Zhang, Wangchunshu Zhou, and Yuchen Eleanor Jiang. Findings of the WMT 2023 shared task on machine translation with terminologies. In *Proceedings of the Eighth Conference on Machine Translation*, pages 663–671. Association for Computational Linguistics, December 2023. URL <https://aclanthology.org/2023.wmt-1.54/>.
- Nikhil Sharma, Kenton Murray, and Ziang Xiao. Faux polyglot: A study on information disparity in multilingual large language models. *arXiv preprint arXiv:2407.05502*, 2024.
- Freda Shi, Mirac Suzgun, Markus Freitag, Xuezhi Wang, Suraj Srivats, Soroush Vosoughi, Hyung Won Chung, Yi Tay, Sebastian Ruder, Denny Zhou, et al. Language models are multilingual chain-of-thought reasoners. In *The Eleventh International Conference on Learning Representations*, 2022.
- Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H Chi, Nathanael Schärli, and Denny Zhou. Large language models can be easily distracted by irrelevant context. In *International Conference on Machine Learning*, pages 31210–31227. PMLR, 2023.

- William Siler and James J Buckley. *Fuzzy expert systems and fuzzy reasoning*. John Wiley & Sons, 2005.
- Harman Singh, Nitish Gupta, Shikhar Bharadwaj, Dinesh Tewari, and Partha Talukdar. Indicgen-bench: A multilingual benchmark to evaluate generation capabilities of llms on indic languages. *arXiv preprint arXiv:2404.16816*, 2024.
- Shivalika Singh, Angelika Romanou, Cl  mentine Fourier, David Ifeoluwa Adelani, Jian Gang Ngui, Daniel Vila-Suero, Peerat Limkonchotiwat, Kelly Marchisio, Wei Qi Leong, Yosephine Susanto, Raymond Ng, Shayne Longpre, Sebastian Ruder, Wei-Yin Ko, Antoine Bosselut, Alice Oh, Andre Martins, Leshem Choshen, Daphne Ippolito, Enzo Ferrante, Marzieh Fadaee, Beyza Ermis, and Sara Hooker. Global MMLU: Understanding and Addressing Cultural and Linguistic Biases in Multilingual Evaluation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 18761–18799. Association for Computational Linguistics, July 2025. ISBN 9798891762510.
- Luca Soldaini, Rodney Kinney, Akshita Bhagia, Dustin Schwenk, et al. Dolma: an open corpus of three trillion tokens for language model pretraining research. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15725–15788, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.840. URL <https://aclanthology.org/2024.acl-long.840/>.
- Harold Somers. Review article: Example-based machine translation. *Machine Translation*, 14(2): 113–157, 1999. ISSN 09226567, 15730573. URL <http://www.jstor.org/stable/40007076>.
- Zhongxiang Sun, Xiaoxue Zang, Kai Zheng, Yang Song, Jun Xu, Xiao Zhang, Weijie Yu, and Han Li. Redeeep: Detecting hallucination in retrieval-augmented generation via mechanistic interpretability. *arXiv preprint arXiv:2410.11414*, 2024.
- Mirac Suzgun, Nathan Scales, Nathanael Scharli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc V. Le, Ed Huai hsin Chi, Denny Zhou, and Jason Wei. Challenging big-bench tasks and whether chain-of-thought can solve them. In *Annual Meeting of the Association for Computational Linguistics*, 2022. URL <https://api.semanticscholar.org/CorpusID:252917648>.
- Weiting Tan, Haoran Xu, Lingfeng Shen, Shuyue Stella Li, Kenton Murray, Philipp Koehn, Benjamin Van Durme, and Yunmo Chen. Narrowing the gap between zero- and few-shot machine translation by matching styles. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 490–502. Association for Computational Linguistics, June 2024. URL <https://aclanthology.org/2024.findings-naacl.33/>.
- Cemal Burak Tansel. *Authoritarian neoliberalism and democratic backsliding in turkey: Beyond the narratives of progress*, 2018.

- Yan Tao, Olga Viberg, Ryan S Baker, and Rene F Kizilcec. Auditing and mitigating cultural bias in llms. *arXiv preprint arXiv:2311.14096*, 2023.
- Nandan Thakur, Luiz Bonifacio, Xinyu Zhang, Odunayo Ogundepo, Ehsan Kamalloo, David Alfonso-Hermelo, Xiaoguang Li, Qun Liu, Boxing Chen, Mehdi Rezagholizadeh, and Jimmy Lin. Nomiracl: Knowing when you don’t know for robust multilingual retrieval-augmented generation, 2024. URL <https://arxiv.org/abs/2312.11361>.
- Marion Thaler, Abdullatif Köksal, Alina Leidinger, Anna Korhonen, and Hinrich Schütze. How far can bias go?—tracing bias from pretraining data to alignment. *arXiv preprint arXiv:2411.19240*, 2024.
- Brian Thompson, Jeremy Gwinnup, Huda Khayrallah, Kevin Duh, and Philipp Koehn. Overcoming catastrophic forgetting during domain adaptation of neural machine translation. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 2062–2068. Association for Computational Linguistics, June 2019. URL <https://aclanthology.org/N19-1209/>.
- G Richard Tucker. A global perspective on bilingualism and bilingual education. *Washington, DC*, 1999.
- Ahmet Üstün, Viraat Aryabumi, Zheng Yong, Wei-Yin Ko, Daniel D’souza, Gbemileke Onilude, Neel Bhandari, Shivalika Singh, Hui-Lee Ooi, Amr Kayid, Freddie Vargus, Phil Blunsom, Shayne Longpre, Niklas Muennighoff, Marzieh Fadaee, Julia Kreutzer, and Sara Hooker. Aya model: An instruction finetuned open-access multilingual language model. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024.
- Ashish Vaswani, Noam M. Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Neural Information Processing Systems*, 2017. URL <https://api.semanticscholar.org/CorpusID:13756489>.
- David Vilar, Markus Freitag, Colin Cherry, Jiaming Luo, Viresh Ratnakar, and George Foster. Prompting PaLM for translation: Assessing strategies and performance. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15406–15427. Association for Computational Linguistics, July 2023. URL <https://aclanthology.org/2023.acl-long.859/>.
- Christopher Walker and Robert W Orttung. Breaking the news: The role of state-run media. *Journal of Democracy*, 25(1):71–85, 2014.
- Jenny S Wang, Samar Haider, Amir Tohidi, Anushkaa Gupta, Yuxuan Zhang, Chris Callison-Burch, David Rothschild, and Duncan J Watts. Media bias detector: Designing and implementing a tool for real-time selection and framing bias analysis in news coverage. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, pages 1–27, 2025a.

- Jiaan Wang, Yunlong Liang, Zengkui Sun, Yuxuan Cao, Jiarong Xu, and Fandong Meng. Cross-lingual knowledge editing in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11676–11686, 2024a.
- Mingyang Wang, Heike Adel, Lukas Lange, Yihong Liu, Ercong Nie, Jannik Strötgen, and Hinrich Schuetze. Lost in Multilinguality: Dissecting Cross-lingual Factual Inconsistency in Transformer Language Models. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5075–5094. Association for Computational Linguistics, July 2025b. ISBN 9798891762510.
- Shuo Wang, Peng Li, Zhixing Tan, Zhaopeng Tu, Maosong Sun, and Yang Liu. A template-based method for constrained neural machine translation. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3665–3679. Association for Computational Linguistics, December 2022a. URL <https://aclanthology.org/2022.emnlp-main.240/>.
- Song Wang, Yaochen Zhu, Haochen Liu, Zaiyi Zheng, Chen Chen, and Jundong Li. Knowledge editing for large language models: A survey. *ACM Computing Surveys*, 57(3):1–37, 2024b.
- Xuezhi Wang, Haohan Wang, and Diyi Yang. Measure and Improve Robustness in NLP Models: A Survey. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4569–4586. Association for Computational Linguistics, 2022b.
- Duncan J Watts, David M Rothschild, and Markus Mobius. Measuring the news and its impact on democracy. *Proceedings of the National Academy of Sciences*, 118(15):e1912443118, 2021.
- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, et al. Emergent abilities of large language models. *Transactions on Machine Learning Research*, 2022a.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35:24824–24837, 2022b.
- Orion Weller, Aleem Khan, Nathaniel Weir, Dawn Lawrie, and Benjamin Van Durme. Defending against disinformation attacks in open-domain question answering. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 402–417, 2024.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander M. Rush. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 Conference on Empirical Methods in Nat-*

- ural Language Processing: System Demonstrations*, pages 38–45. Association for Computational Linguistics, October 2020. URL <https://www.aclweb.org/anthology/2020.emnlp-demos.6>.
- Ellen Woolford. Bilingual code-switching and syntactic theory. *Linguistic inquiry*, 14(3):520–536, 1983.
- Hua Wu, Haifeng Wang, and Chengqing Zong. Domain adaptation for statistical machine translation with domain dictionary and monolingual corpora. In *Proceedings of the 22nd International Conference on Computational Linguistics (Coling 2008)*, pages 993–1000. Coling 2008 Organizing Committee, August 2008. URL <https://aclanthology.org/C08-1125/>.
- Suhang Wu, Jialong Tang, Baosong Yang, Ante Wang, Kaidi Jia, Jiawei Yu, Junfeng Yao, and Jinsong Su. Not all languages are equal: Insights into multilingual retrieval-augmented generation. *arXiv preprint arXiv:2410.21970*, 2024.
- Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, and Yu Su. Adaptive chameleon or stubborn sloth: Revealing the behavior of large language models in knowledge conflicts. In *The Twelfth International Conference on Learning Representations*, 2023.
- Rongwu Xu, Zehan Qi, Zhijiang Guo, Cunxiang Wang, Hongru Wang, Yue Zhang, and Wei Xu. Knowledge Conflicts for LLMs: A Survey. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 8541–8565. Association for Computational Linguistics, 2024. URL <https://aclanthology.org/2024.emnlp-main.486/>.
- Linting Xue, Noah Constant, Adam Roberts, Mihir Kale, Rami Al-Rfou, Aditya Siddhant, Aditya Barua, and Colin Raffel. mt5: A massively multilingual pre-trained text-to-text transformer. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 483–498, 2021.
- Masaru Yamada. The effect of translation memory databases on productivity. *Translation research projects*, 3:63–73, 2011.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jing Zhou, Jingren Zhou, Junyang Lin, Kai Dang, Keqin Bao, Kexin Yang, Le Yu, Lianghao Deng, Mei Li, Mingfeng Xue, Mingze Li, Pei Zhang, Peng Wang, Qin Zhu, Rui Men, Ruize Gao, Shixuan Liu, Shuang Luo, Tianhao Li, Tianyi Tang, Wenbiao Yin, Xingzhang Ren, Xinyu Wang, Xinyu Zhang, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yinger Zhang, Yu Wan, Yuqiong Liu, Zekun Wang, Zeyu Cui, Zhenru Zhang, Zhipeng Zhou, and Zihan Qiu. Qwen3 technical report, 2025. URL <https://arxiv.org/abs/2505.09388>.
- Zhewei Yao, Reza Yazdani Aminabadi, Olatunji Ruwase, Samyam Rajbhandari, Xiaoxia Wu, Ammar Ahmad Awan, Jeff Rasley, Minjia Zhang, Conglong Li, Connor Holmes, Zhongzhu Zhou, Michael Wyatt, Molly Smith, Lev Kurilenko, Heyang Qin, Masahiro Tanaka, Shuai Che, Shuai-

- wen Leon Song, and Yuxiong He. DeepSpeed-Chat: Easy, Fast and Affordable RLHF Training of ChatGPT-like Models at All Scales. *arXiv preprint arXiv:2308.01320*, 2023.
- Da Yin, Hritik Bansal, Masoud Monajatipoor, Liunian Harold Li, and Kai-Wei Chang. GeoM-LAMA: Geo-diverse commonsense probing on multilingual pre-trained language models. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 2039–2055. Association for Computational Linguistics, December 2022. URL <https://aclanthology.org/2022.emnlp-main.132/>.
- Yu Yu, Abdul Rafae Khan, and Jia Xu. Measuring Robustness for NLP. In *Proceedings of the 29th International Conference on Computational Linguistics*, pages 3908–3916. International Committee on Computational Linguistics, October 2022.
- Biao Zhang, Barry Haddow, and Alexandra Birch. Prompting Large Language Model for Machine Translation: A Case Study. In *Proceedings of the 40th International Conference on Machine Learning*, pages 41092–41110. PMLR, 2023a. URL <https://proceedings.mlr.press/v202/zhang23m.html>.
- Li Zhang, Hainiu Xu, Yue Yang, Shuyan Zhou, Weiqiu You, Manni Arora, and Chris Callison-Burch. Causal reasoning of entities and events in procedural texts. In *Findings of the Association for Computational Linguistics: EACL 2023*, pages 415–431. Association for Computational Linguistics, May 2023b. URL <https://aclanthology.org/2023.findings-eacl.31>.
- Michael Zhang and Eunsol Choi. Mitigating temporal misalignment by discarding outdated facts. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 14213–14226, 2023.
- Xinyu Zhang, Nandan Thakur, Odunayo Ogundepo, Ehsan Kamalloo, David Alfonso-Hermelo, Xiaoguang Li, Qun Liu, Mehdi Rezagholizadeh, and Jimmy Lin. MIRACL: A Multilingual Retrieval Dataset Covering 18 Diverse Languages. *Transactions of the Association for Computational Linguistics*, 11:1114–1131, 09 2023c. ISSN 2307-387X. URL https://doi.org/10.1162/tacl_a_00595.
- Yiran Zhao, Wenxuan Zhang, Guizhen Chen, Kenji Kawaguchi, and Lidong Bing. How do Large Language Models Handle Multilingualism?, November 2024.
- Wanjuan Zhong, Siyuan Wang, Duyu Tang, Zenan Xu, Daya Guo, Yining Chen, Jiahai Wang, Jian Yin, Ming Zhou, and Nan Duan. Analytical reasoning of text. In Marine Carpuat, Marie-Catherine de Marneffe, and Ivan Vladimir Meza Ruiz, editors, *Findings of the Association for Computational Linguistics: NAACL 2022*, pages 2306–2319, Seattle, United States, July 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.findings-naacl.177. URL <https://aclanthology.org/2022.findings-naacl.177/>.
- Wenhao Zhu, Hongyi Liu, Qingxiu Dong, Jingjing Xu, Shujian Huang, Lingpeng Kong, Jiajun Chen, and Lei Li. Multilingual machine translation with large language models: Empirical results and analysis. In Kevin Duh, Helena Gomez, and Steven Bethard, editors, *Findings of the Association*

for Computational Linguistics: NAACL 2024, pages 2765–2781, Mexico City, Mexico, June 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-naacl.176. URL <https://aclanthology.org/2024.findings-naacl.176/>.